OLAP with Tableau

Enrico Gallinucci <enrico.gallinucci@unibo.it>

DFM – Foodmart

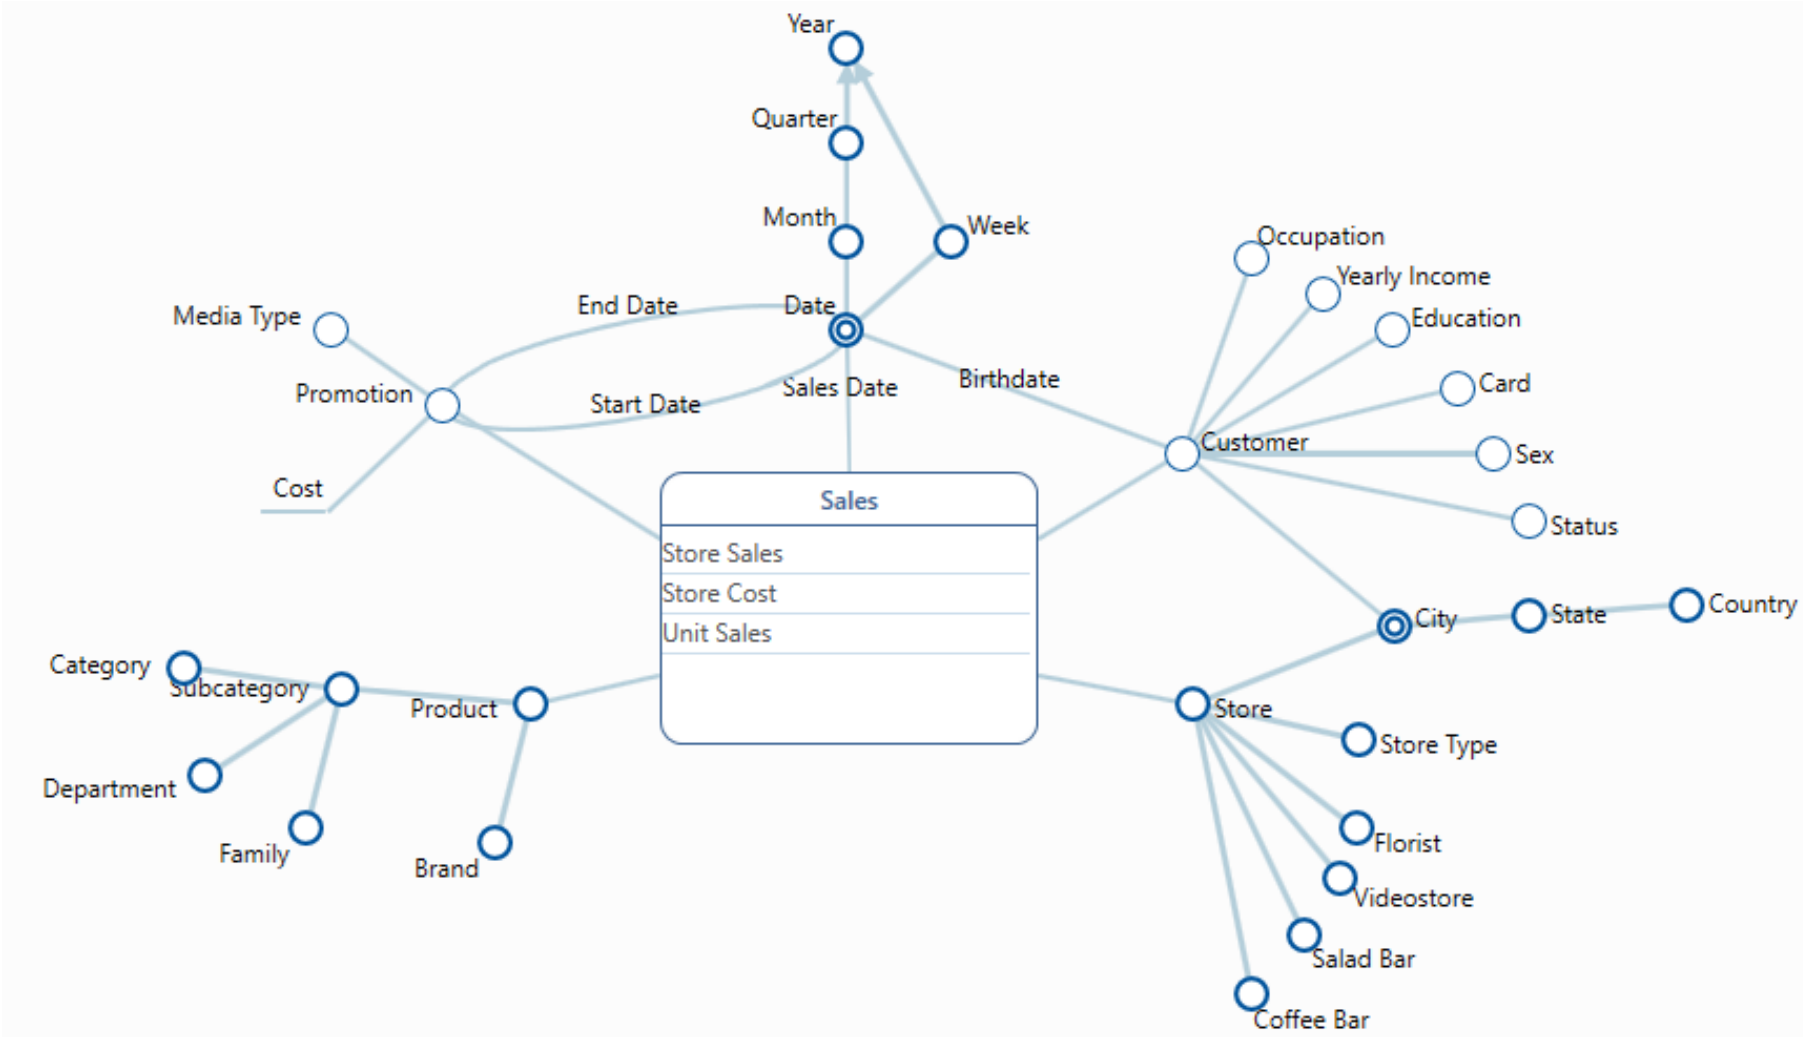

Database Structure – Foodmart (Sales)

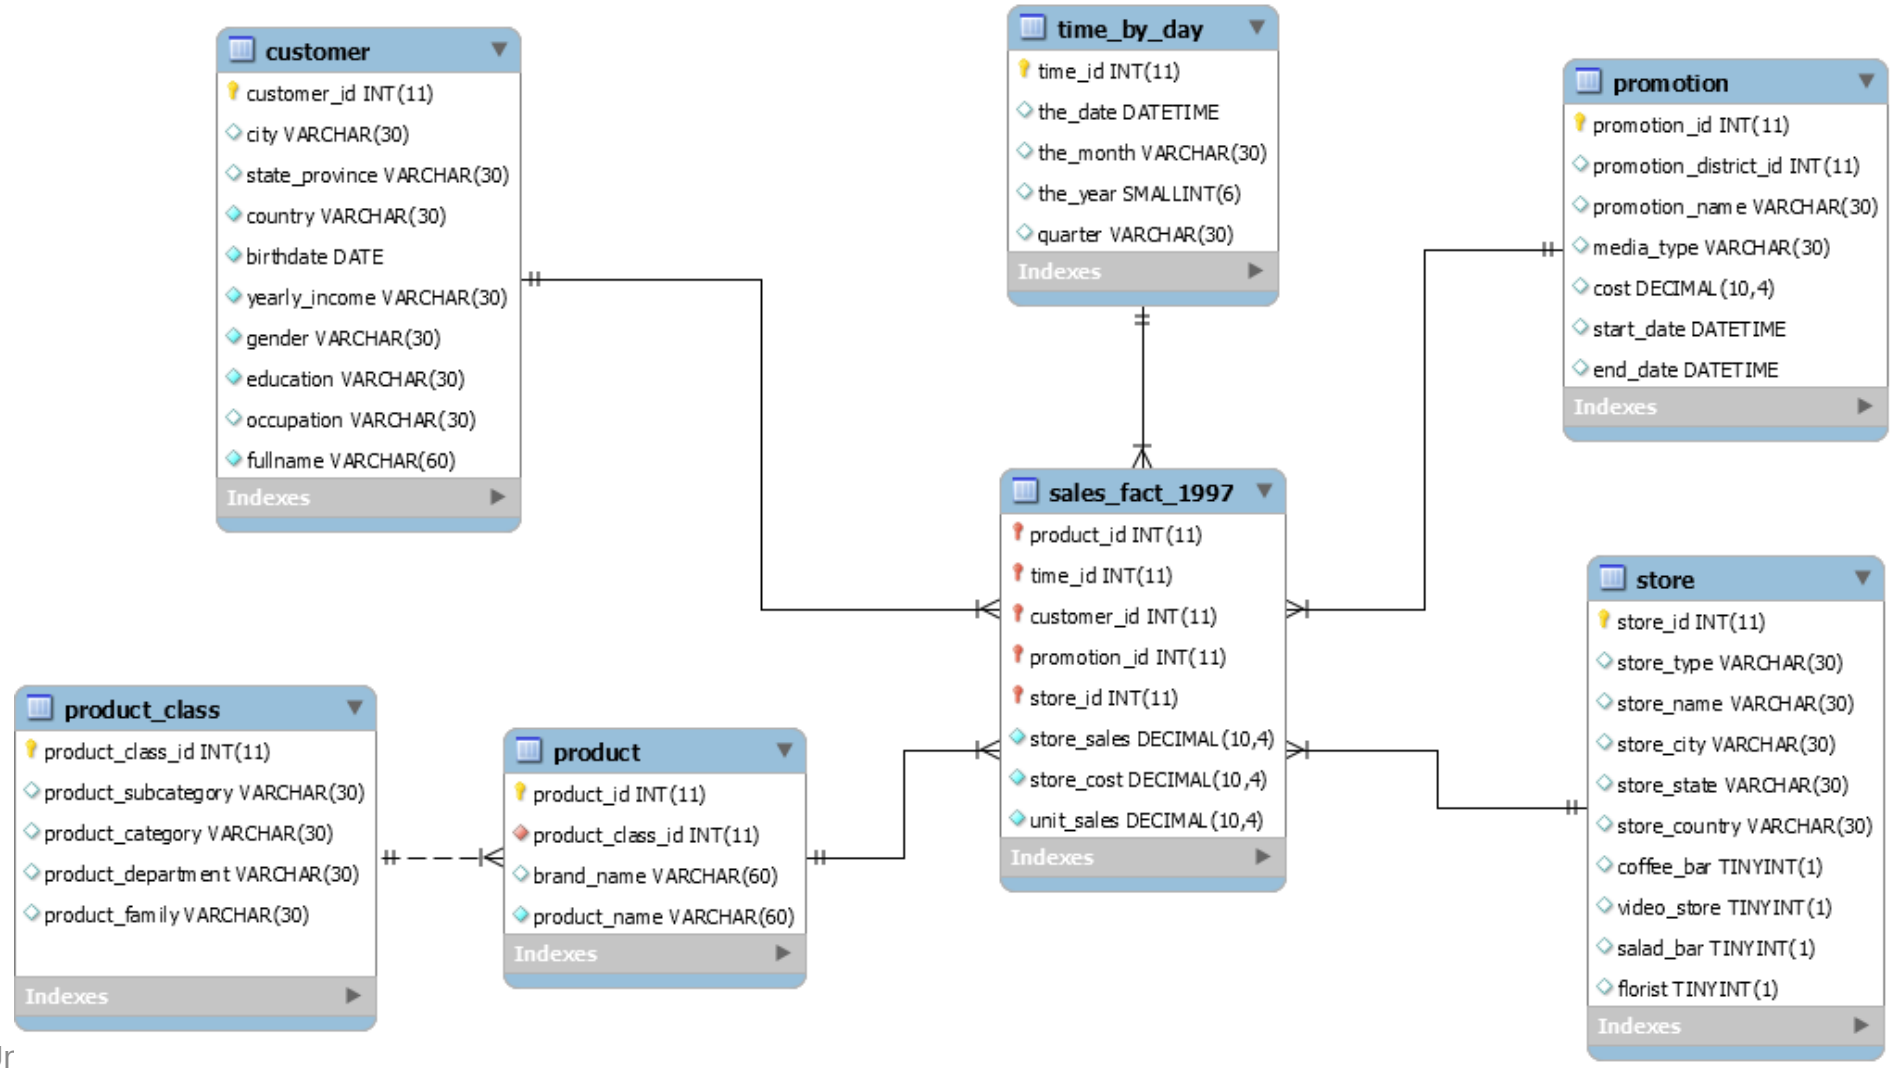

Database Structure – Foodmart (Sales)

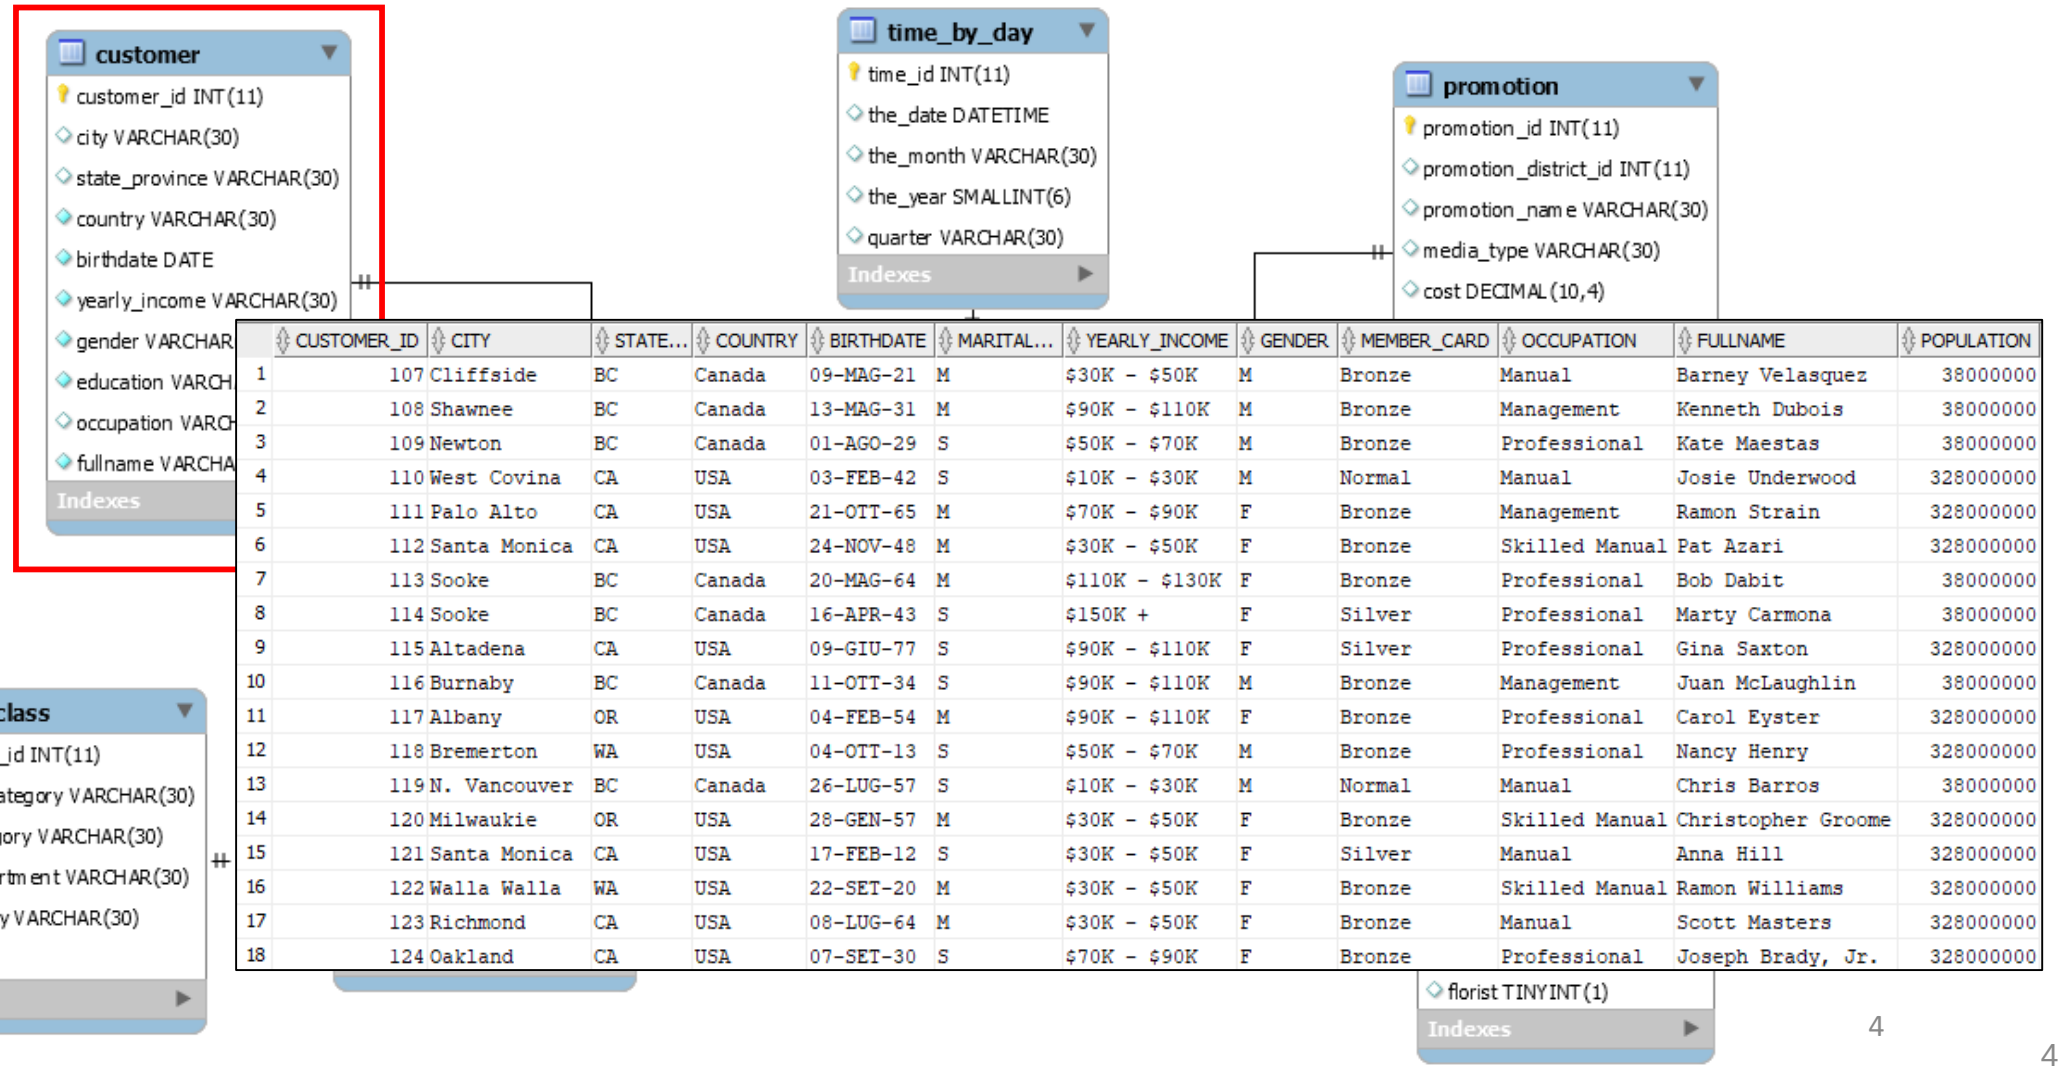

	CUSTOMER_ID	CITY	STATE...	COUNTRY	BIRTHDATE	MARITAL...	YEARLY_INCOME	GENDER	MEMBER_CARD	OCCUPATION	FULLNAME	POPULATION
1	107	Cliffside	BC	Canada	09-MAG-21	M	$30K - $50K	M	Bronze	Manual	Barney Velasquez	38000000
2	108	Shawnee	BC	Canada	13-MAG-31	M	$90K - $110K	M	Bronze	Management	Kenneth Dubois	38000000
3	109	Newton	BC	Canada	01-AGO-29	S	$50K - $70K	M	Bronze	Professional	Kate Maestas	38000000
4	110	West Covina	CA	USA	03-FEB-42	S	$10K - $30K	M	Normal	Manual	Josie Underwood	328000000
5	111	Palo Alto	CA	USA	21-OTT-65	M	$70K - $90K	F	Bronze	Management	Ramon Strain	328000000
6	112	Santa Monica	CA	USA	24-NOV-48	M	$30K - $50K	F	Bronze	Skilled Manual	Pat Azari	328000000
7	113	Sooke	BC	Canada	20-MAG-64	M	$110K - $130K	F	Bronze	Professional	Bob Dabit	38000000
8	114	Sooke	BC	Canada	16-APR-43	S	$150K +	F	Silver	Professional	Marty Carmona	38000000
9	115	Altadena	CA	USA	09-GIU-77	S	$90K - $110K	F	Silver	Professional	Gina Saxton	328000000
10	116	Burnaby	BC	Canada	11-OTT-34	S	$90K - $110K	M	Bronze	Management	Juan McLaughlin	38000000
11	117	Albany	OR	USA	04-FEB-54	M	$90K - $110K	F	Bronze	Professional	Carol Eyster	328000000
12	118	Bremerton	WA	USA	04-OTT-13	S	$50K - $70K	M	Bronze	Professional	Nancy Henry	328000000
13	119	N. Vancouver	BC	Canada	26-LUG-57	S	$10K - $30K	M	Normal	Manual	Chris Barros	38000000
14	120	Milwaukie	OR	USA	28-GEN-57	M	$30K - $50K	F	Bronze	Skilled Manual	Christopher Groome	328000000
15	121	Santa Monica	CA	USA	17-FEB-12	S	$30K - $50K	F	Silver	Manual	Anna Hill	328000000
16	122	Walla Walla	WA	USA	22-SET-20	M	$30K - $50K	F	Bronze	Skilled Manual	Ramon Williams	328000000
17	123	Richmond	CA	USA	08-LUG-64	M	$30K - $50K	F	Bronze	Manual	Scott Masters	328000000
18	124	Oakland	CA	USA	07-SET-30	S	$70K - $90K	F	Bronze	Professional	Joseph Brady, Jr.	328000000

4

Database Structure – Foodmart (Sales)

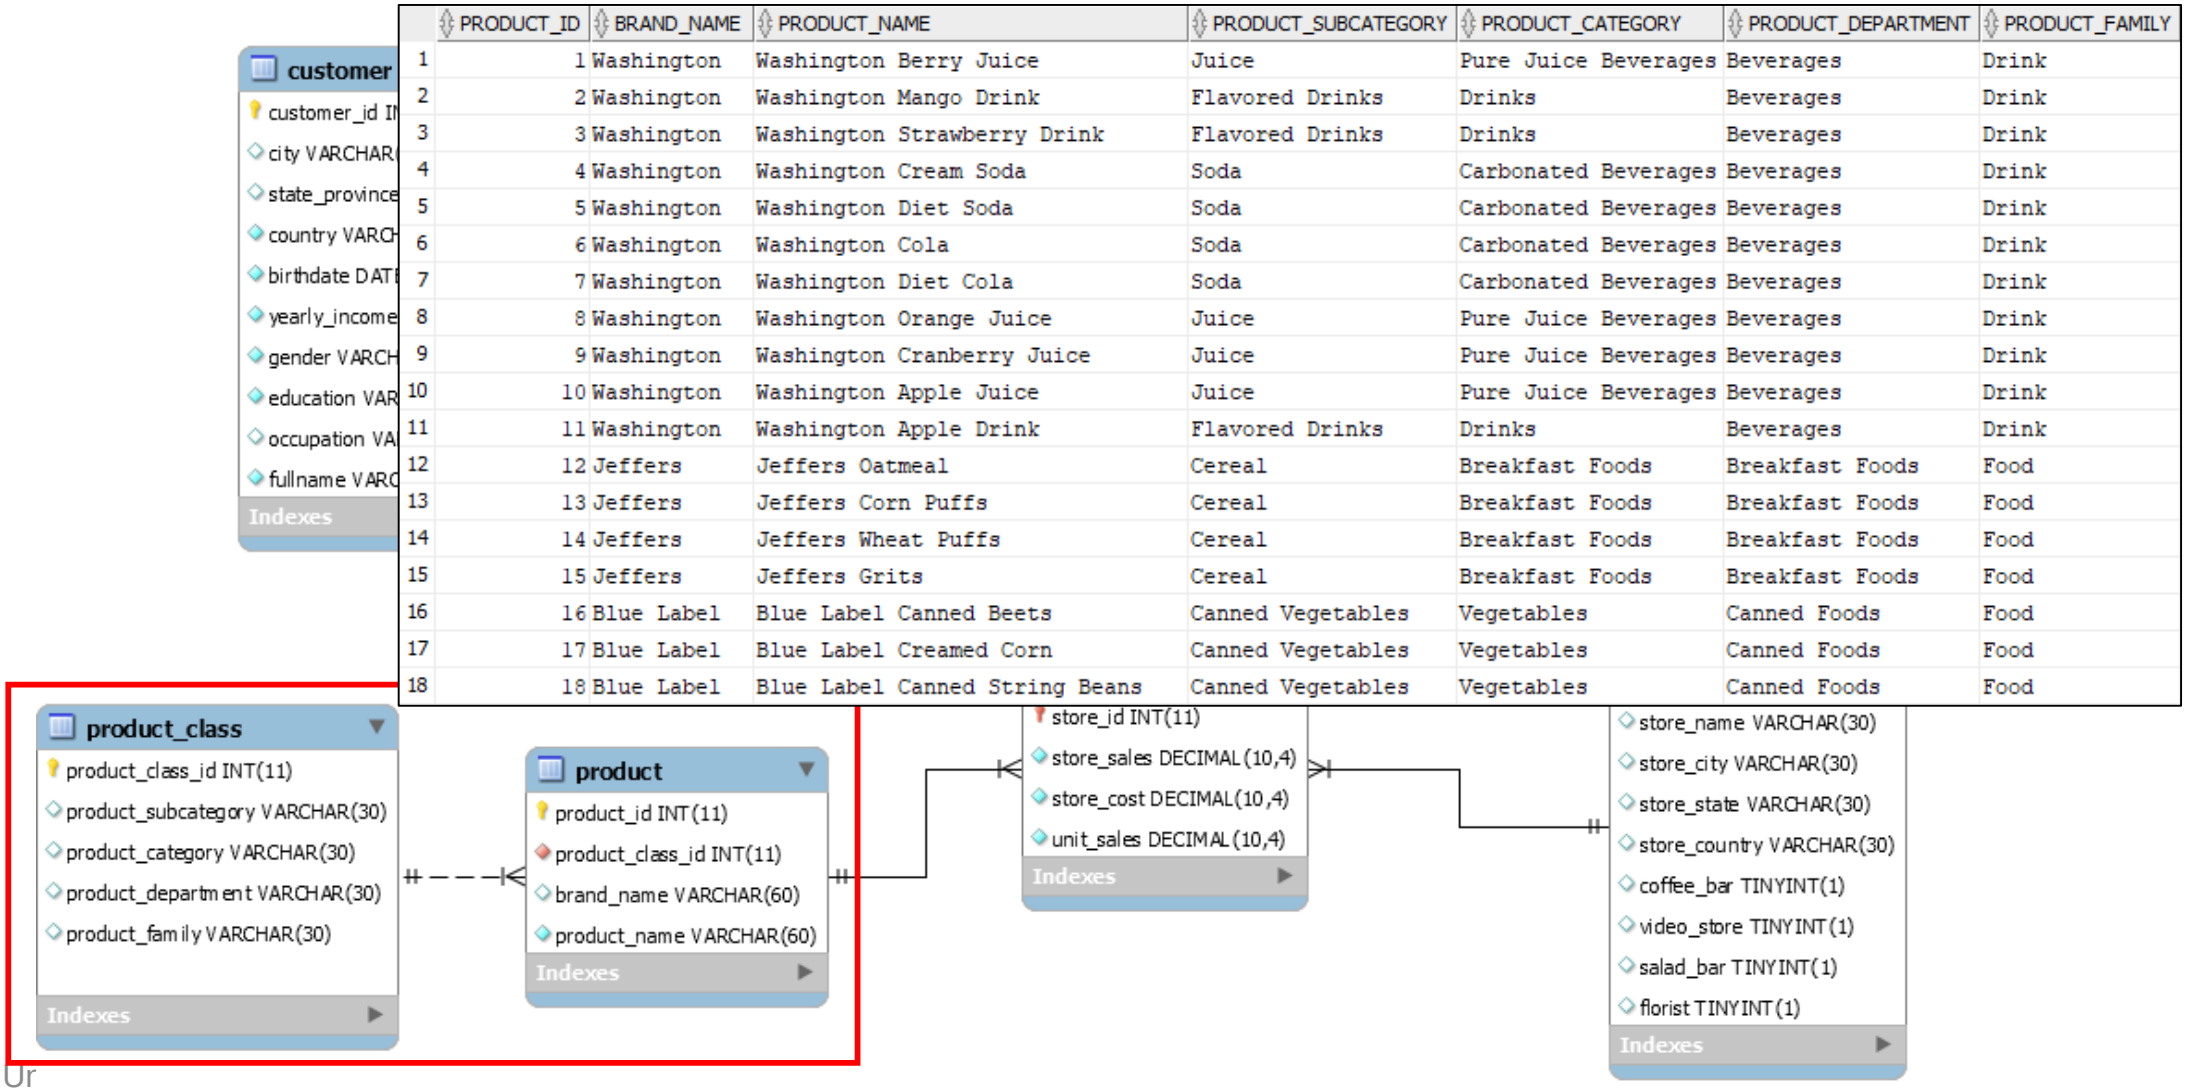

	PRODUCT_ID	BRAND_NAME	PRODUCT_NAME	PRODUCT_SUBCATEGORY	PRODUCT_CATEGORY	PRODUCT_DEPARTMENT	PRODUCT_FAMILY
1	1	Washington	Washington Berry Juice	Juice	Pure Juice Beverages	Beverages	Drink
2	2	Washington	Washington Mango Drink	Flavored Drinks	Drinks	Beverages	Drink
3	3	Washington	Washington Strawberry Drink	Flavored Drinks	Drinks	Beverages	Drink
4	4	Washington	Washington Cream Soda	Soda	Carbonated Beverages	Beverages	Drink
5	5	Washington	Washington Diet Soda	Soda	Carbonated Beverages	Beverages	Drink
6	6	Washington	Washington Cola	Soda	Carbonated Beverages	Beverages	Drink
7	7	Washington	Washington Diet Cola	Soda	Carbonated Beverages	Beverages	Drink
8	8	Washington	Washington Orange Juice	Juice	Pure Juice Beverages	Beverages	Drink
9	9	Washington	Washington Cranberry Juice	Juice	Pure Juice Beverages	Beverages	Drink
10	10	Washington	Washington Apple Juice	Juice	Pure Juice Beverages	Beverages	Drink
11	11	Washington	Washington Apple Drink	Flavored Drinks	Drinks	Beverages	Drink
12	12	Jeffers	Jeffers Oatmeal	Cereal	Breakfast Foods	Breakfast Foods	Food
13	13	Jeffers	Jeffers Corn Puffs	Cereal	Breakfast Foods	Breakfast Foods	Food
14	14	Jeffers	Jeffers Wheat Puffs	Cereal	Breakfast Foods	Breakfast Foods	Food
15	15	Jeffers	Jeffers Grits	Cereal	Breakfast Foods	Breakfast Foods	Food
16	16	Blue Label	Blue Label Canned Beets	Canned Vegetables	Vegetables	Canned Foods	Food
17	17	Blue Label	Blue Label Creamed Corn	Canned Vegetables	Vegetables	Canned Foods	Food
18	18	Blue Label	Blue Label Canned String Beans	Canned Vegetables	Vegetables	Canned Foods	Food

Database Structure – Foodmart (Sales)

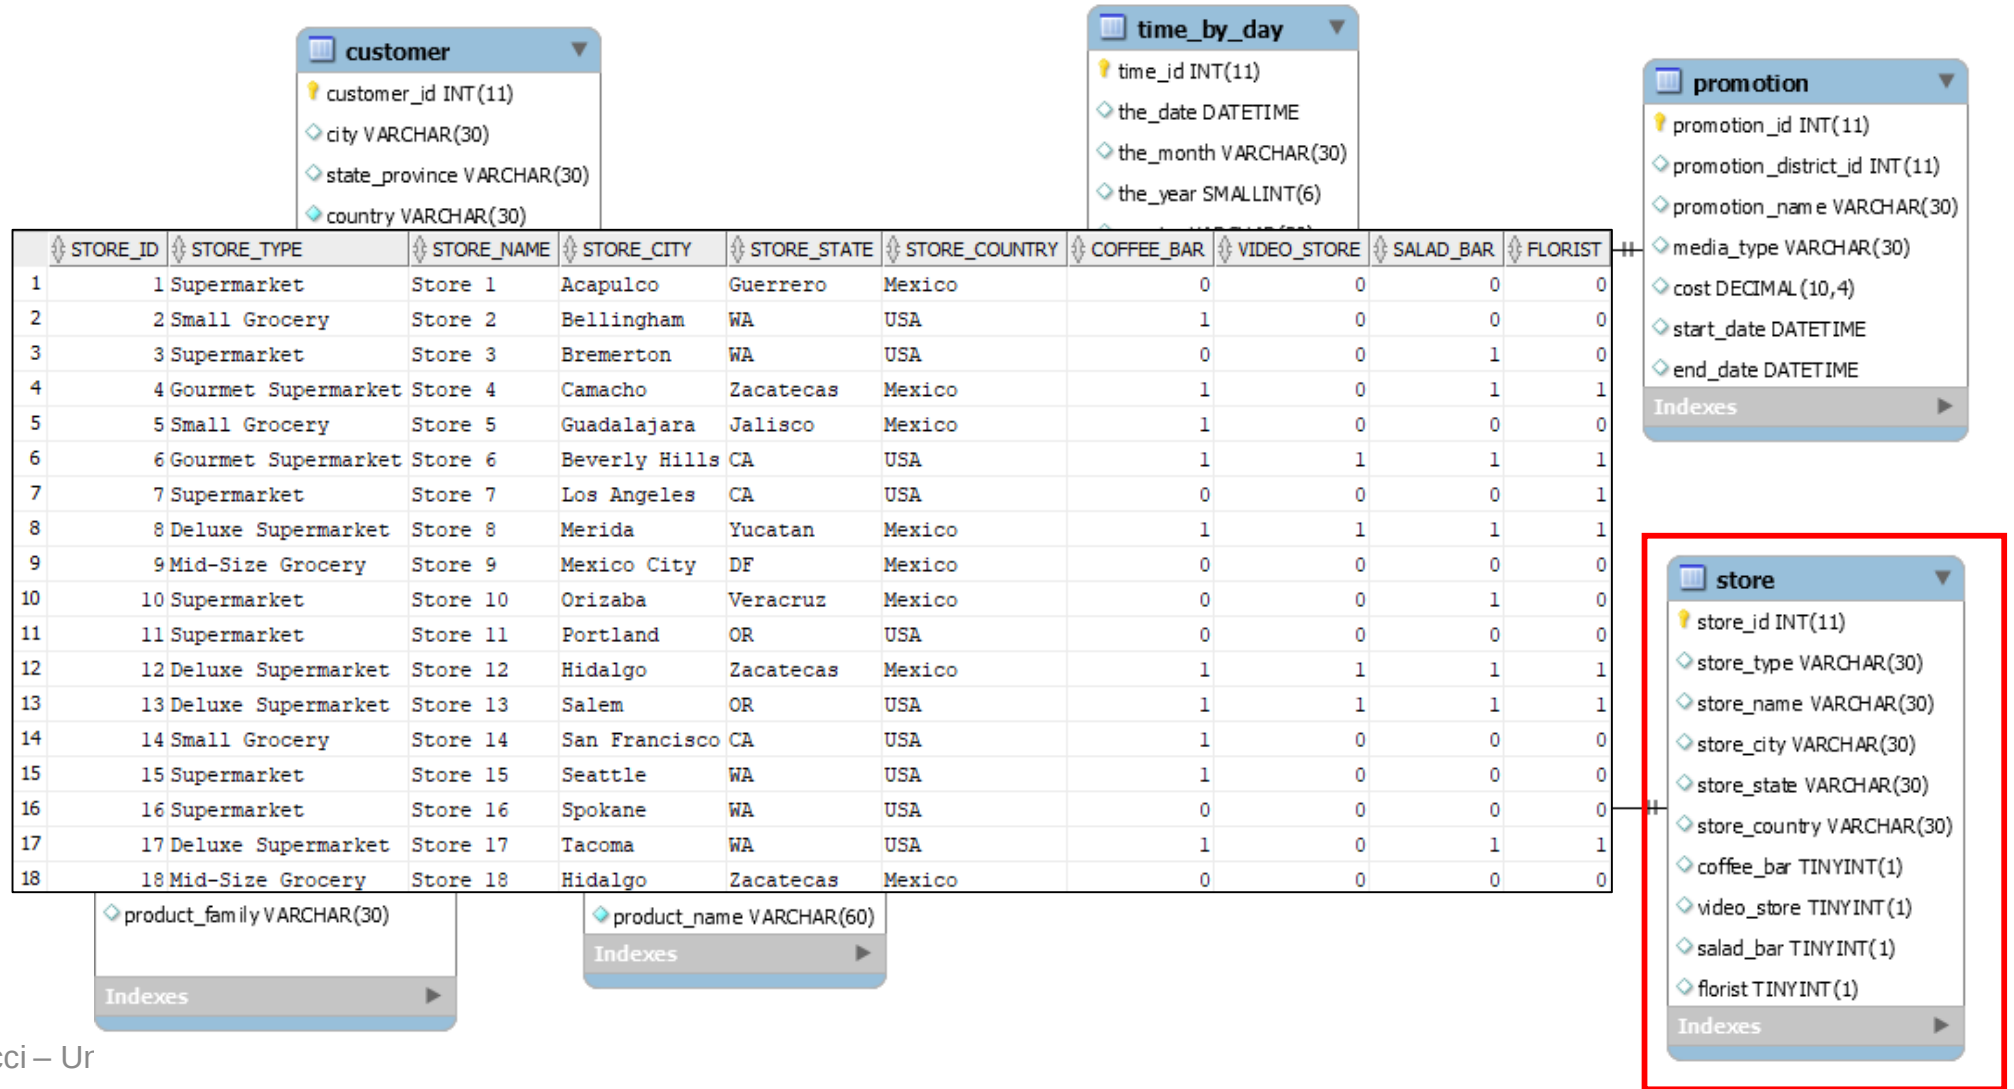

	STORE_ID	STORE_TYPE	STORE_NAME	STORE_CITY	STORE_STATE	STORE_COUNTRY	COFFEE_BAR	VIDEO_STORE	SALAD_BAR	FLORIST
1	1	Supermarket	Store 1	Acapulco	Guerrero	Mexico	0	0	0	0
2	2	Small Grocery	Store 2	Bellingham	WA	USA	1	0	0	0
3	3	Supermarket	Store 3	Bremerton	WA	USA	0	0	1	0
4	4	Gourmet Supermarket	Store 4	Camacho	Zacatecas	Mexico	1	0	1	1
5	5	Small Grocery	Store 5	Guadalajara	Jalisco	Mexico	1	0	0	0
6	6	Gourmet Supermarket	Store 6	Beverly Hills	CA	USA	1	1	1	1
7	7	Supermarket	Store 7	Los Angeles	CA	USA	0	0	0	1
8	8	Deluxe Supermarket	Store 8	Merida	Yucatan	Mexico	1	1	1	1
9	9	Mid-Size Grocery	Store 9	Mexico City	DF	Mexico	0	0	0	0
10	10	Supermarket	Store 10	Orizaba	Veracruz	Mexico	0	0	1	0
11	11	Supermarket	Store 11	Portland	OR	USA	0	0	0	0
12	12	Deluxe Supermarket	Store 12	Hidalgo	Zacatecas	Mexico	1	1	1	1
13	13	Deluxe Supermarket	Store 13	Salem	OR	USA	1	1	1	1
14	14	Small Grocery	Store 14	San Francisco	CA	USA	1	0	0	0
15	15	Supermarket	Store 15	Seattle	WA	USA	1	0	0	0
16	16	Supermarket	Store 16	Spokane	WA	USA	0	0	0	0
17	17	Deluxe Supermarket	Store 17	Tacoma	WA	USA	1	0	1	1
18	18	Mid-Size Grocery	Store 18	Hidalgo	Zacatecas	Mexico	0	0	0	0

Database Structure – Foodmart (Sales)

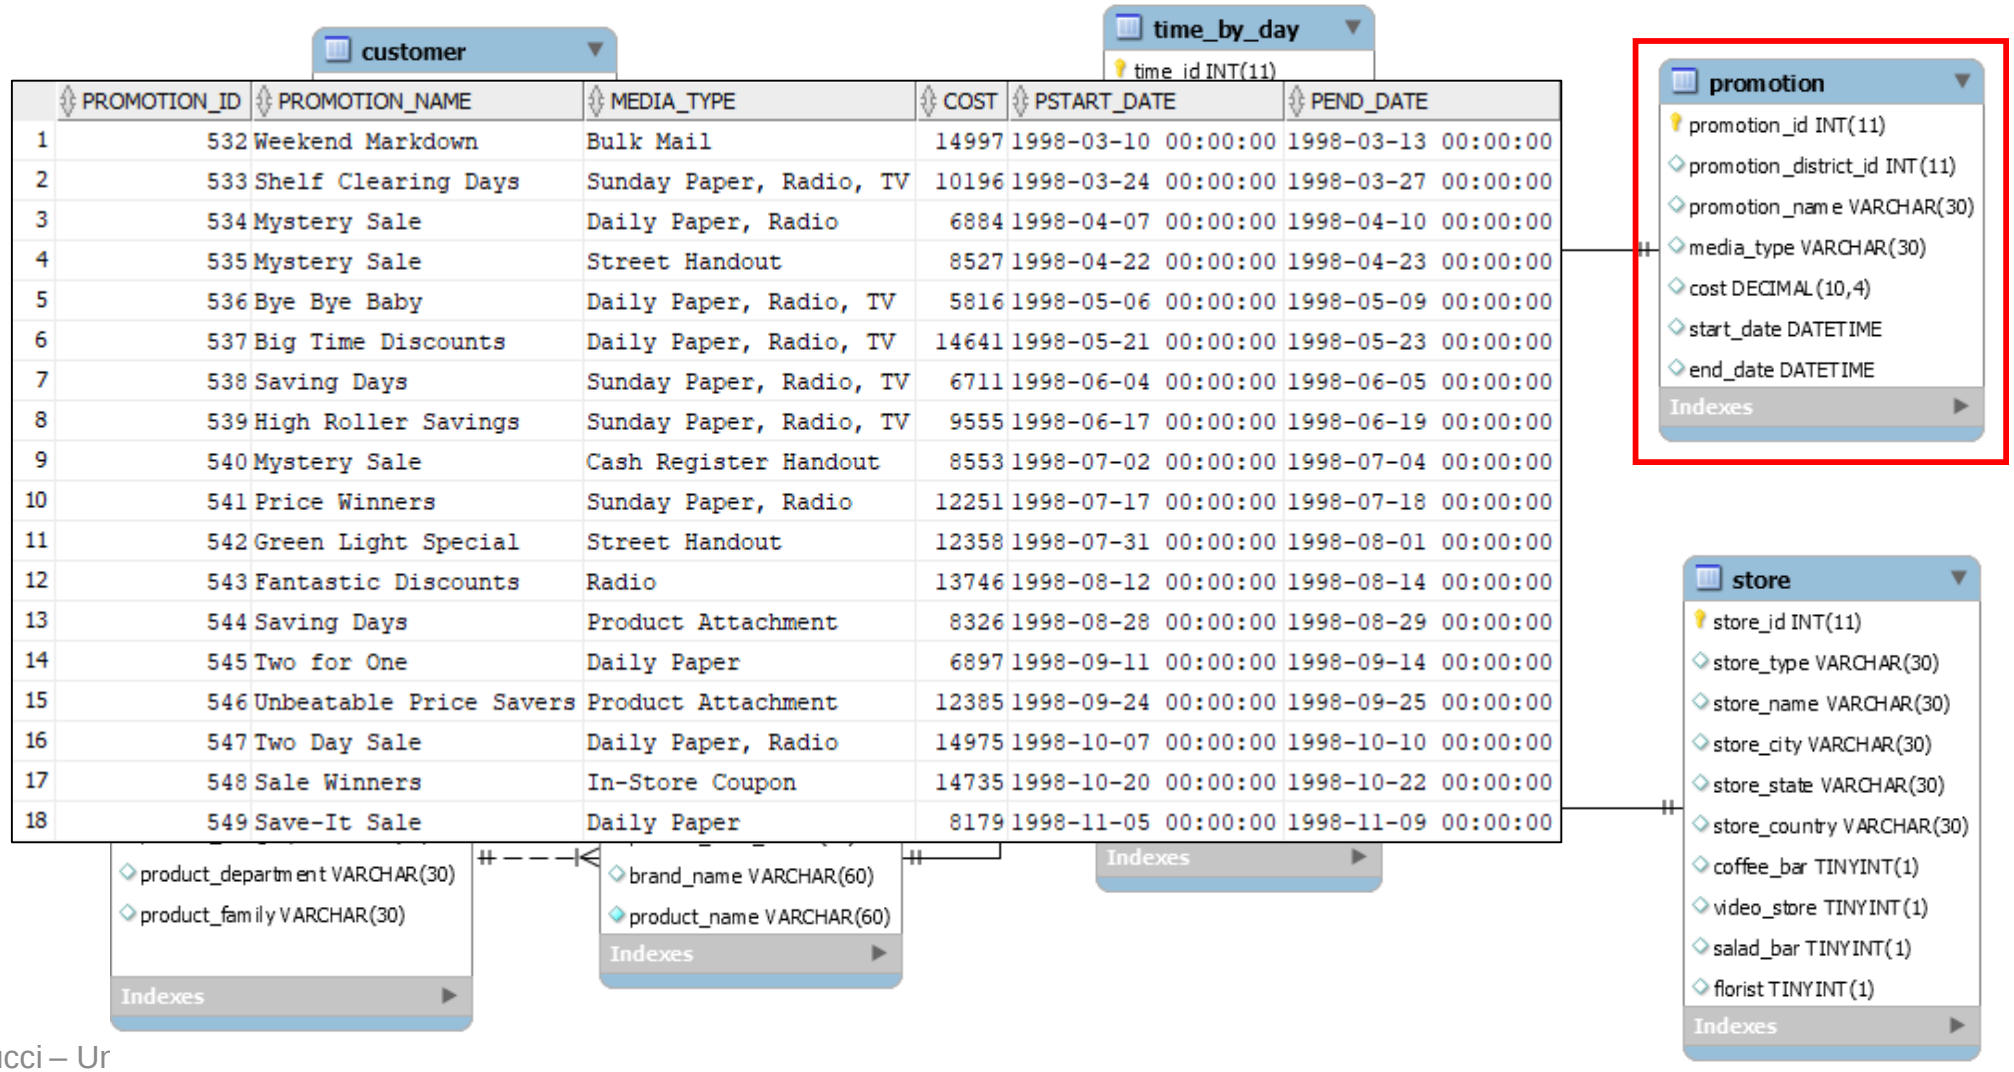

	PROMOTION_ID	PROMOTION_NAME	MEDIA_TYPE	COST	PSTART_DATE	PEND_DATE
1	532	Weekend Markdown	Bulk Mail	14997	1998-03-10 00:00:00	1998-03-13 00:00:00
2	533	Shelf Clearing Days	Sunday Paper, Radio, TV	10196	1998-03-24 00:00:00	1998-03-27 00:00:00
3	534	Mystery Sale	Daily Paper, Radio	6884	1998-04-07 00:00:00	1998-04-10 00:00:00
4	535	Mystery Sale	Street Handout	8527	1998-04-22 00:00:00	1998-04-23 00:00:00
5	536	Bye Bye Baby	Daily Paper, Radio, TV	5816	1998-05-06 00:00:00	1998-05-09 00:00:00
6	537	Big Time Discounts	Daily Paper, Radio, TV	14641	1998-05-21 00:00:00	1998-05-23 00:00:00
7	538	Saving Days	Sunday Paper, Radio, TV	6711	1998-06-04 00:00:00	1998-06-05 00:00:00
8	539	High Roller Savings	Sunday Paper, Radio, TV	9555	1998-06-17 00:00:00	1998-06-19 00:00:00
9	540	Mystery Sale	Cash Register Handout	8553	1998-07-02 00:00:00	1998-07-04 00:00:00
10	541	Price Winners	Sunday Paper, Radio	12251	1998-07-17 00:00:00	1998-07-18 00:00:00
11	542	Green Light Special	Street Handout	12358	1998-07-31 00:00:00	1998-08-01 00:00:00
12	543	Fantastic Discounts	Radio	13746	1998-08-12 00:00:00	1998-08-14 00:00:00
13	544	Saving Days	Product Attachment	8326	1998-08-28 00:00:00	1998-08-29 00:00:00
14	545	Two for One	Daily Paper	6897	1998-09-11 00:00:00	1998-09-14 00:00:00
15	546	Unbeatable Price Savers	Product Attachment	12385	1998-09-24 00:00:00	1998-09-25 00:00:00
16	547	Two Day Sale	Daily Paper, Radio	14975	1998-10-07 00:00:00	1998-10-10 00:00:00
17	548	Sale Winners	In-Store Coupon	14735	1998-10-20 00:00:00	1998-10-22 00:00:00
18	549	Save-It Sale	Daily Paper	8179	1998-11-05 00:00:00	1998-11-09 00:00:00

Database Structure – Foodmart (Sales)

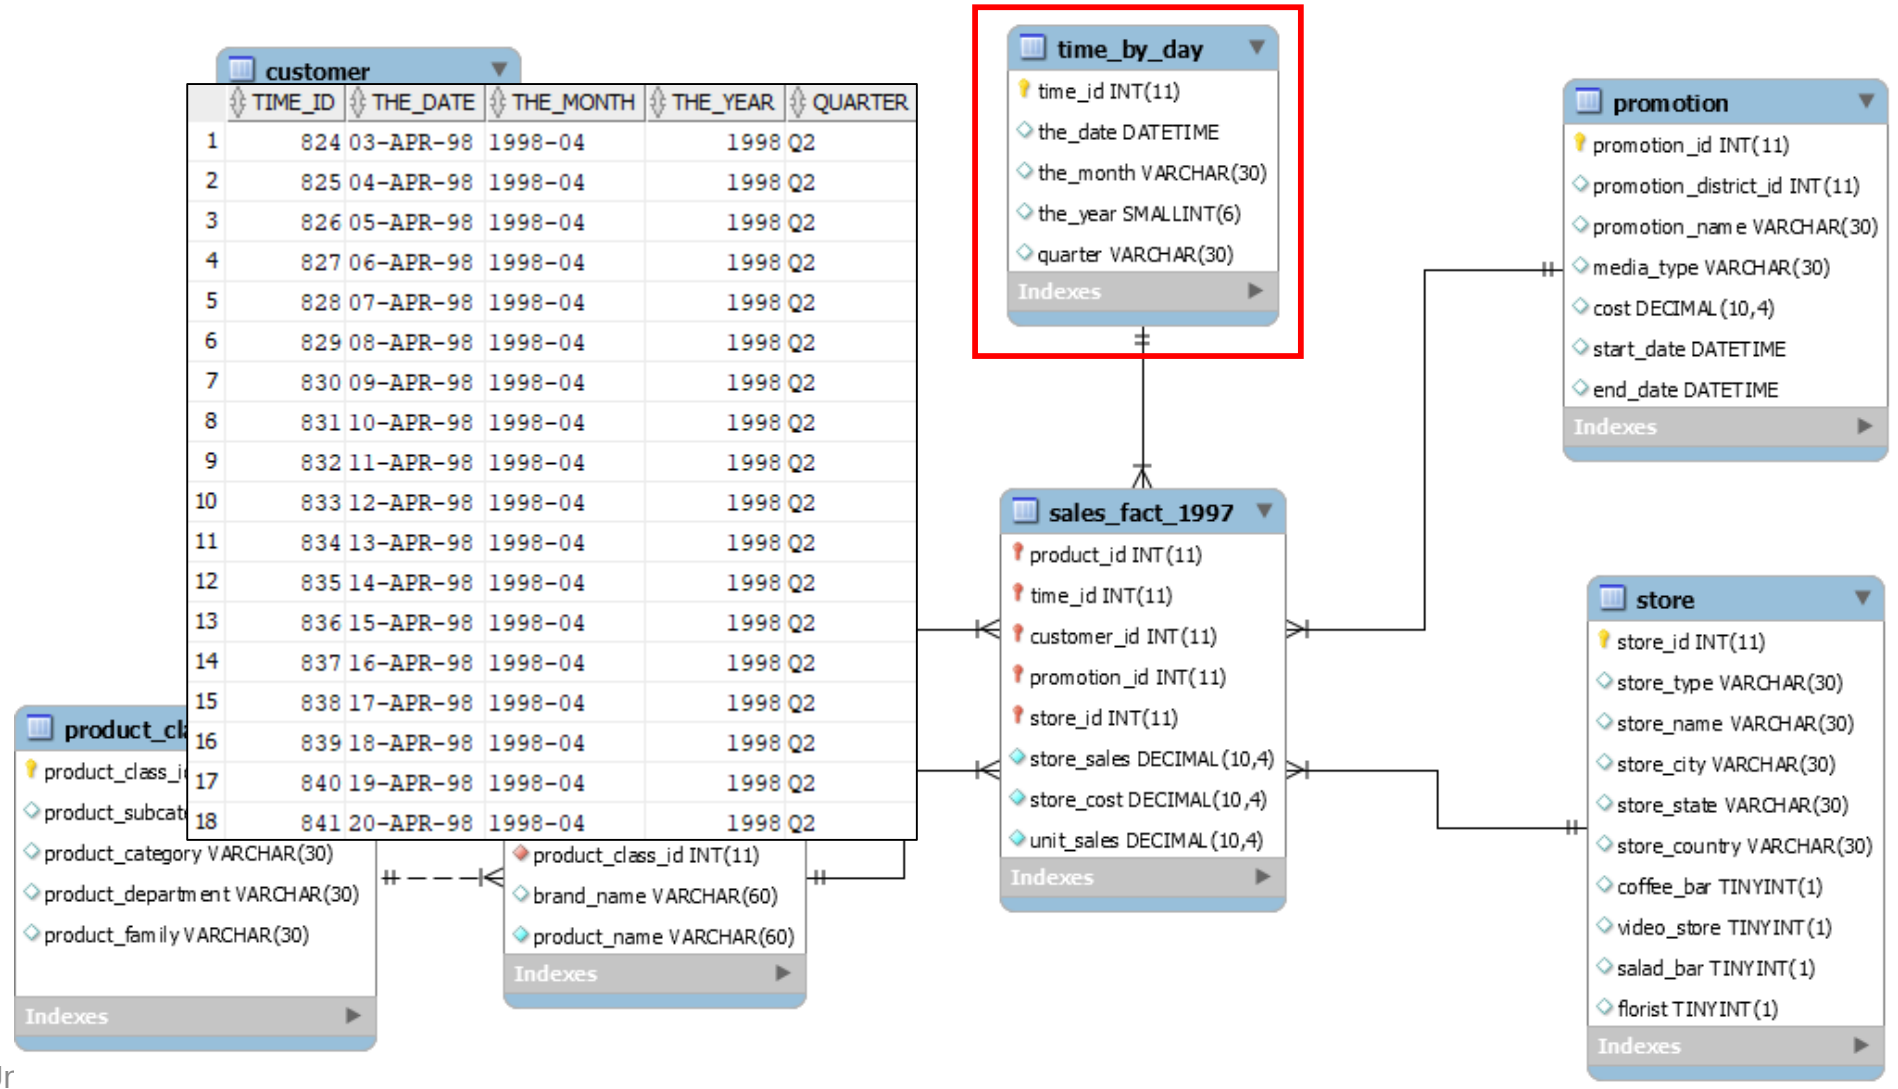

Database Structure – Foodmart (Sales)

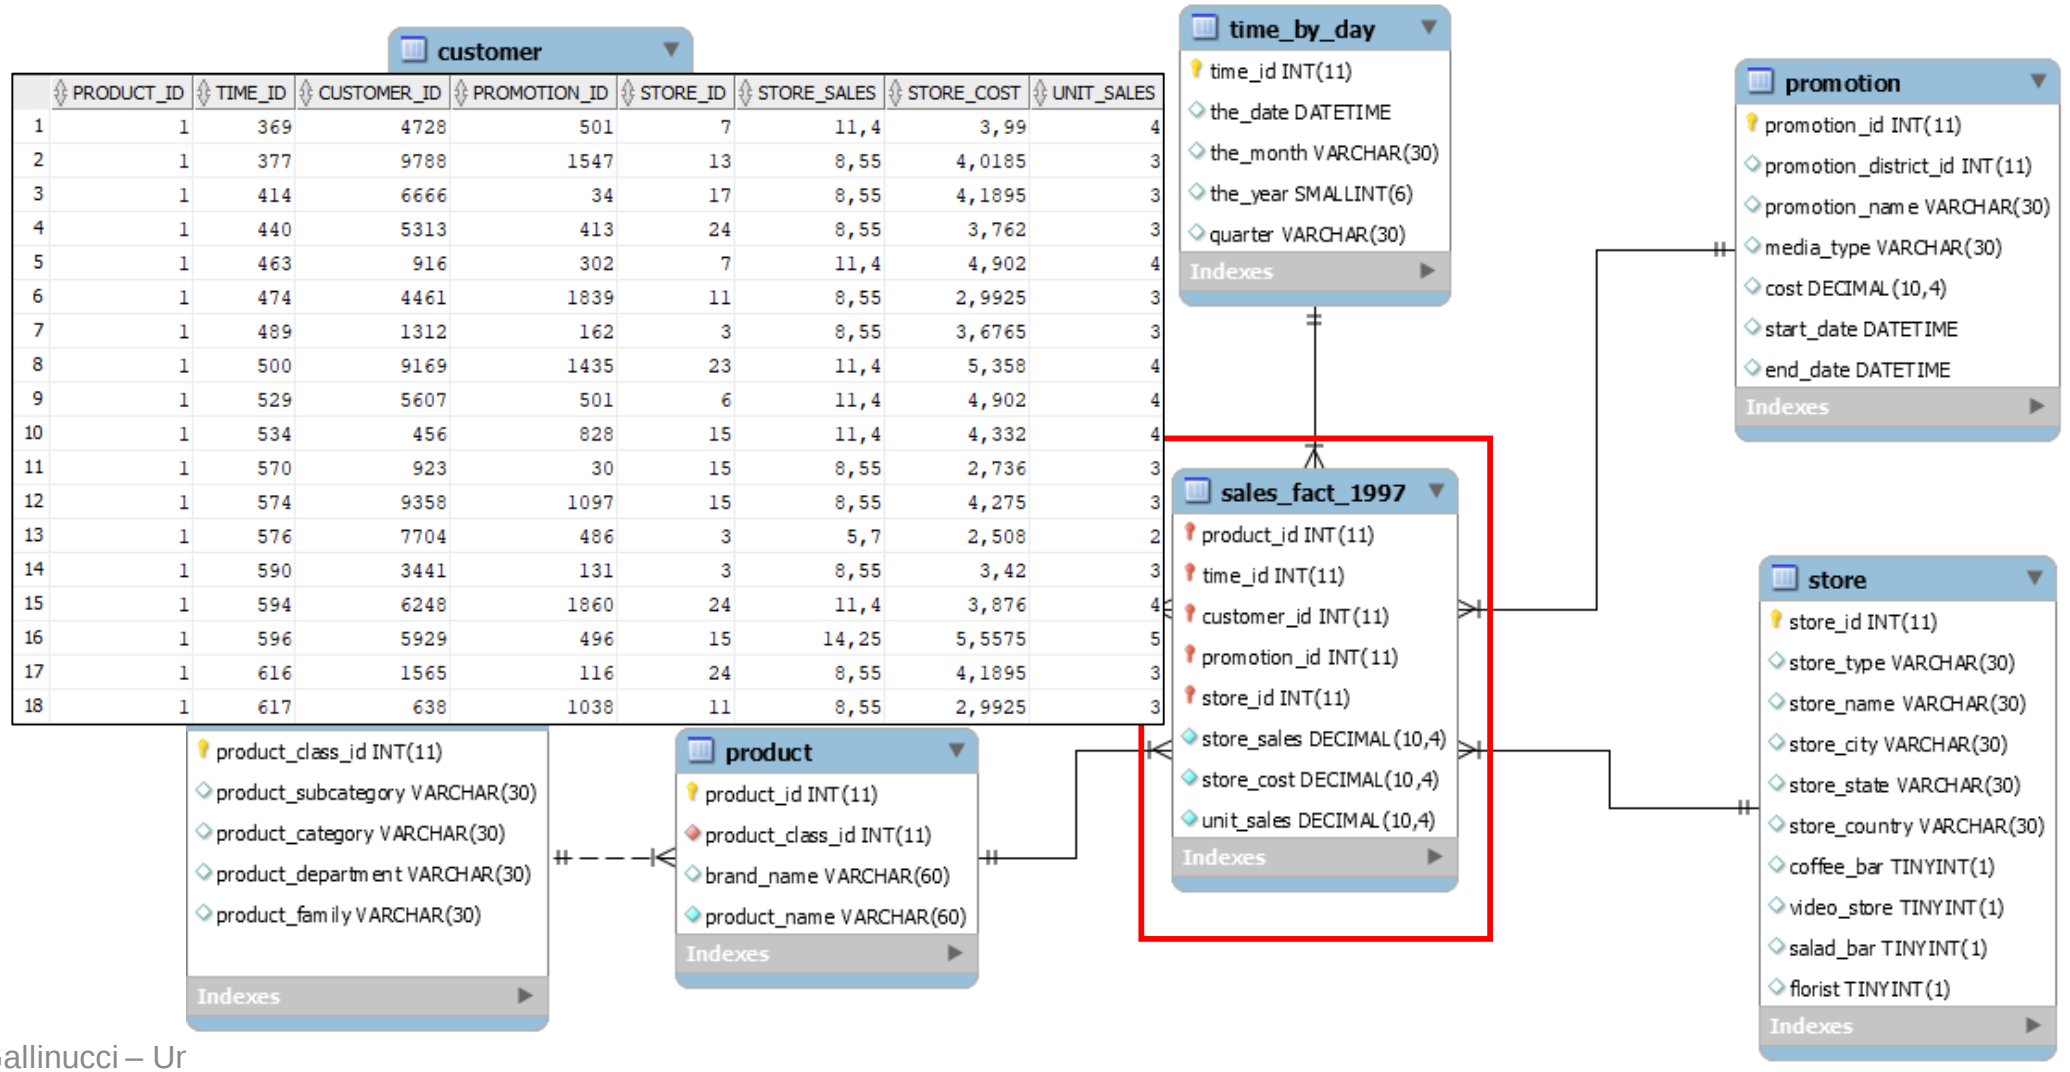

OLAP query

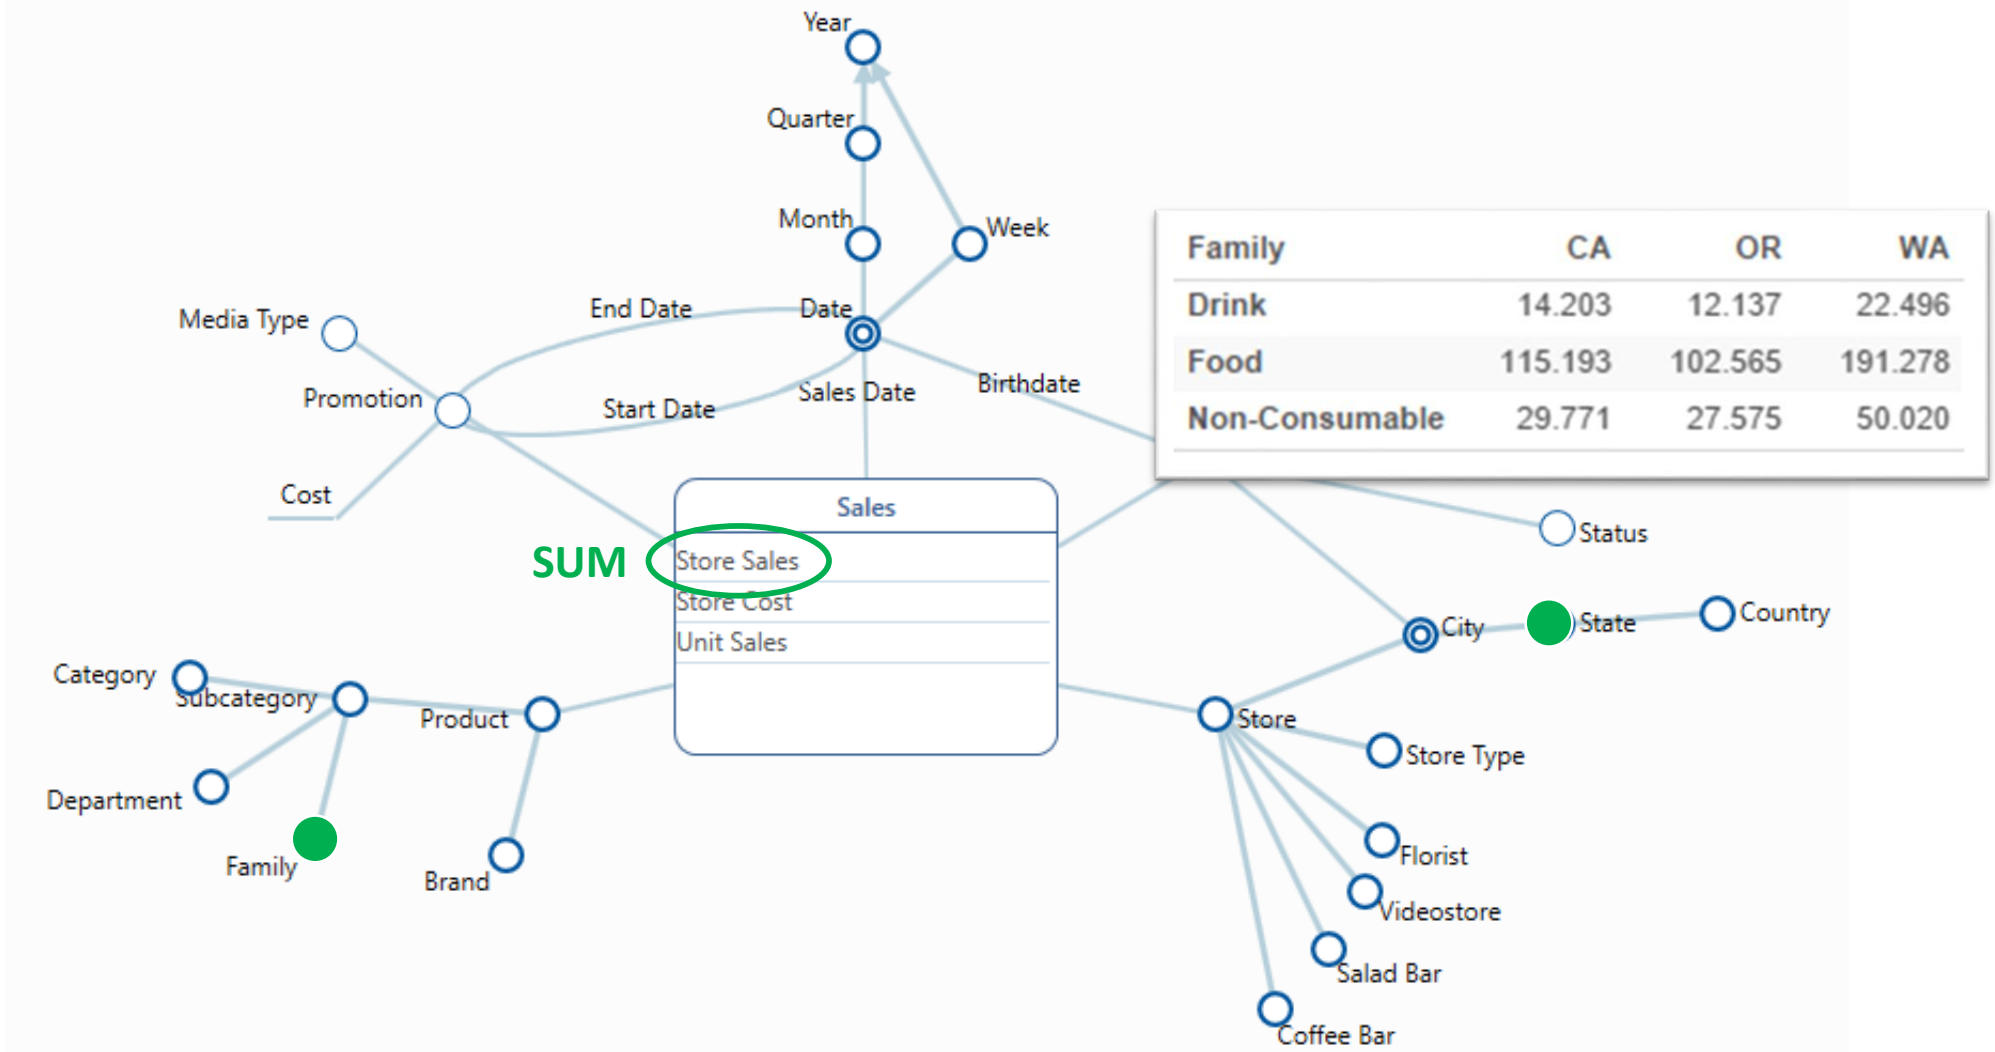

Family	CA	OR	WA
Drink	14.203	12.137	22.496
Food	115.193	102.565	191.278
Non-Consumable	29.771	27.575	50.020

OLAP Operators: Roll-Up

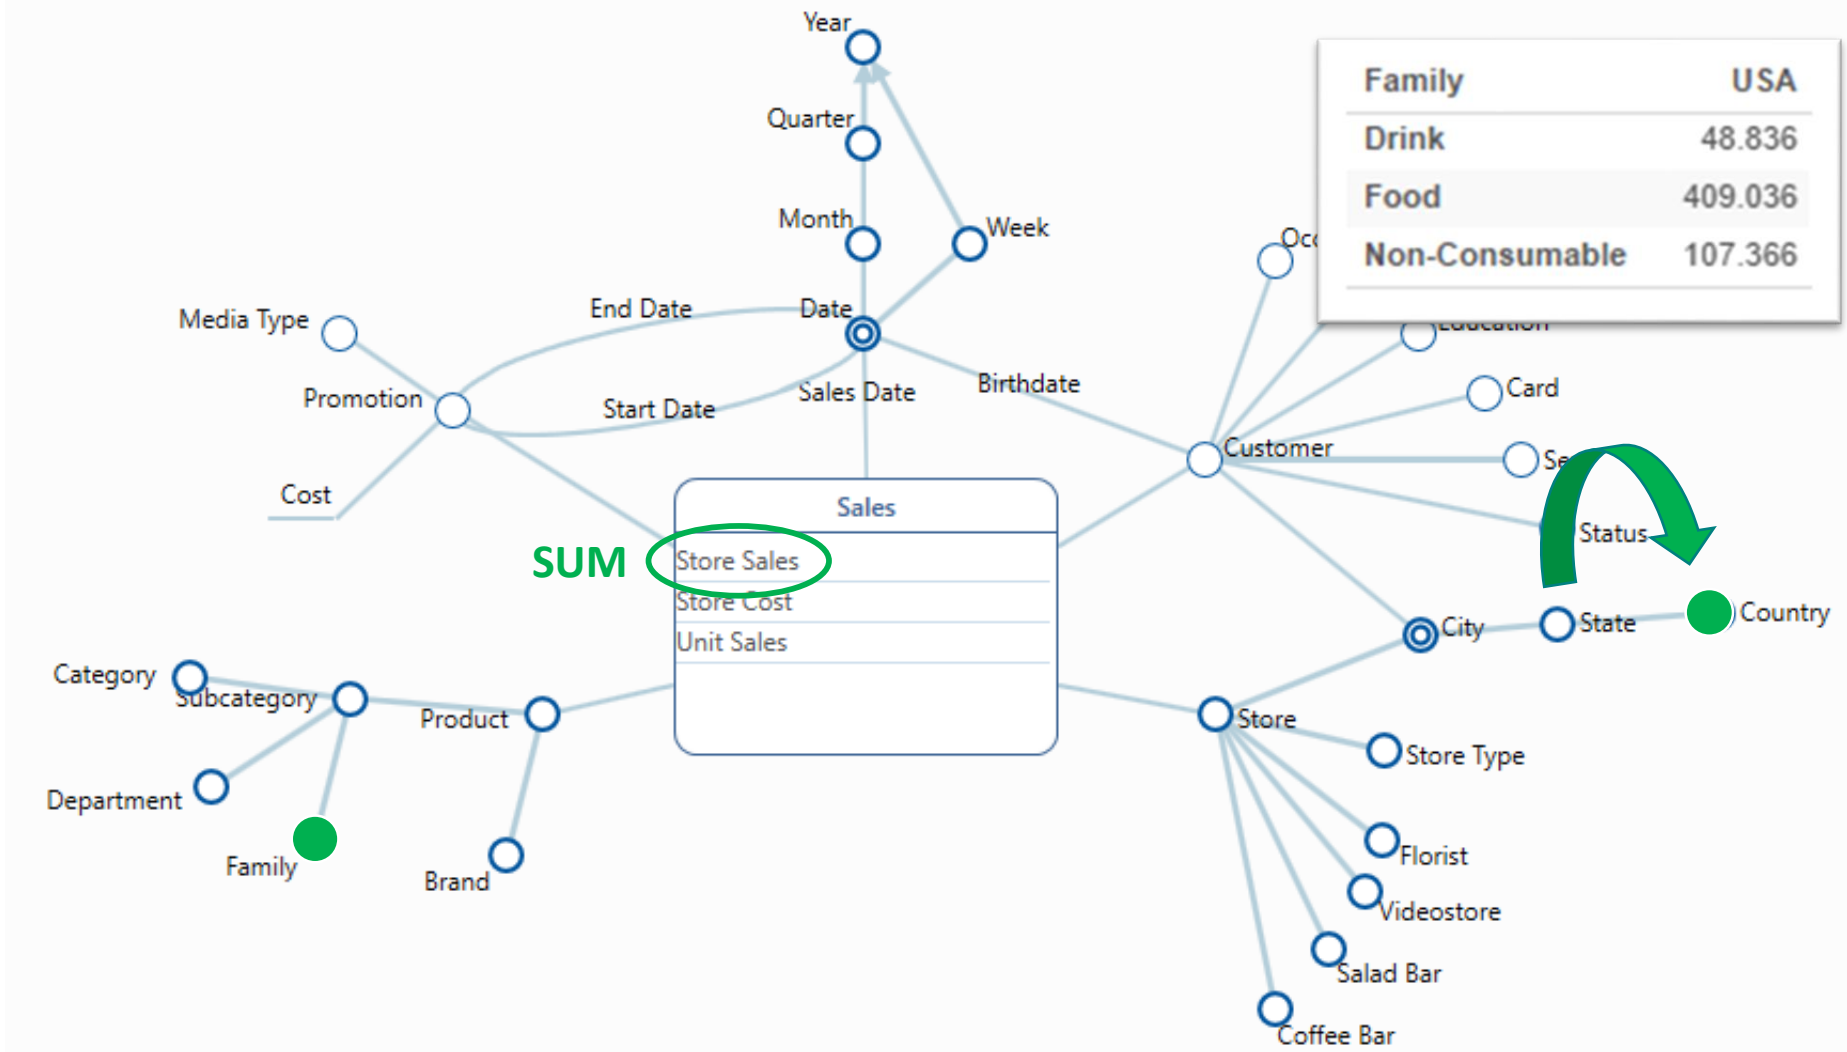

OLAP Operators: Drill-down

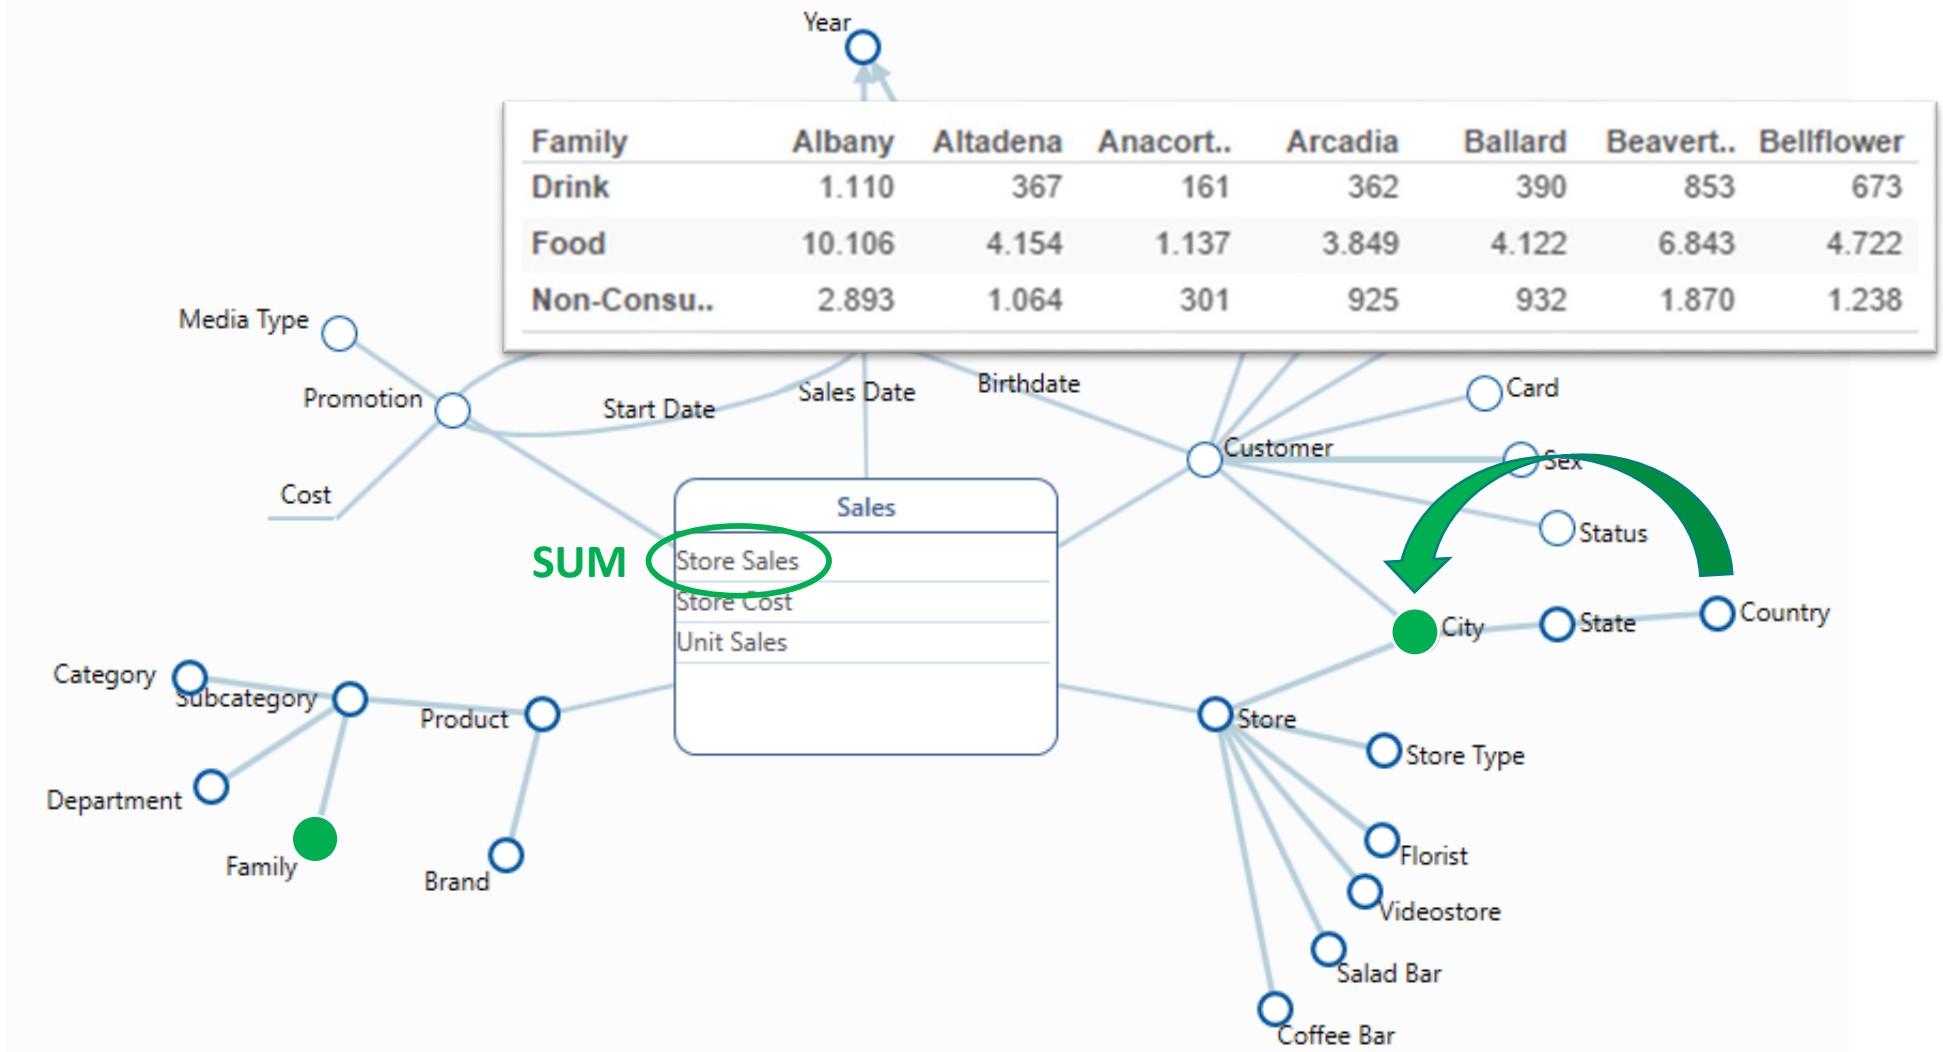

OLAP Operators: Slice & Dice

Family	Albany	Altadena	Anacort..	Arcadia	Ballard	Beavert..	Bellflower
Drink	568	158	81	153	257	351	367
Food	4.270	1.883	555	1.960	1.999	3.026	2.477
Non-Consu..	1.387	502	146	437	443	911	571

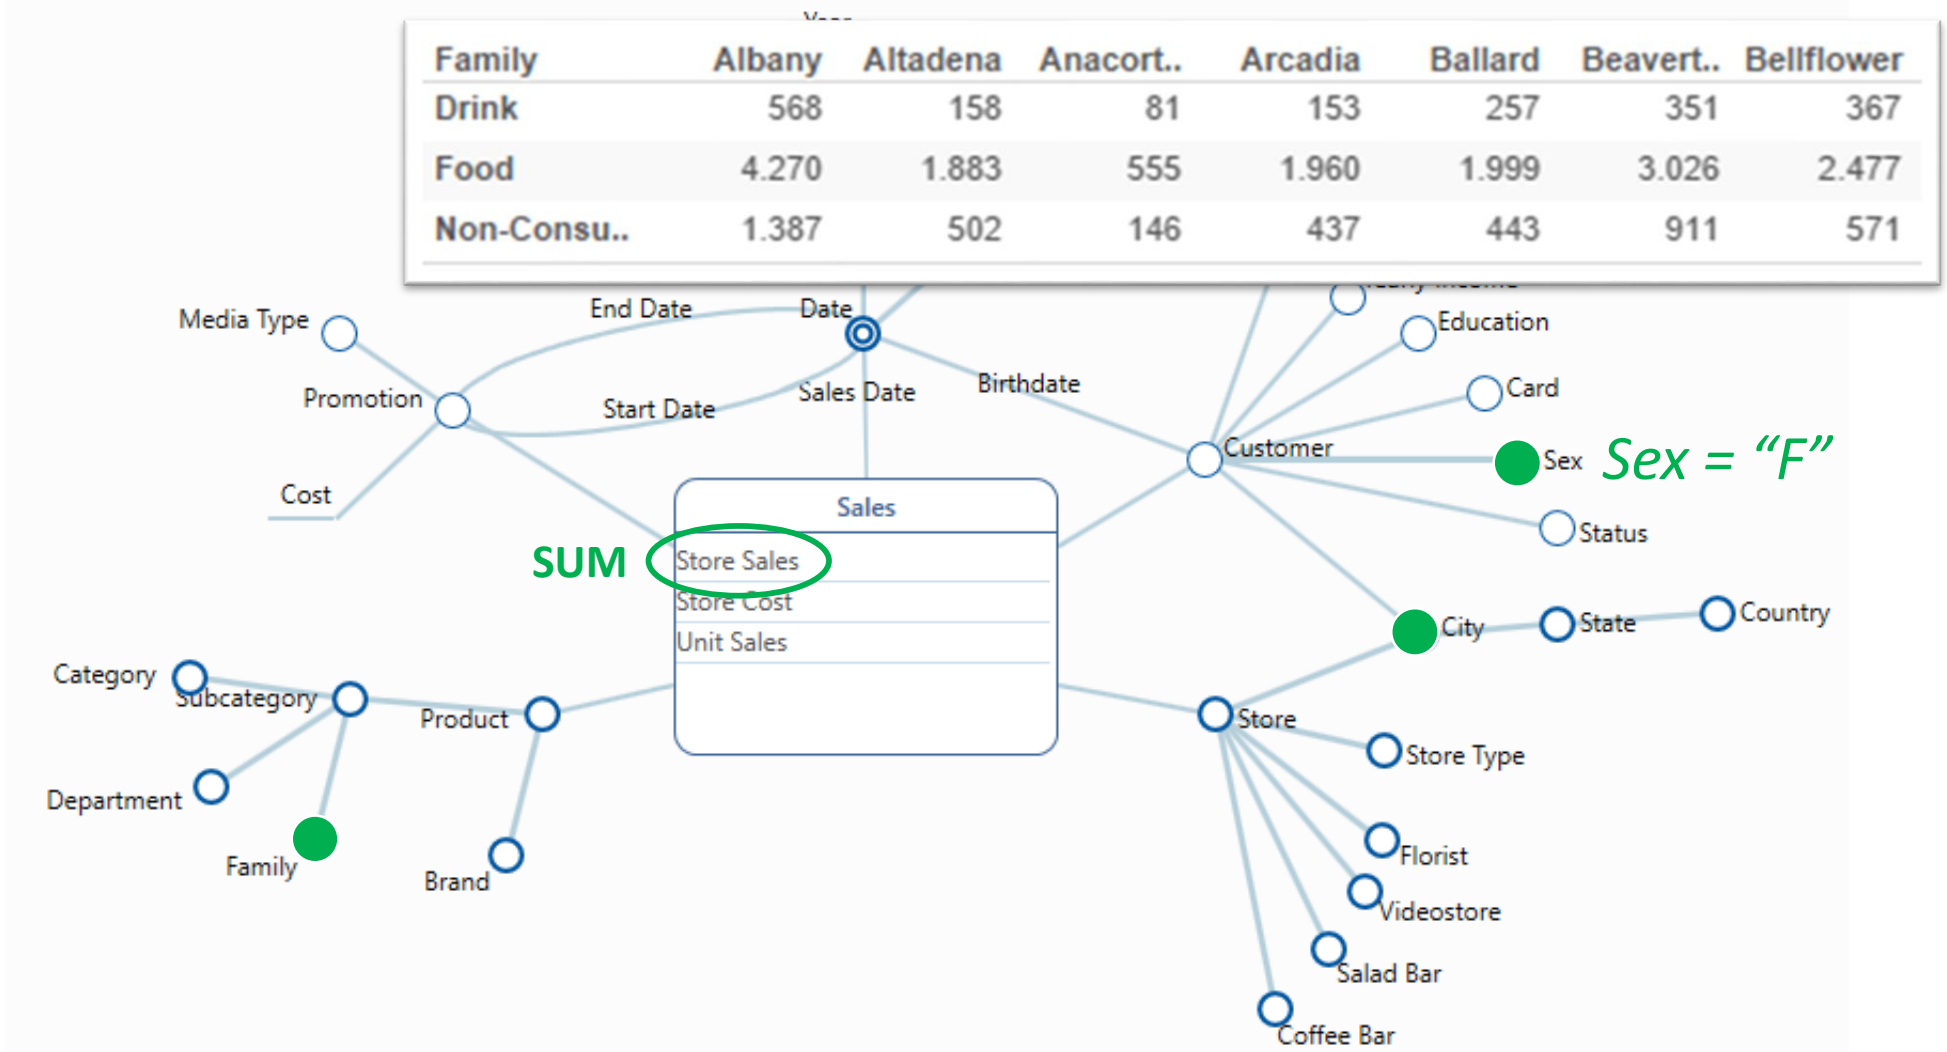

OLAP Operators: Pivoting

Family	Albany	Altadena	Anacort..	Arcadia	Ballard	Beavert..	Bellflower
Drink	568	158	81	153	257	351	367
Food	4.270	1.883	555	1.960	1.999	3.026	2.477
Non-Consu..	1.387	502	146	437	443	911	571

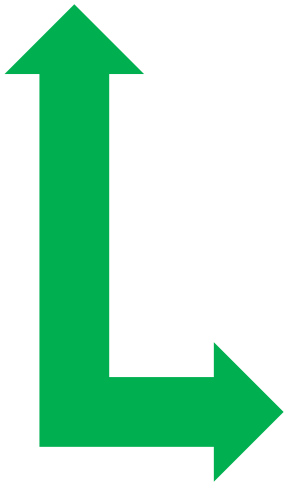

C. City	Drink	Food	Non-Con..
Albany	568	4.270	1.387
Altadena	158	1.883	502
Anacortes	81	555	146
Arcadia	153	1.960	437
Ballard	257	1.999	443
Beaverton	351	3.026	911
Bellflower	367	2.477	571

Tableau: Visual Analytics

tableau

Data

Query

MySQL

TXT

ODBC

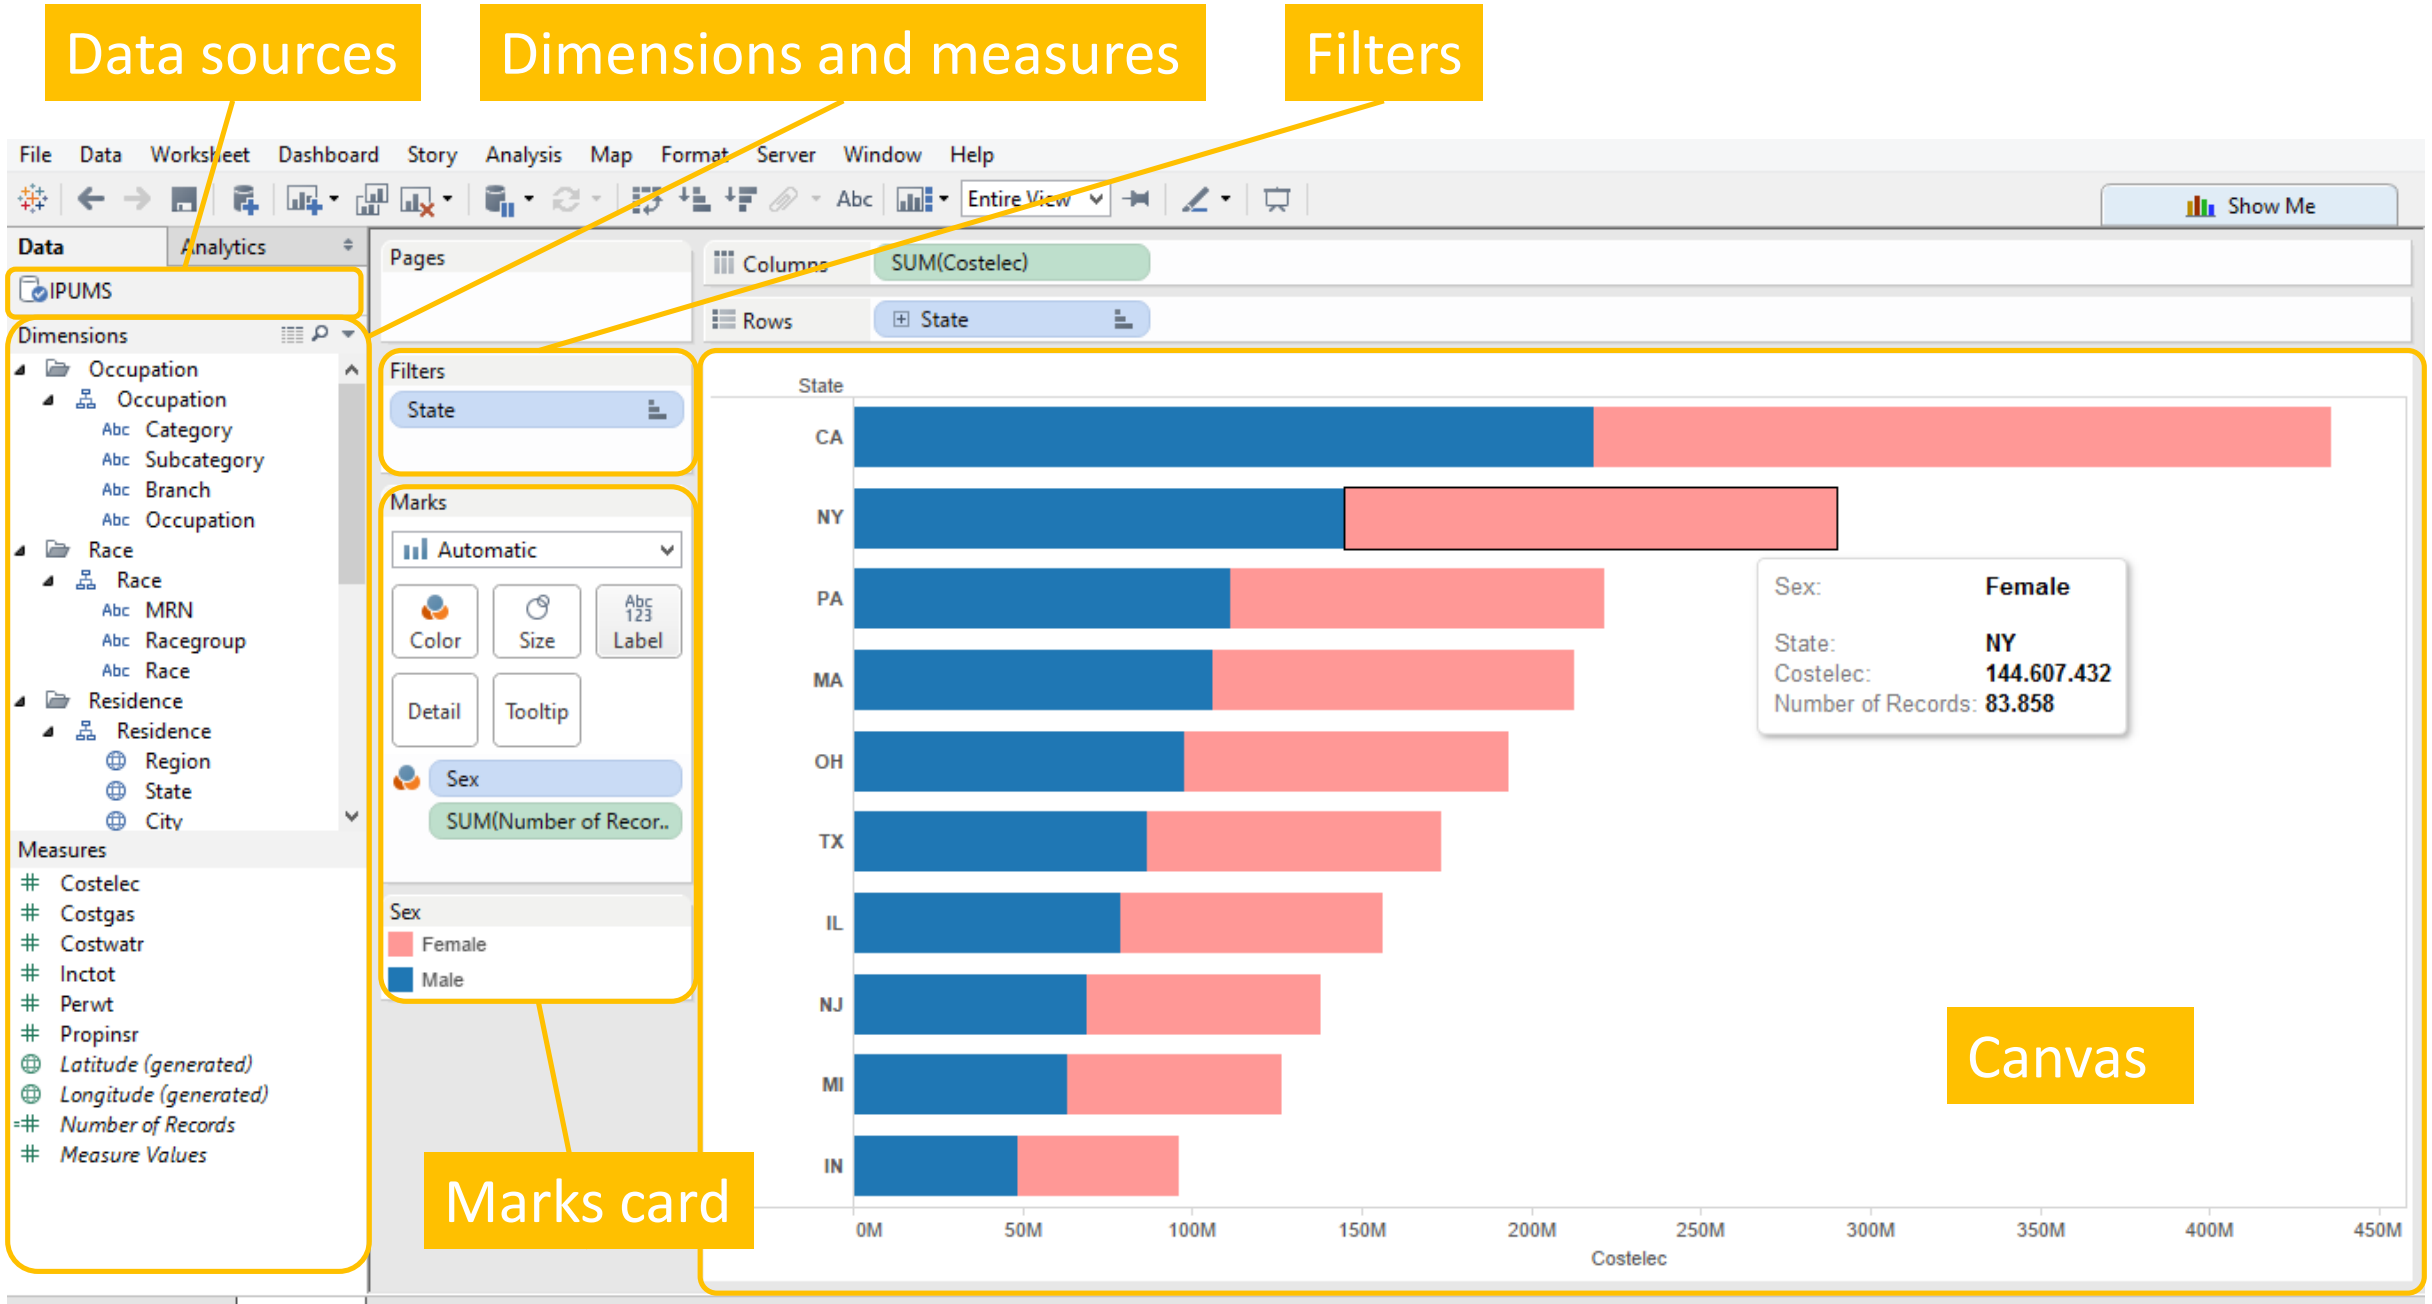
Data sources
Dimensions and measures
Filters
File Data Worksheet Dashboard Story Analysis Map Format Server Window Help
Show Me
Data
Analytics
IPUMS
Dimensions
Occupation
Occupation
Category
Subcategory
Branch
Occupation
Race
Race
MRN
Racegroup
Race
Residence
Residence
Region
State
City
Measures
Costelec
Costgas
Costwatr
Inctot
Perwt
Propinsr
Latitude (generated)
Longitude (generated)
Number of Records
Measure Values
Pages
Columns
SUM(Costelec)
Rows
State
Filters
State
Marks
Automatic
Color
Size
Label
Detail
Tooltip
Sex
SUM(Number of Recor..
Sex
Female
Male
State
CA
NY
PA
MA
OH
TX
IL
NJ
MI
IN
0M
50M
100M
150M
200M
250M
300M
350M
400M
450M
Costelec
Sex: Female
State: NY
Costelec: 144.607.432
Number of Records: 83.858
Canvas
Marks card
Data Source
Sheet 1
20 marks
10 rows by 1 column
SUM(Costelec): 2.043.964.564

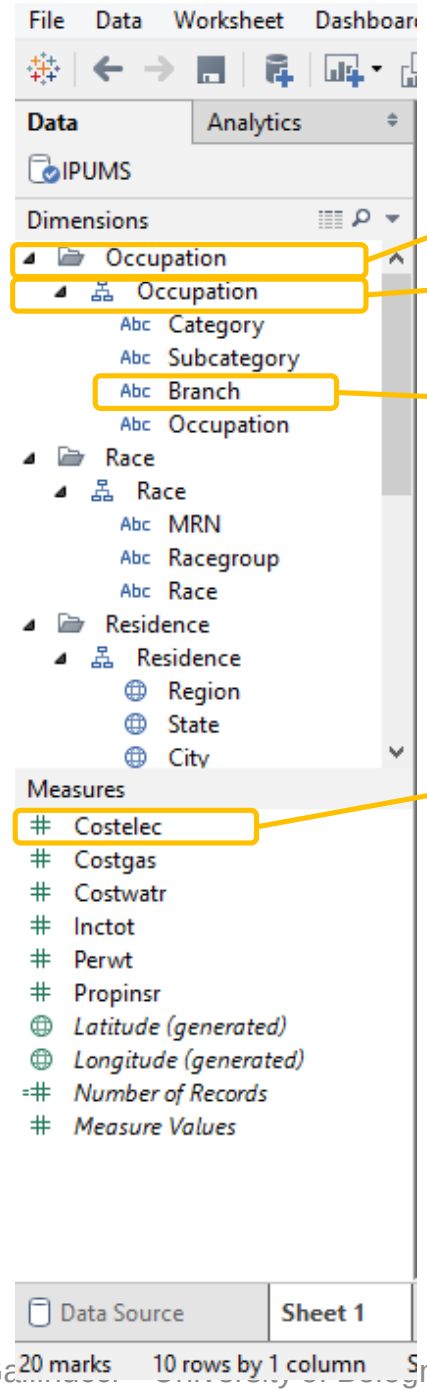

Dimension

Hierarchy

Dimensional attribute

Measure

- Conventionally, folders are used to represent dimensions; in general they are simply a way to group elements

Dimension VS Measure

With Tableau the definitions of dimension and measure are quite loose. Indeed, every field can be used both as a dimension and as a measure

Regardless, it is a good idea to give an initial classification to the available fields by following these guidelines

- A dimension is any *independent* field, such as *city*, *product*, etc.
- A measure is any field that is *dependent* on other fields, such as *profit*
- Usually measures are numerical while dimensions are categorical (but not always!)

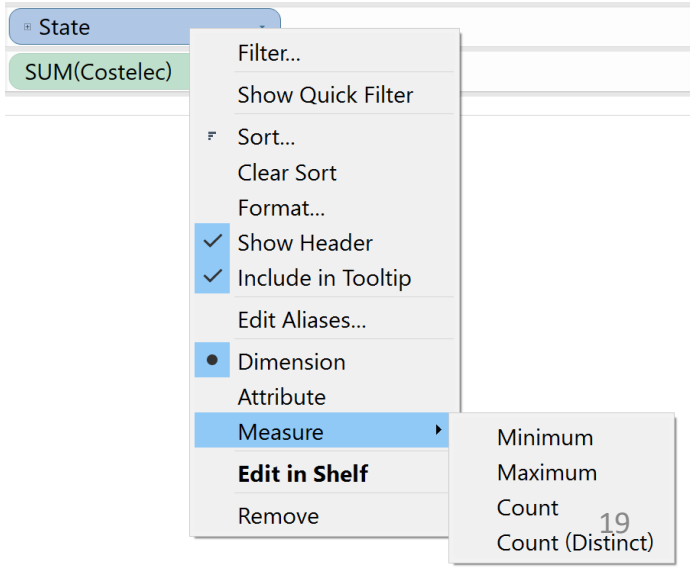

Hierarchies

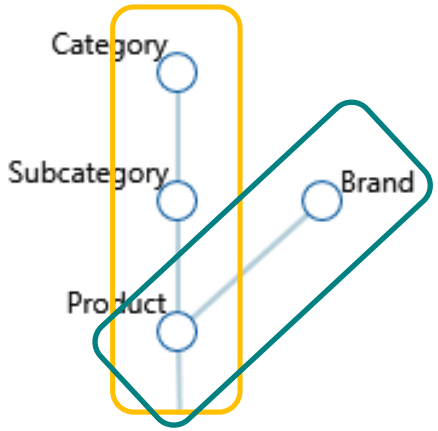

- Each path from leaves to root becomes a different hierarchy; shared attributes are duplicated

- Shared hierarchies are duplicated

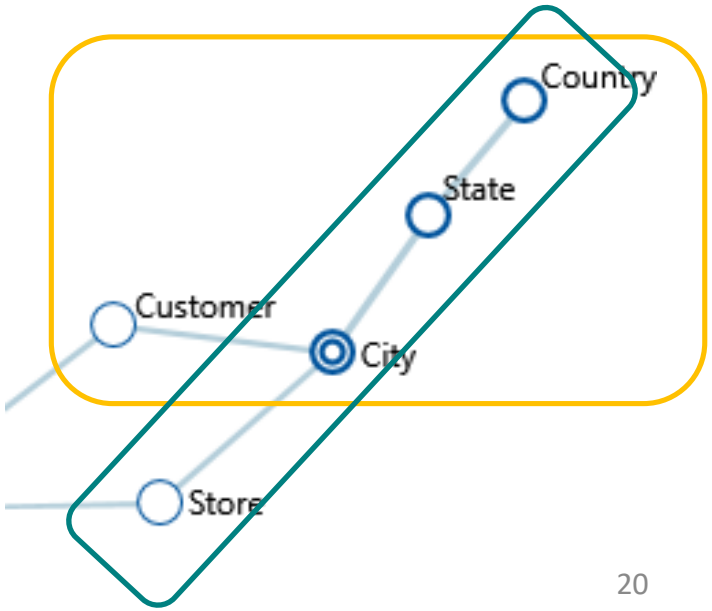

Green VS Blue

With Tableau, the greed colour is associated to *continuous* fields, while the blue colour refers to *discrete* ones

Often (but not always) measures are continuous fields, while dimensions are discrete

Continuous and discrete fields behave in different ways

- When they are used on rows and columns
- When a filter is applied on them
- When they are associated to colours (i.e., with a colour marker)

Green VS Blue (Rows and Columns)

When they are used on rows and columns

- A discrete field generates a *header* where each value has a its own label
- A *continuous* field generates an *axis*

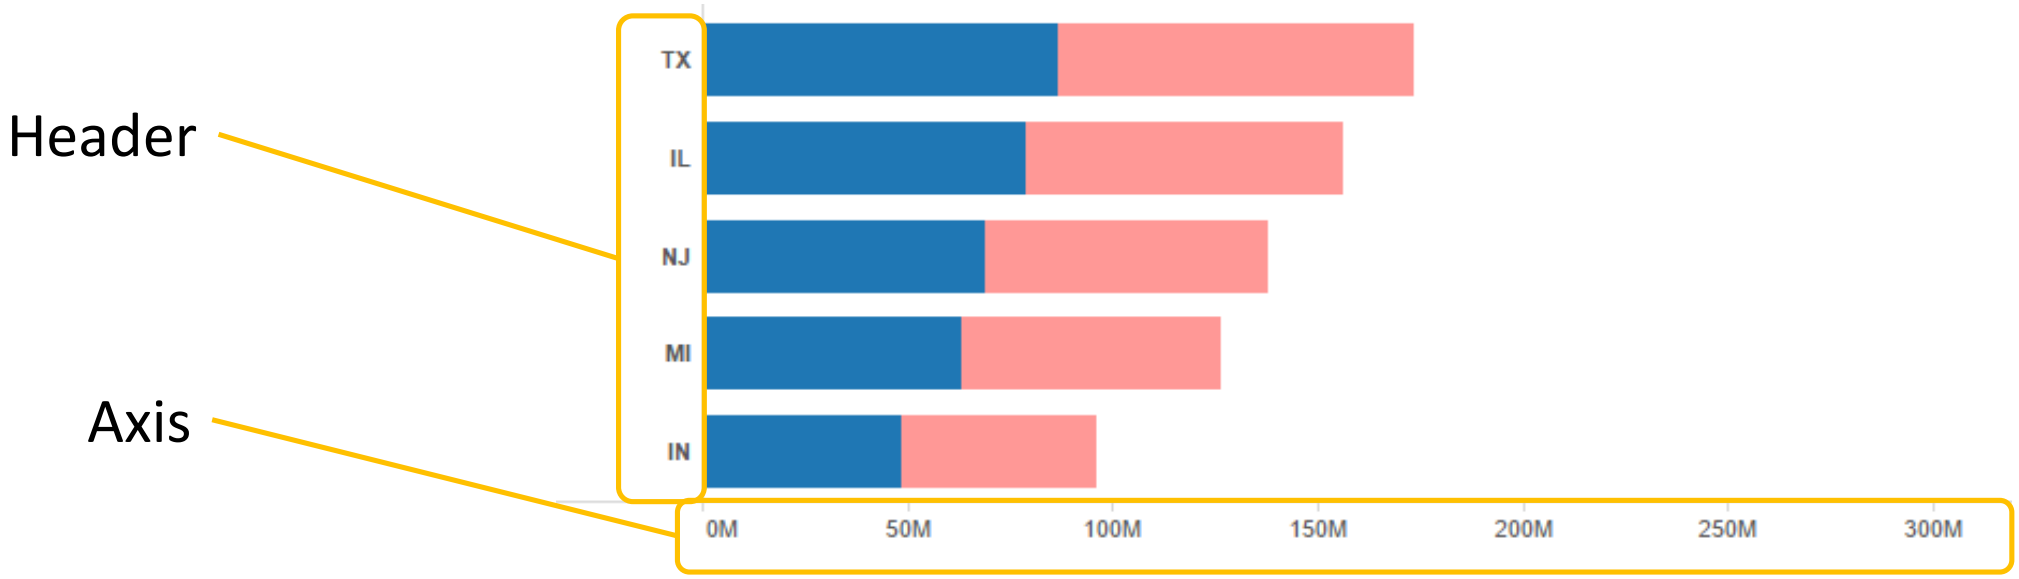

Green VS Blue (Filters)

- With a continuous field it is possible to specify a *range*
  - On the values at the most detailed level (i.e., not aggregated)
  - Or at a particular aggregation level of the base values

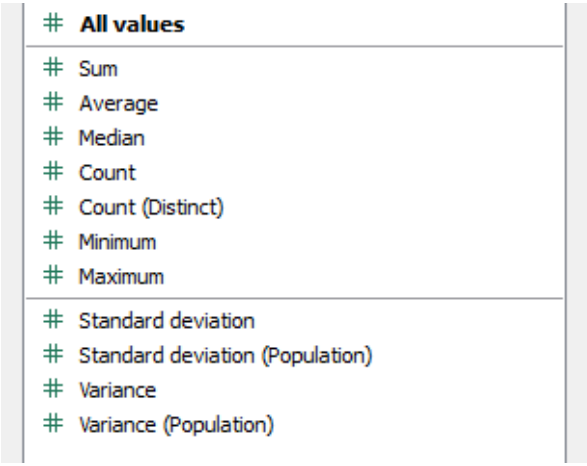

- With a discrete field it is also possible to select specific values one by one

Green VS Blue (Colours)

- A continuous field is associated to a sequence of related colours; e.g., different shades of the same base colour: darker and brighter shades may refer to respectively higher and lower values

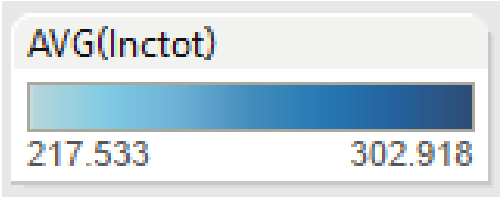

- With a discrete field each value is associated to a single arbitrary colour without any constraint

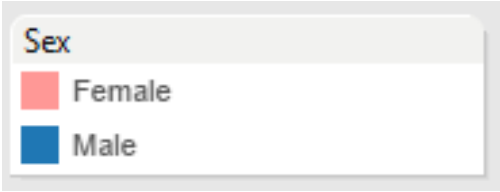

Canvas

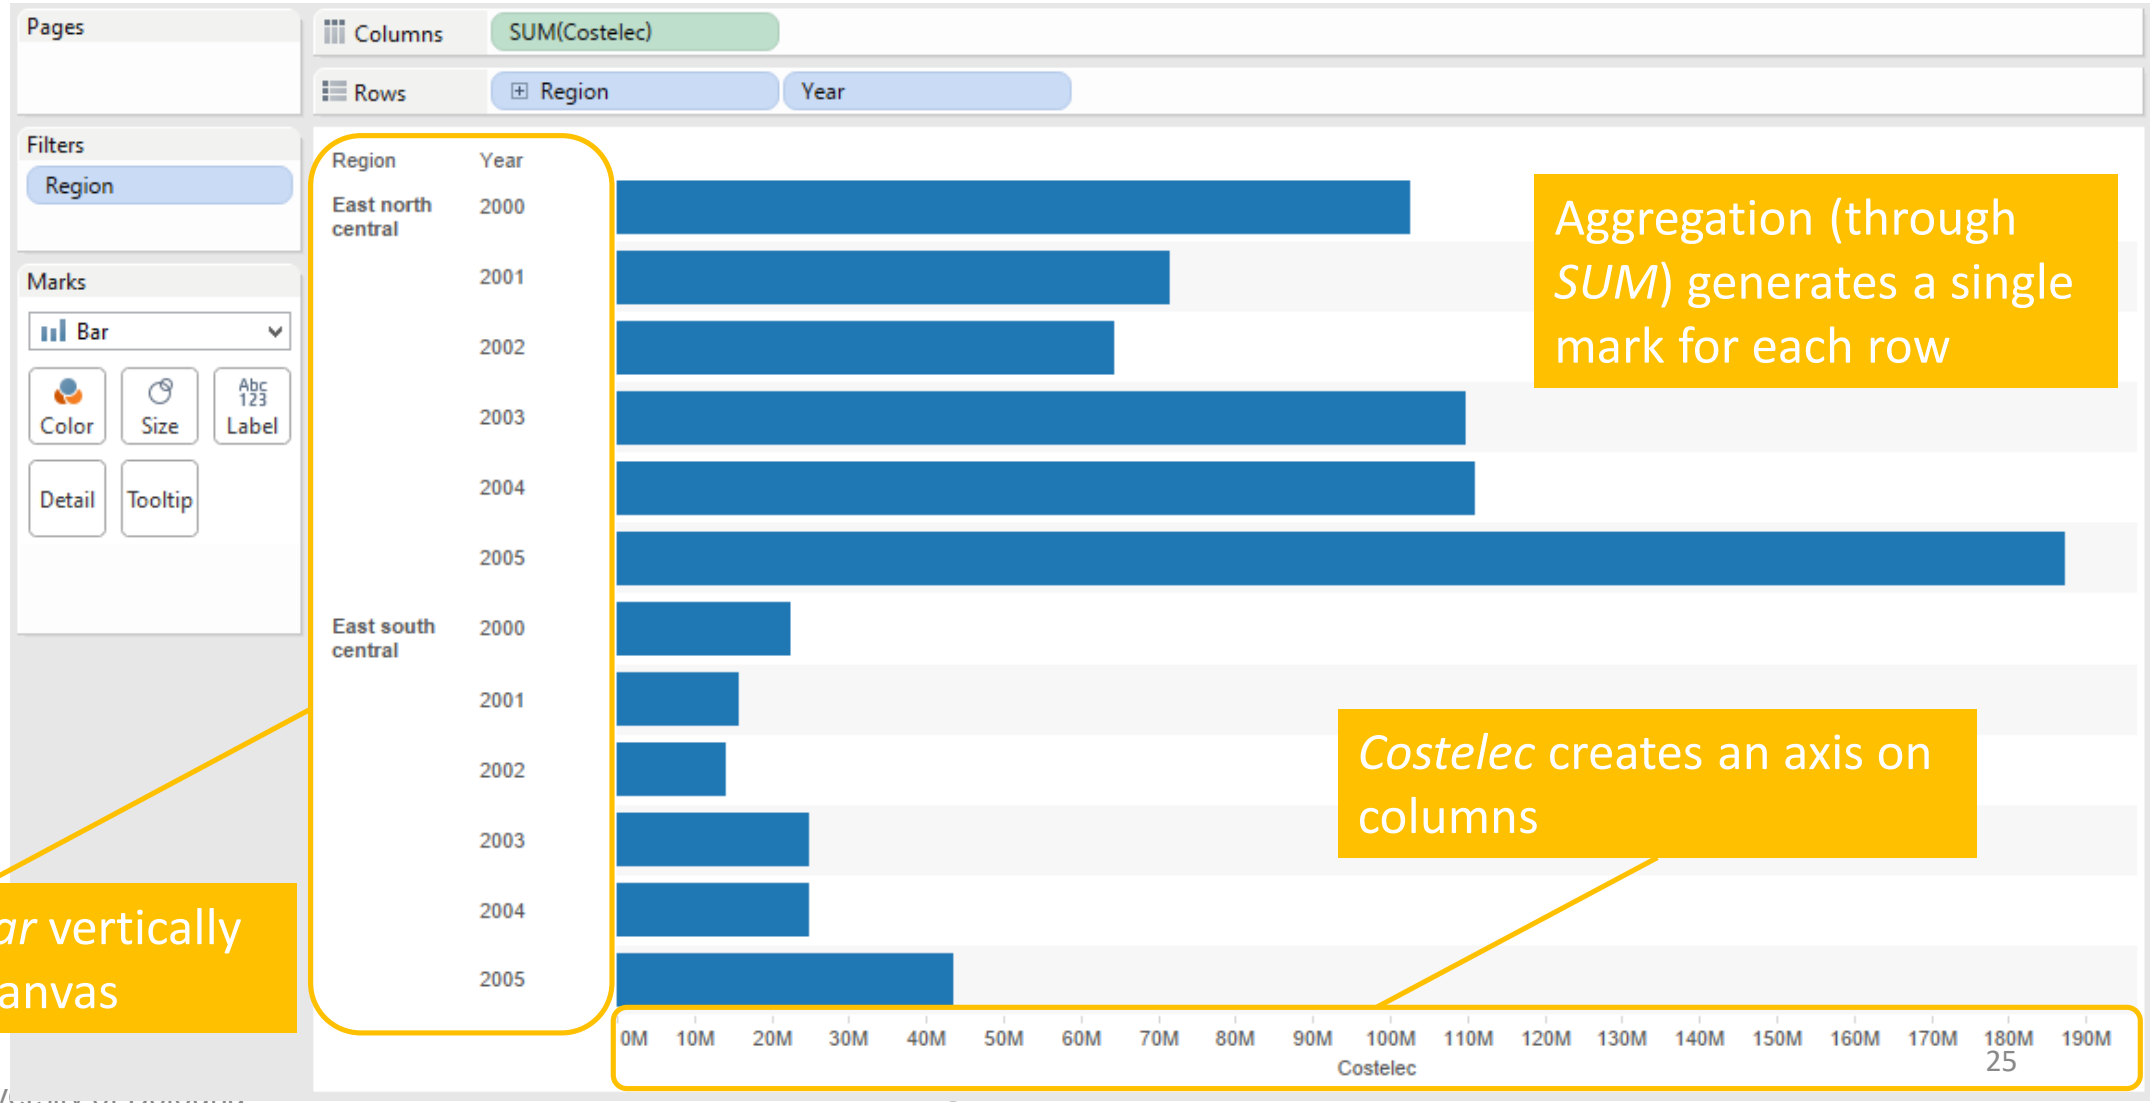
Pages
Columns
SUM(Costelec)
Rows
Region
Year
Filters
Region
Marks
Bar
Color
Size
Label
Detail
Tooltip
Region
Year
East north central
2000
2001
2002
2003
2004
2005
East south central
2000
2001
2002
2003
2004
2005
0M
10M
20M
30M
40M
50M
60M
70M
80M
90M
100M
110M
120M
130M
140M
150M
160M
170M
180M
190M
Costelec
Aggregation (through *SUM*) generates a single mark for each row
Costelec creates an axis on columns
Region and *Year* vertically partition the canvas

Mark

Data are visualised inside the canvas through the use of *marks*

There are different types of marks (*bar*, *line*, *text*, etc.)

- Each mark type has differ
- At each property can be a pend on whether the field is continuous or discrete and

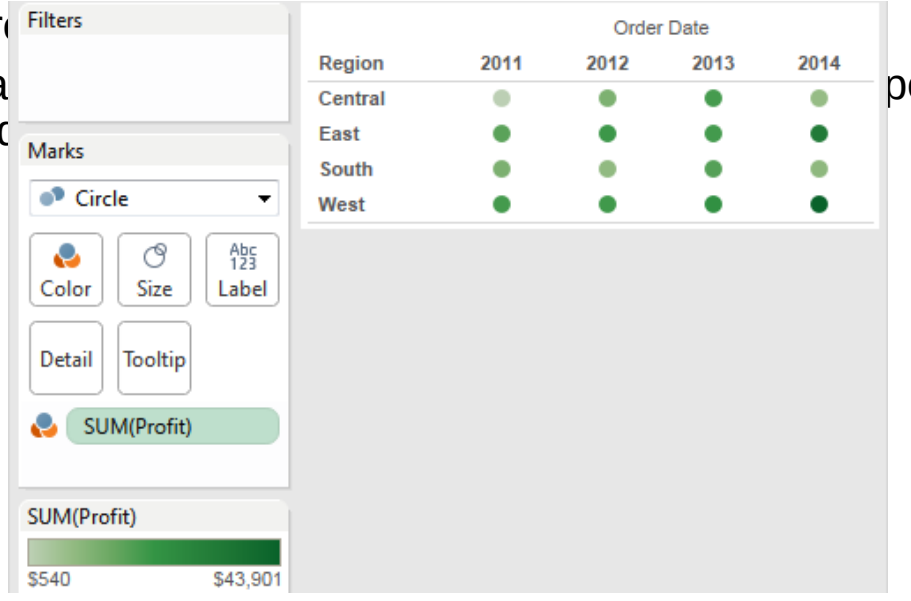

Mark (2)

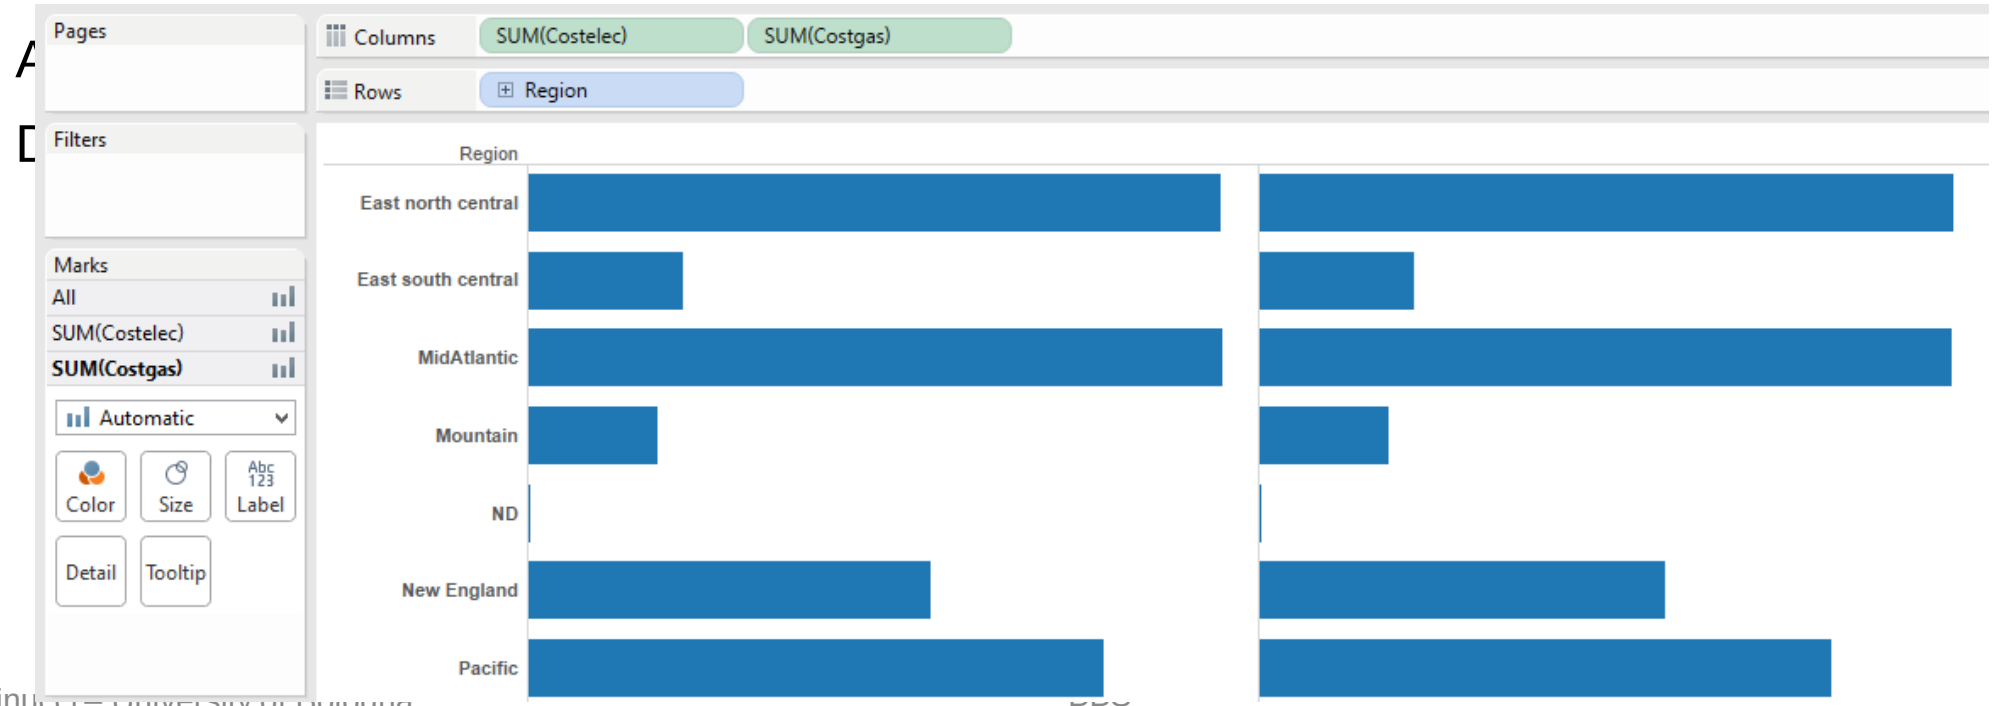
Pages
Filters
Marks
All
SUM(Costelec)
SUM(Costgas)
Automatic
Color
Size
Label
Detail
Tooltip
Columns
SUM(Costelec)
SUM(Costgas)
Rows
Region
Region
East north central
East south central
MidAtlantic
Mountain
ND
New England
Pacific

Ordering

Ordering does not break hierarchies (although there exist workarounds)

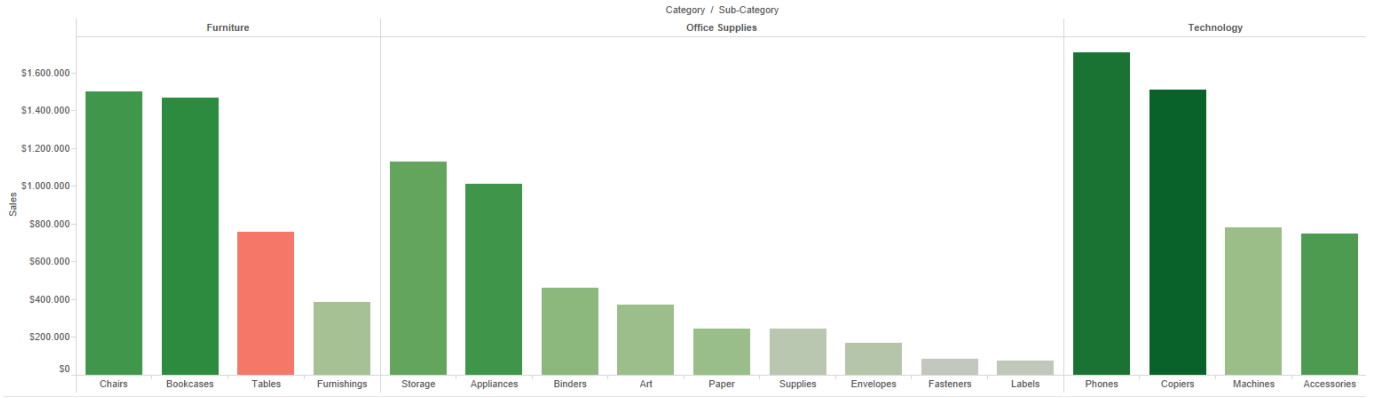

There are different types of ordering

- *Alphabetical*
- *Manual*: the order is fixed and manually chosen by the user
- *Computed*: the order is based on a computation (e.g., the sum of specific measure)
- *Data Source Order*: the order is the same as on the datasource
- *Nested: useful* in the presence of two (or more) independent fields

Ordering (2)

Education	
Partial High School	79.155
High School Degree	78.664
Bachelors Degree	68.839
Partial College	24.545
Graduate Degree	15.570

on ordered

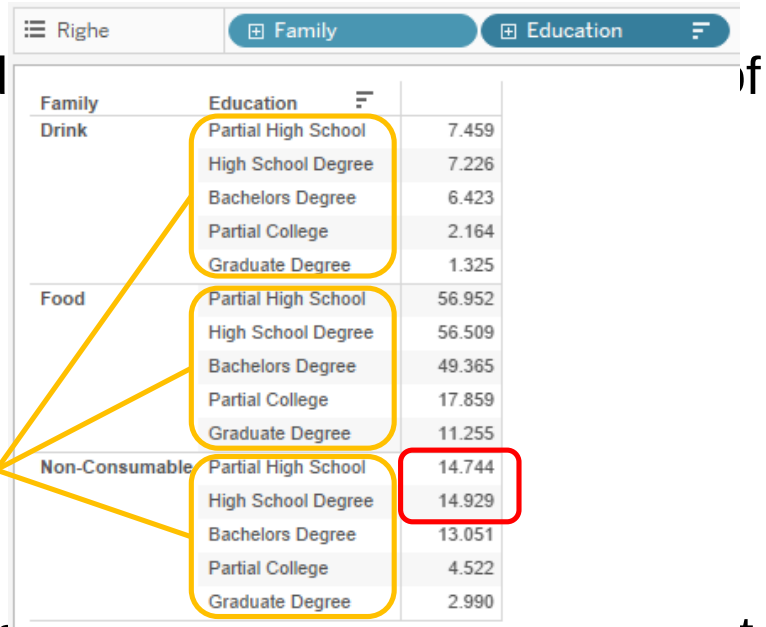

Family	Education	
Drink	Partial High School	7.459
	High School Degree	7.226
	Bachelors Degree	6.423
	Partial College	2.164
	Graduate Degree	1.325
Food	Partial High School	56.952
	High School Degree	56.509
	Bachelors Degree	49.365
	Partial College	17.859
	Graduate Degree	11.255
Non-Consumable	Partial High School	14.744
	High School Degree	14.929
	Bachelors Degree	13.051
	Partial College	4.522
	Graduate Degree	2.990

The order used for Education is the one obtained by evaluating Education independently!

Family	Education	
Drink	Partial High School	7.459
	High School Degree	7.226
	Bachelors Degree	6.423
	Partial College	2.164
	Graduate Degree	1.325
Food	Partial High School	56.952
	High School Degree	56.509
	Bachelors Degree	49.365
	Partial College	17.859
	Graduate Degree	11.255
Non-Consumable	High School Degree	14.929
	Partial High School	14.744
	Bachelors Degree	13.051
	Partial College	4.522
	Graduate Degree	2.990

Right: Education ordered by nested field (desc sum of unit sales)

View Data

Through the *View Data* option it is possible to visualise the set of records (i.e., the data at the most detailed level of aggregation) used to compute a certain mark

View Data can be used as a *light* version of the *Drill Through* operation

It is particularly useful for testing and debugging complex visualisations

View Data (2)

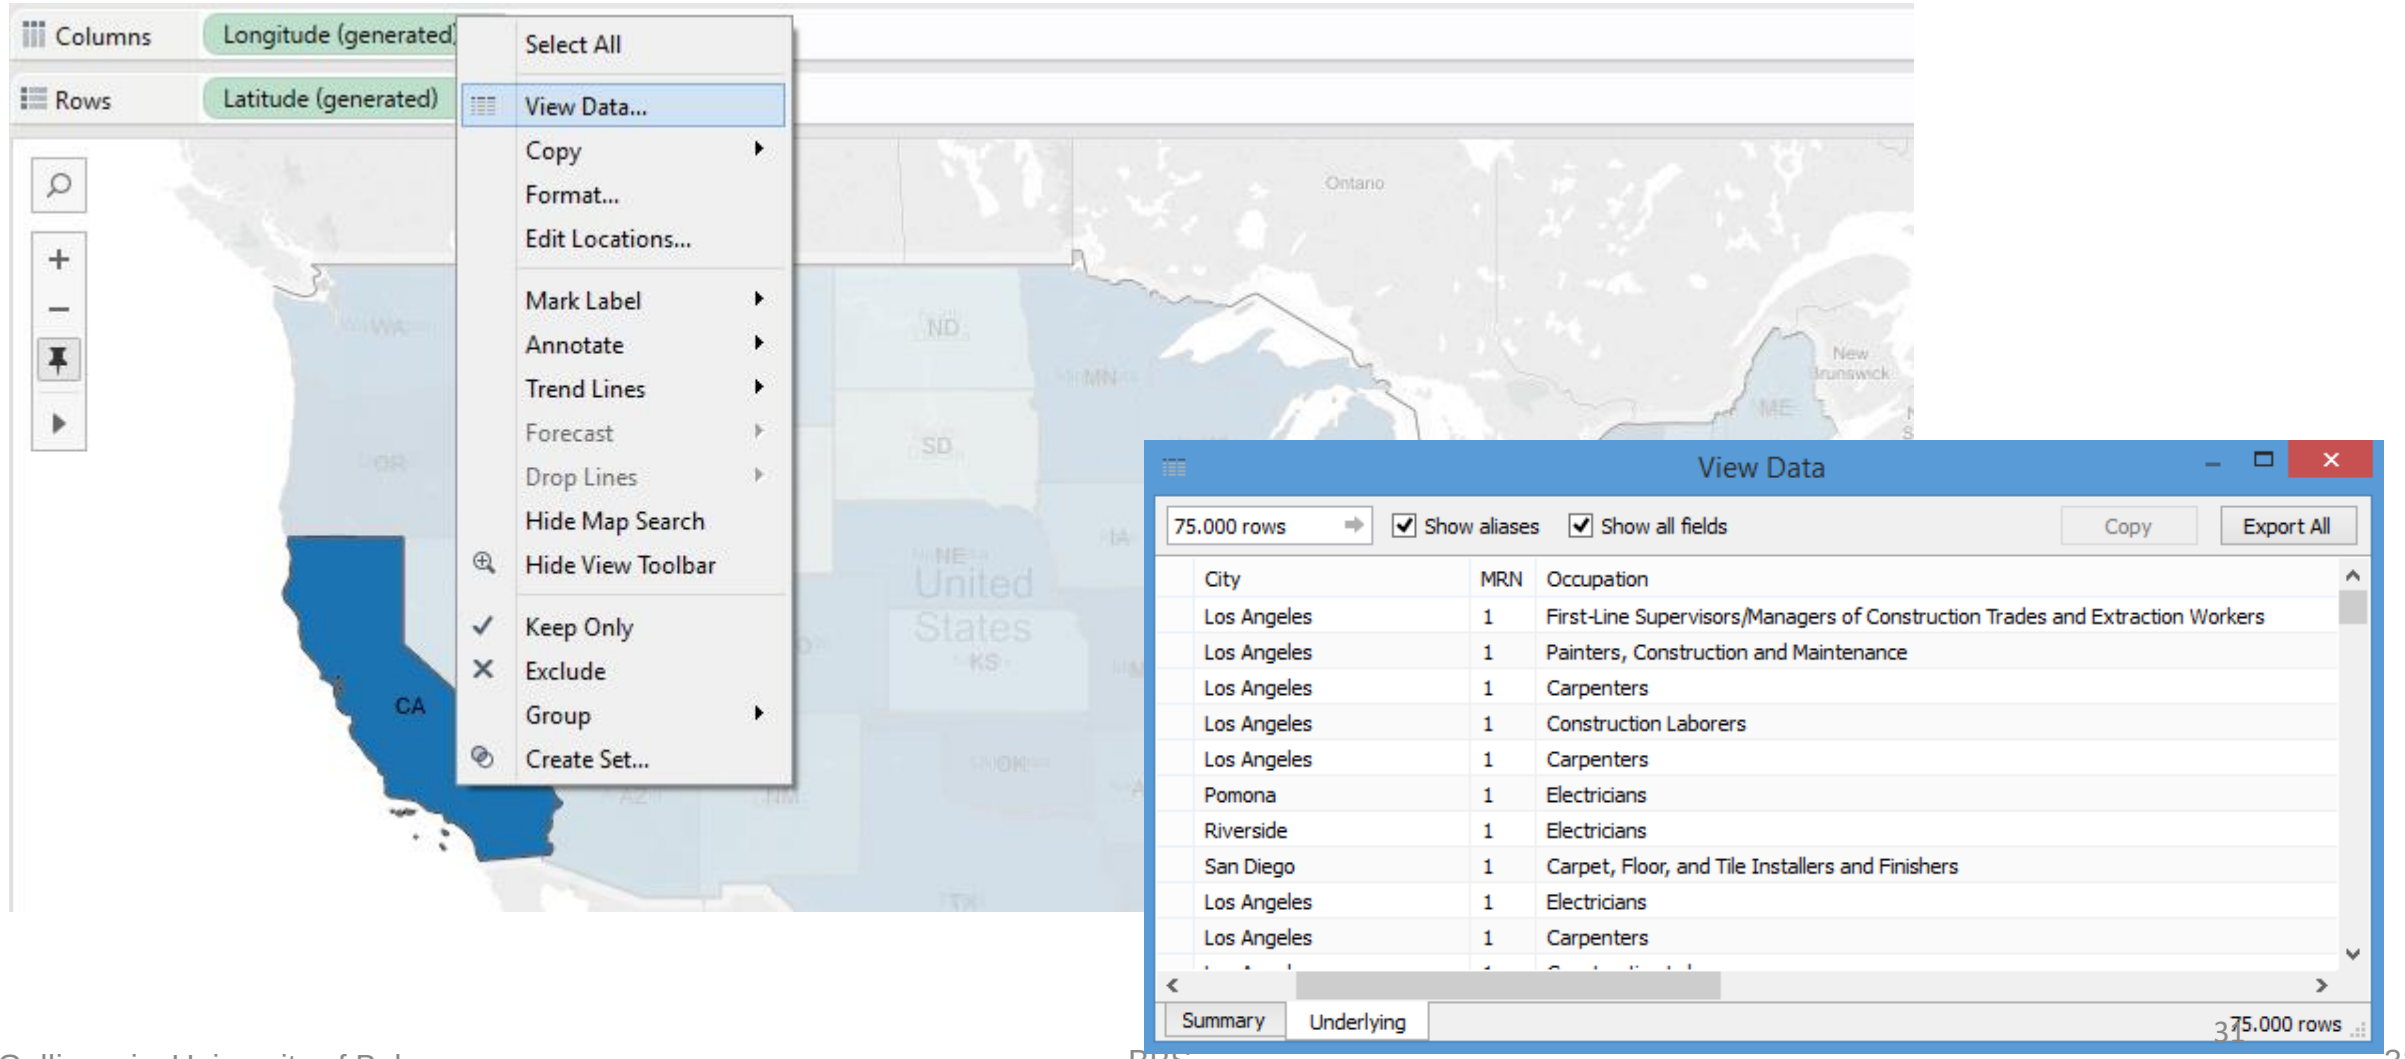
Columns
Longitude (generated)
Rows
Latitude (generated)
Select All
View Data...
Copy
Format...
Edit Locations...
Mark Label
Annotate
Trend Lines
Forecast
Drop Lines
Hide Map Search
Hide View Toolbar
Keep Only
Exclude
Group
Create Set...
CA
Ontario
United States
View Data
75.000 rows
Show aliases
Show all fields
Copy
Export All
City
MRN
Occupation
Los Angeles 1 First-Line Supervisors/Managers of Construction Trades and Extraction Workers
Los Angeles 1 Painters, Construction and Maintenance
Los Angeles 1 Carpenters
Los Angeles 1 Construction Laborers
Los Angeles 1 Carpenters
Pomona 1 Electricians
Riverside 1 Electricians
San Diego 1 Carpet, Floor, and Tile Installers and Finishers
Los Angeles 1 Electricians
Los Angeles 1 Carpenters
Summary
Underlying
75.000 rows

Filters

It is possible to apply a filter (i.e., *Slice & Dice*)

- To specific worksheets: the filter is applied only to the specified worksheets
- To a specific data source: the filter is applied to all the worksheets that use data from that source

It is possible to filter

- At *record level*: the view is computed using only the records that satisfy the filter; each filter is computed independently from the others
  - Eg. *Sales* > 100.00$
- At *aggregation level*: after that the view has been computed, all cells containing marks that do not satisfy the aggregation level filter are removed
  - Eg. SUM(*Sales*) > 100.00$

Filters: Context Filter

A *Context Filter* is a particular type of filter that is applied *before* record and aggregation level filters (i.e., the other filters depend on the result of context filters)

A context filter cannot be applied to aggregations (i.e., record level only)

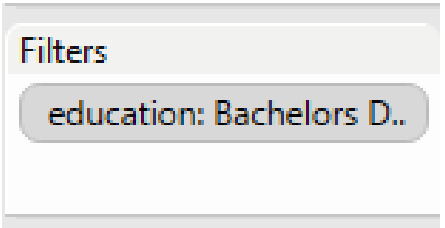

Context filters can be distinguished by their gray colour (for both discrete and continuous fields)

Pivoting

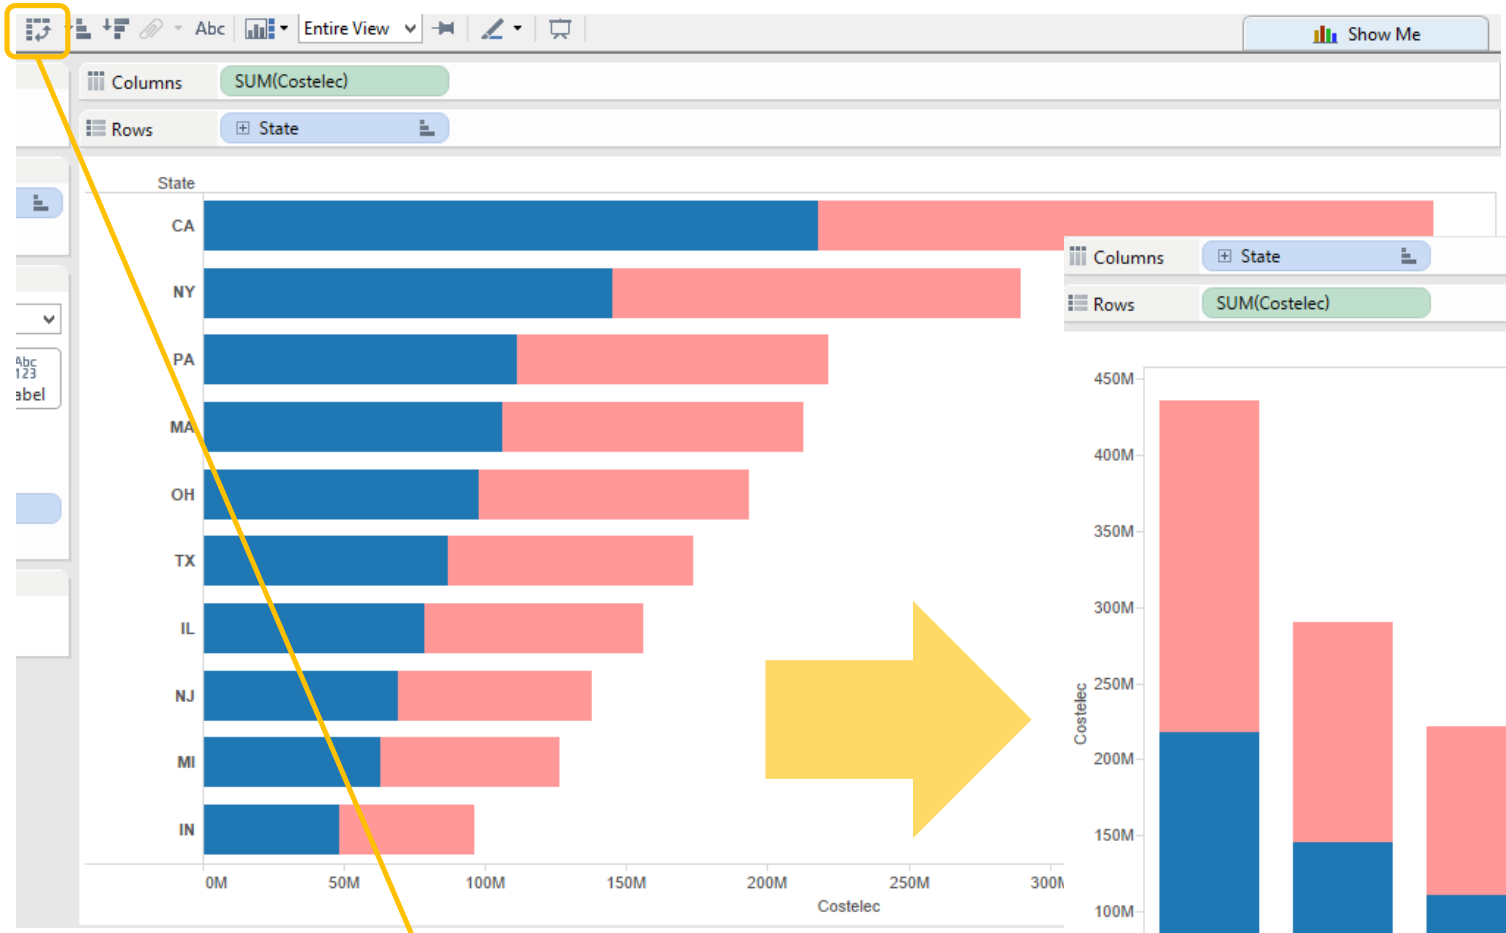

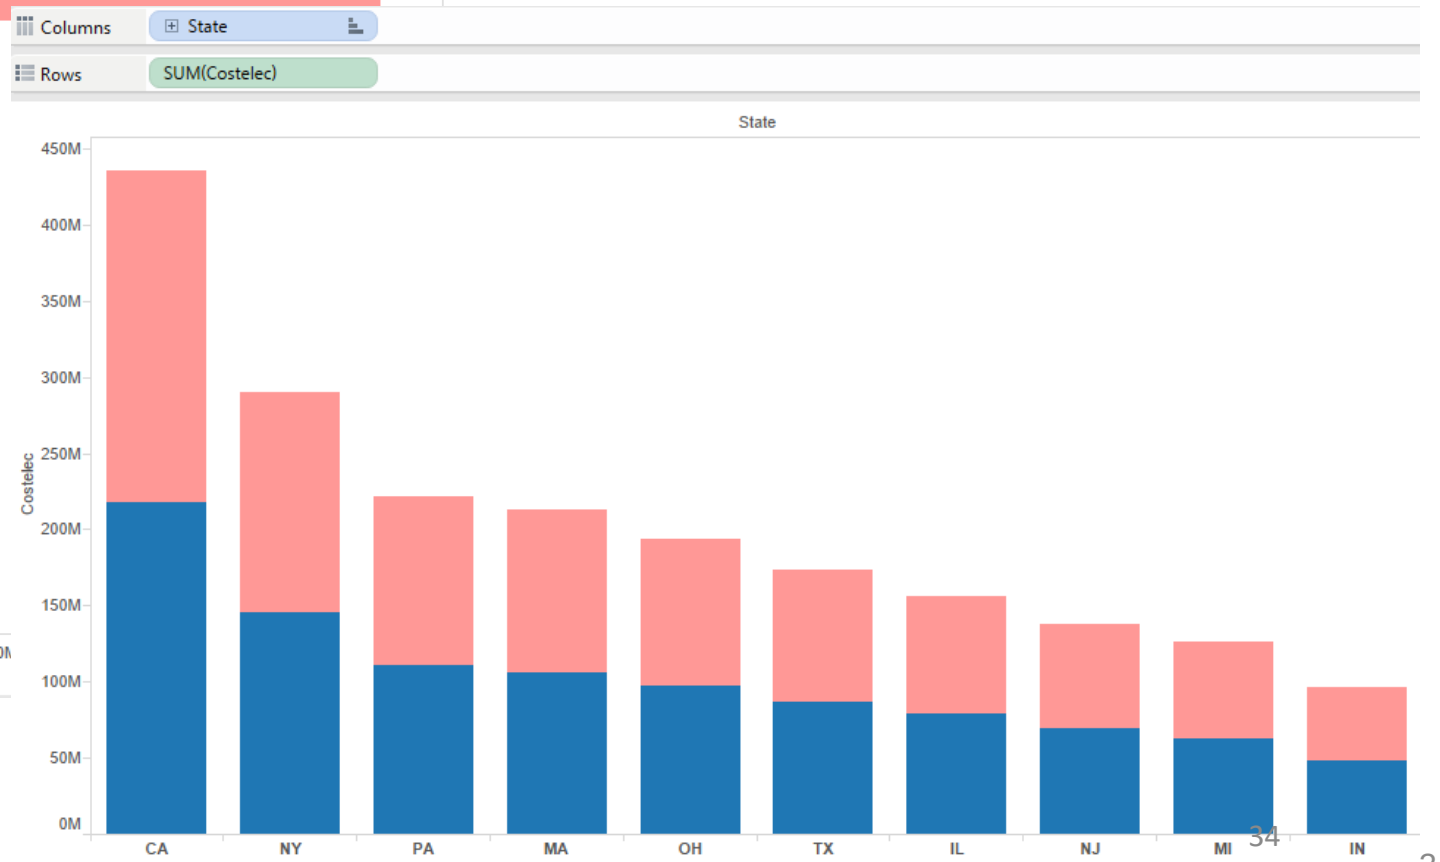

Pivoting (AKA *Swap*)

34

Grand Totals and Sub Totals

Totals are useful to show different aggregation levels in the same view

They can be applied only to discrete fields

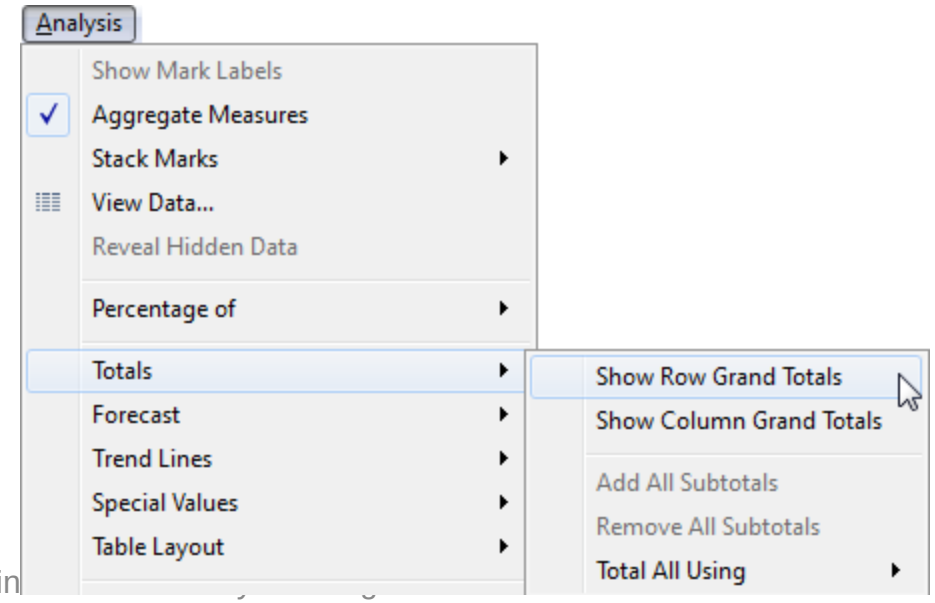

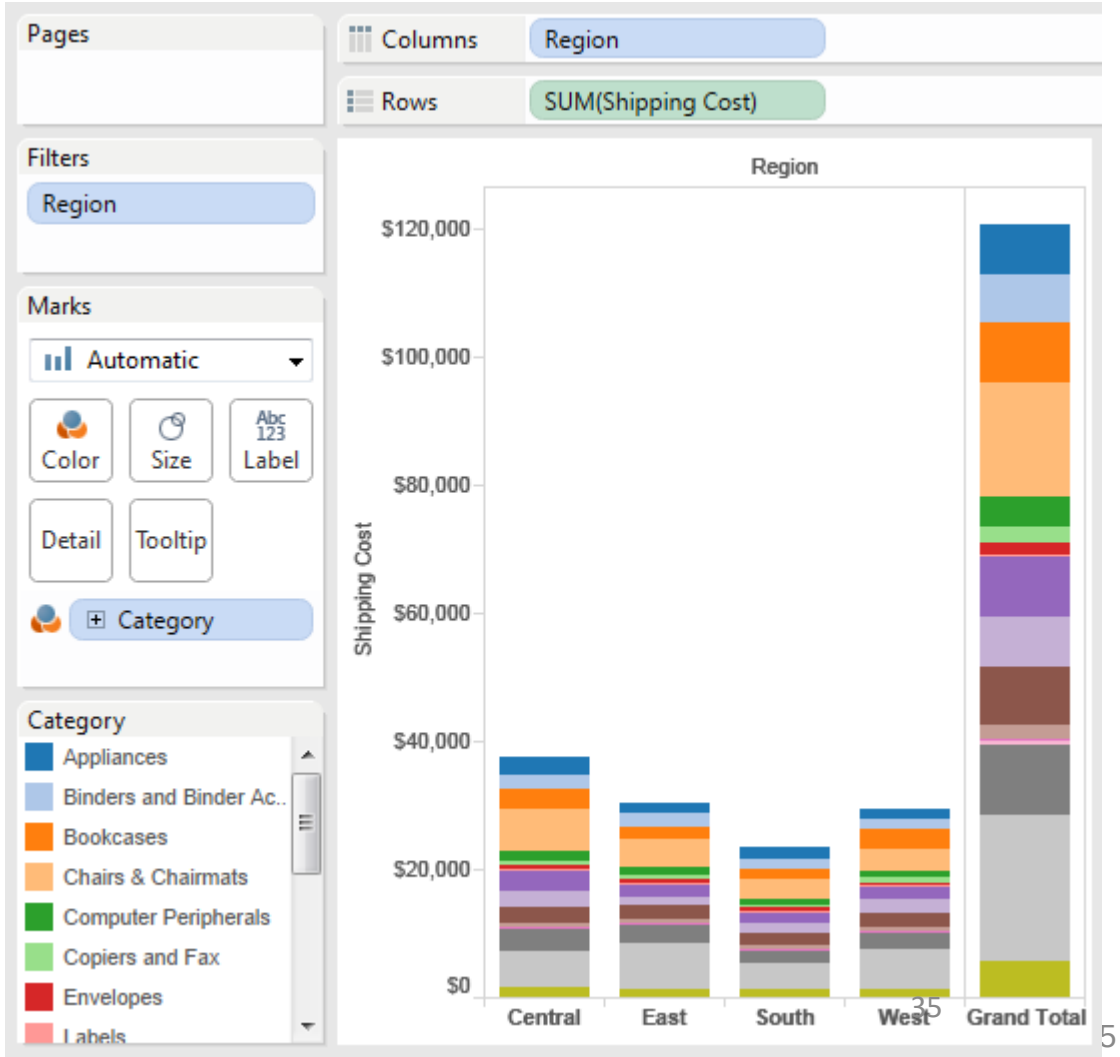

Show Me

The *Show Me* window contains shortcuts to create advanced visualisations starting from a set of dimensions and measures

Each visualisation type in the Show Me window has specific requirements: e.g., a scatter plot can be created only if there are at least continuous fields

Some types of visualisation can be quite difficult to build from scratch, for instance maps and box-plots

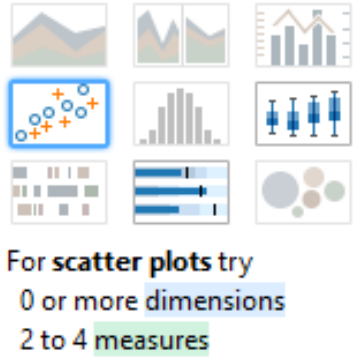

Top N

It is also possible to keep only the first (or last) *N* elements based on a given ordering

- E.g., the first 10 categories by sold amount

Attention: top / bottom filters are applied independently from other filters and from the view

- E.g., selecting a region while using a top 10 filter on categories would result in a visualisation that shows the categories whose sales are globally (i.e., across all the regions) the top 10!

Top N (2)

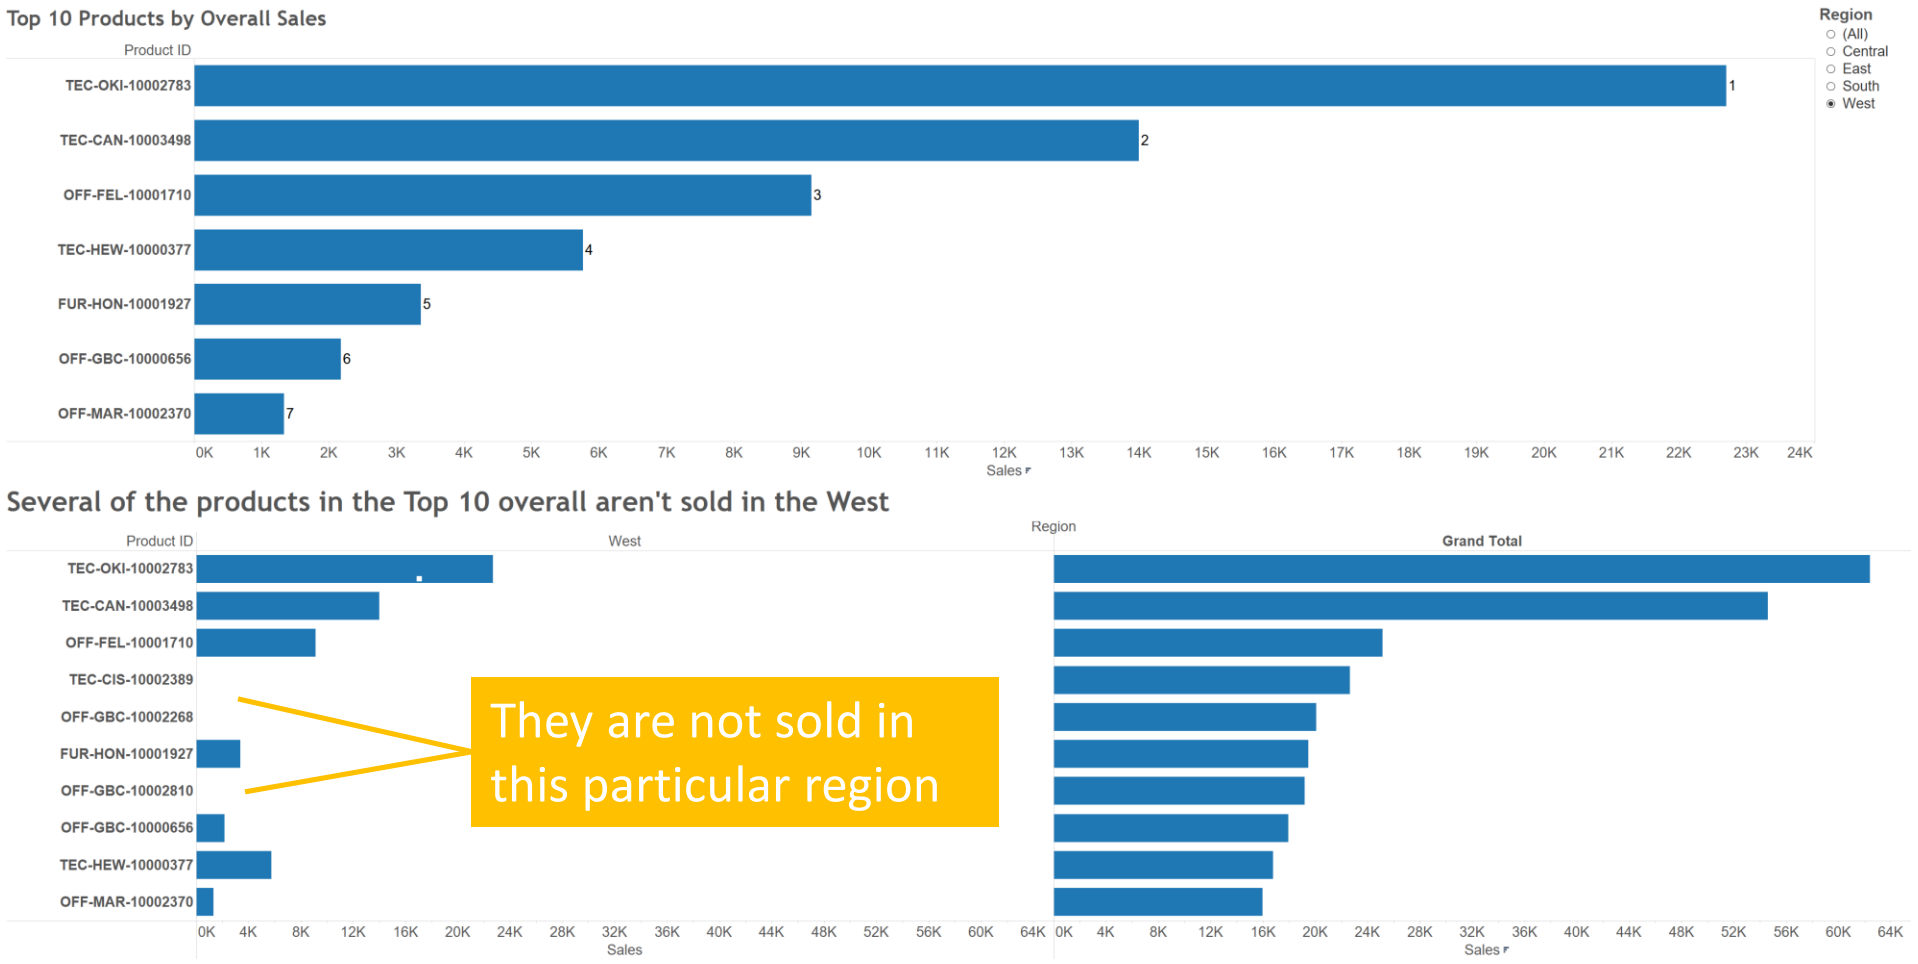
Top 10 Products by Overall Sales
Product ID
TEC-OKI-10002783
TEC-CAN-10003498
OFF-FEL-10001710
TEC-HEW-10000377
FUR-HON-10001927
OFF-GBC-10000656
OFF-MAR-10002370
1
2
3
4
5
6
7
0K 1K 2K 3K 4K 5K 6K 7K 8K 9K 10K 11K 12K 13K 14K 15K 16K 17K 18K 19K 20K 21K 22K 23K 24K
Sales
Region
(All)
Central
East
South
West
Several of the products in the Top 10 overall aren't sold in the West
Region
Product ID
West
Grand Total
TEC-OKI-10002783
TEC-CAN-10003498
OFF-FEL-10001710
TEC-CIS-10002389
OFF-GBC-10002268
FUR-HON-10001927
OFF-GBC-10002810
OFF-GBC-10000656
TEC-HEW-10000377
OFF-MAR-10002370
They are not sold in this particular region
0K 4K 8K 12K 16K 20K 24K 28K 32K 36K 40K 44K 48K 52K 56K 60K 64K
Sales

Binning

Through binning it is possible to create *bins* (or *buckets*), which are numerical ranges over the values of a numerical attribute

- In Tableau, ranges are left-inclusive and right-exclusive: [*start*, *end*)

Binning is useful to *discretize* continuous attributes while controlling the number of shown values; it is also possible to further group discrete numerical attributes to reduce their number (i.e., show fewer labels)

Binning can also be used as a *smoothing* mechanism, that is, to reduce the effect small variations considered noise

Histograms are created by first binning and then counting the number of elements belonging to each bin

Measures are aggregated by bin

Tableau only supports *equi-width* bins, i.e., each bin has the same range length

Binning (2)

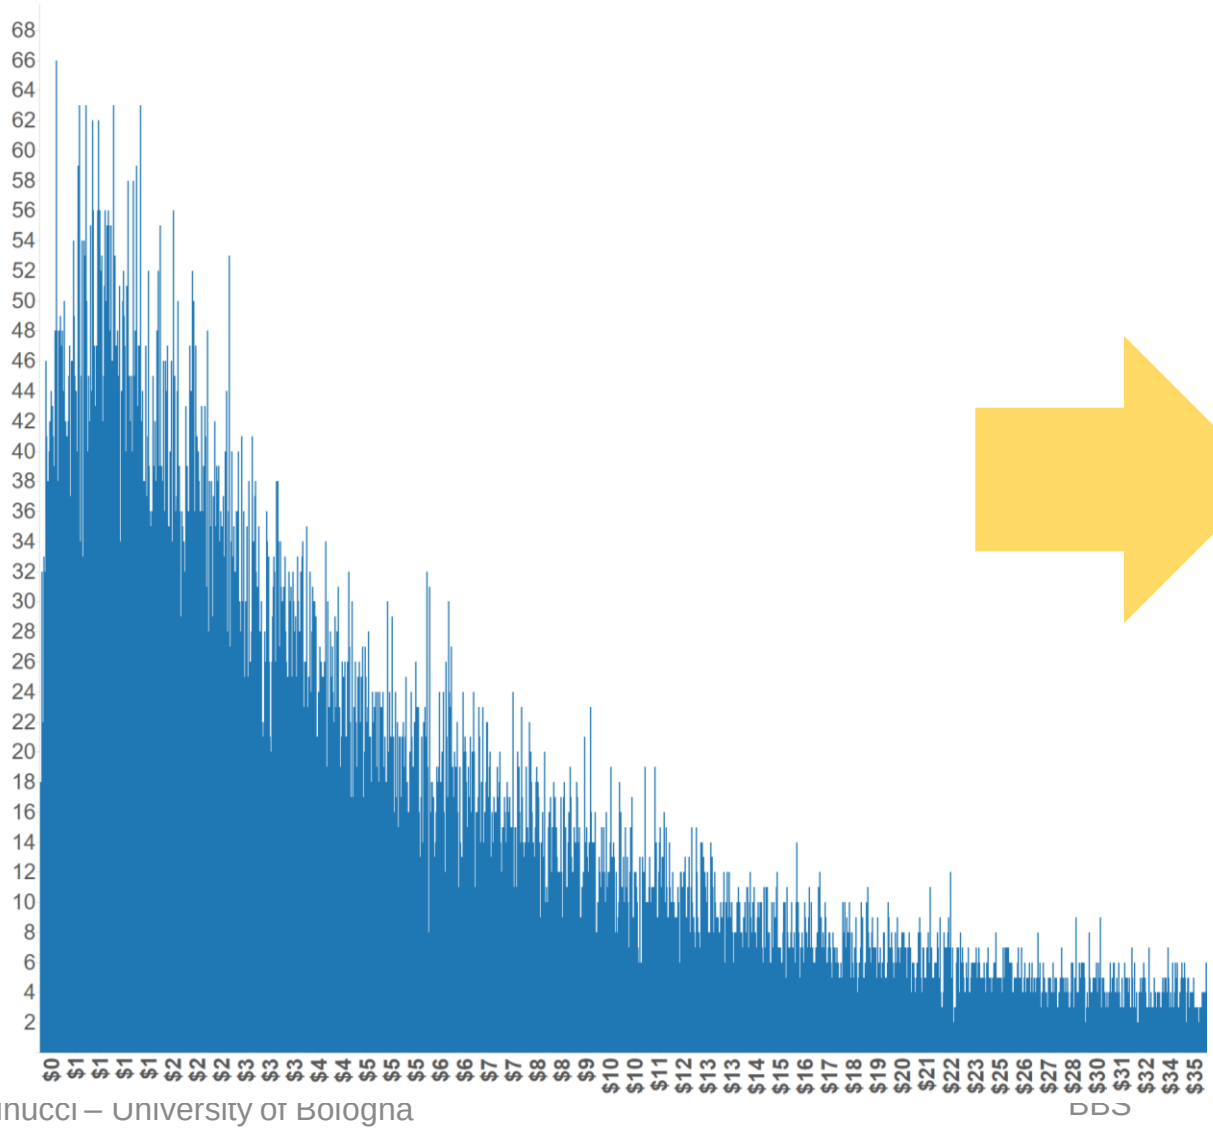

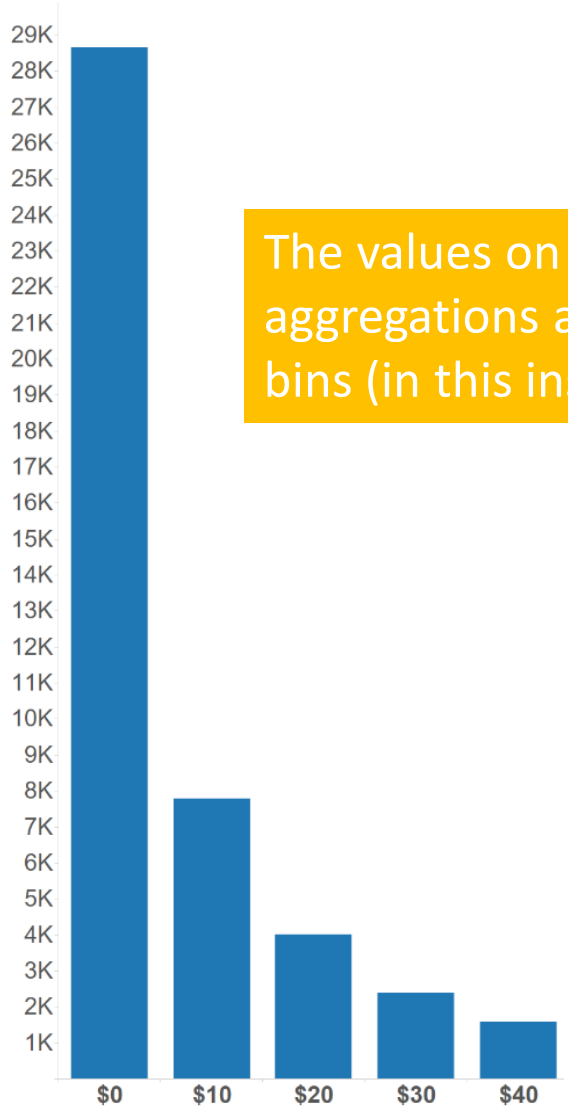

The values on the y-axis are aggregations associated to the bins (in this instance, *sum*)

Group

It is possible to create new attributes by (manually) grouping values of already existing ones; especially useful for

- Grouping values that have the same semantic but are labeled differently
- Creating new intermediate levels of aggregation without modifying the data source

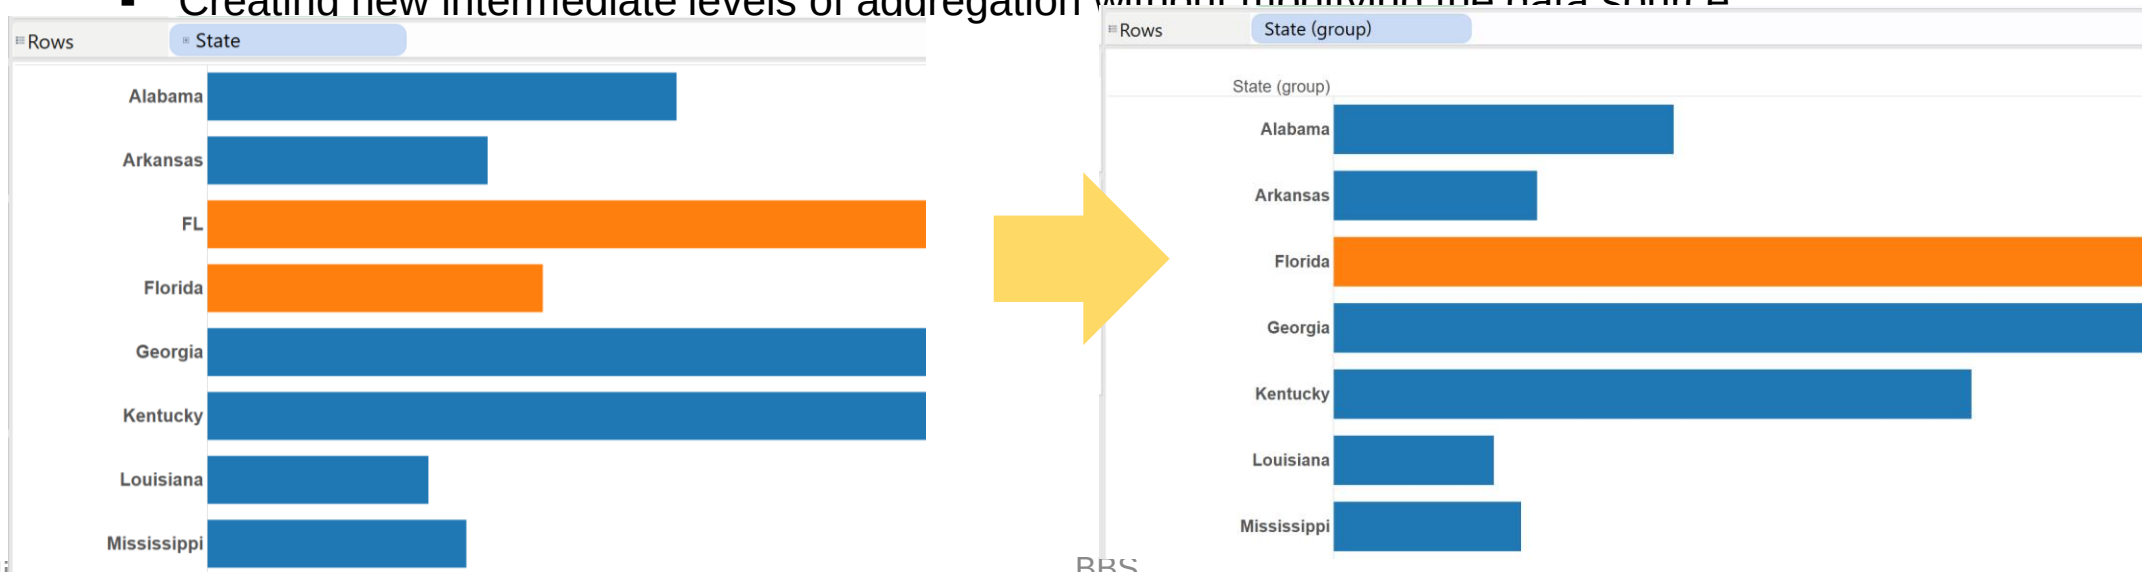

41

Set

Sets are custom fields that allow to create subsets of attribute values (cannot be used with measures) that satisfy a certain condition. Based on the defined condition, sets can be

- *Constant*: the elements of the set must be manually selected; the set remains the same even when underlying data change
- *Computed*: the set is created by defining a condition (e.g., sales > 100) that the elements must satisfy

Sets can be used as filters, hierarchy levels, inside calculated fields, …

Different sets can be *combined* to create new ones

- For instance, it is possible to create the union of two sets

Set (2)

When used as filters, all elements that do not belong to the set are discarded

When used as a dimension, a set creates a new header composed by values *In* and *Out*

Similarly, when assigned to the colour property, markers are split into two (one colour for *In* and one for *Out*)

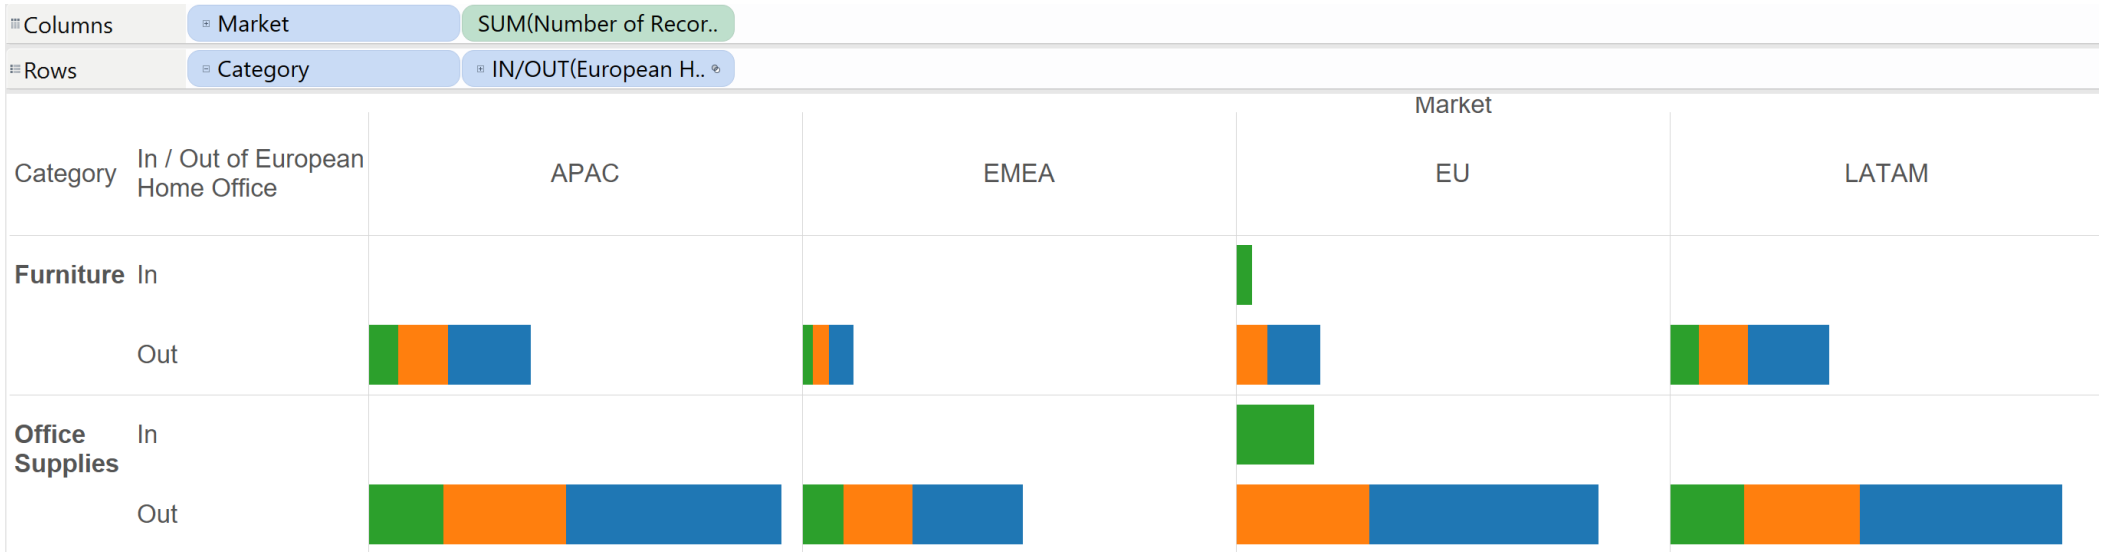

Calculated Field

A *calculated field* is a field defined by a formula that can make use of other fields and various functions such as arithmetic and logical operations

- E.g., the *Profit* field can be defined as *Sales* - *Cost*

Through a calculated field it is possible to create new fields without modifying the data source

A calculated field can be defined at different granularity levels

- *Line Granularity*: the field is computed tuple by tuple (e.g., profit)
- *Aggregated Granularity*: the is computed over other aggregations

A calculated field can be used as any other field; an exception are calculated fields with aggregated granularity, which can be used as filters only if they are continuous fields

Calculated Field: Syntax

To use an existing field inside a formula it is sufficient to type its name surrounded by square brackets

- e.g., `[Sales]`

Conditional statements

```
if [Profit] > 0 then
    'Profitable'
else
    'Not profitable'
end
```

A function can be called by simply typing its name followed by its parameters inside parentheses; the parameters must be separated by commas

- e.g., `MAX([Sales], [Cost])`

Calculated Field: Line Gran

Before computing the view, a new value is computed for each tuple

The SUM aggregation is computed over the calculated field

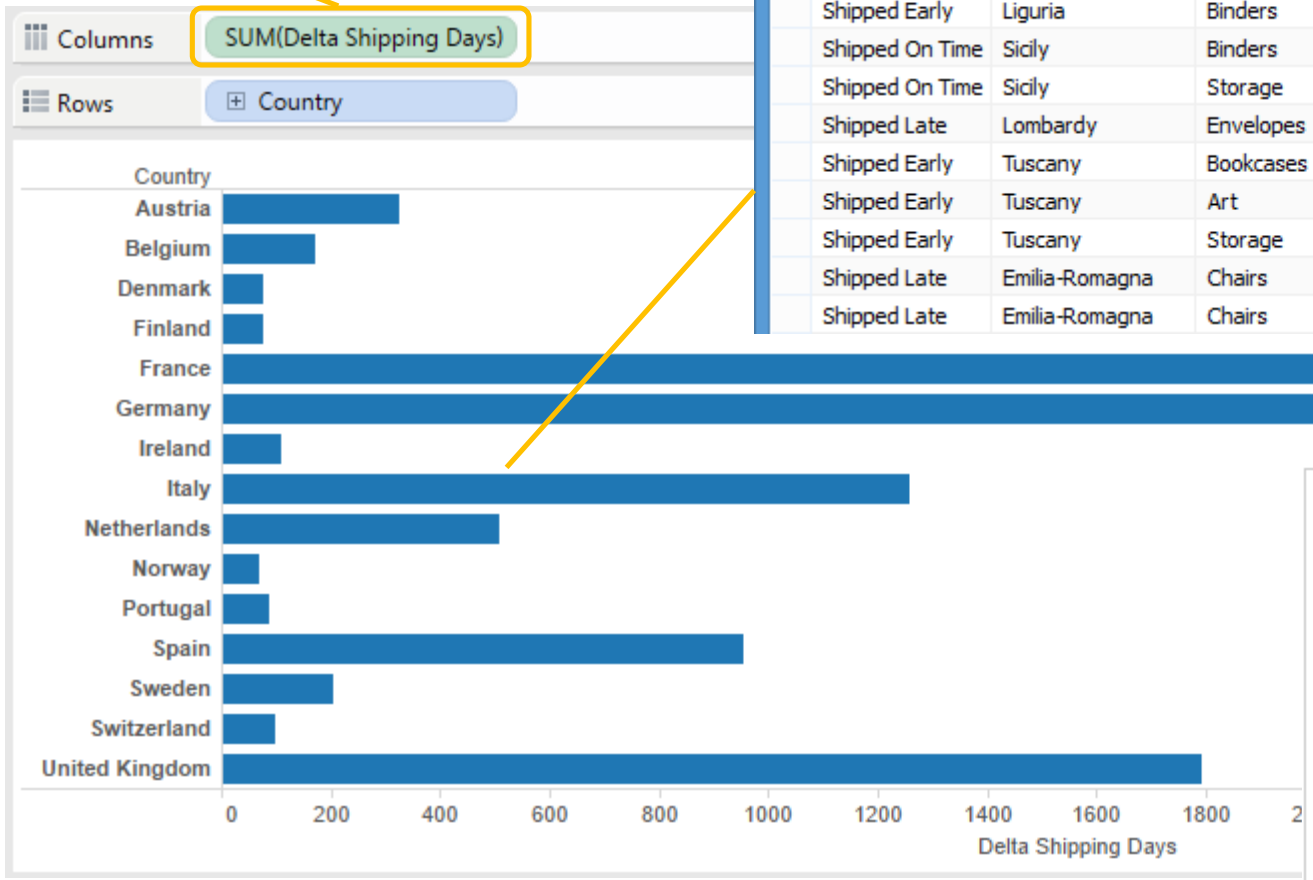

View Data

1,108 rows | Show aliases | Show all fields | Copy | Export All

Ship Status	State	Sub-Category	Days to Ship Actual	Days to Ship Scheduled	Delta Shipping Days	Discount
Shipped Early	Liguria	Binders	5	6	1	0.000000
Shipped On Time	Sicily	Binders	1	1	0	0.000000
Shipped On Time	Sicily	Storage	1	1	0	0.400000
Shipped Late	Lombardy	Envelopes	2	1	1	0.000000
Shipped Early	Tuscany	Bookcases	4	6	2	0.000000
Shipped Early	Tuscany	Art	4	6	2	0.000000
Shipped Early	Tuscany	Storage	4	6	2	0.400000
Shipped Late	Emilia-Romagna	Chairs	2	1	1	0.600000
Shipped Late	Emilia-Romagna	Chairs	2	1	1	0.600000

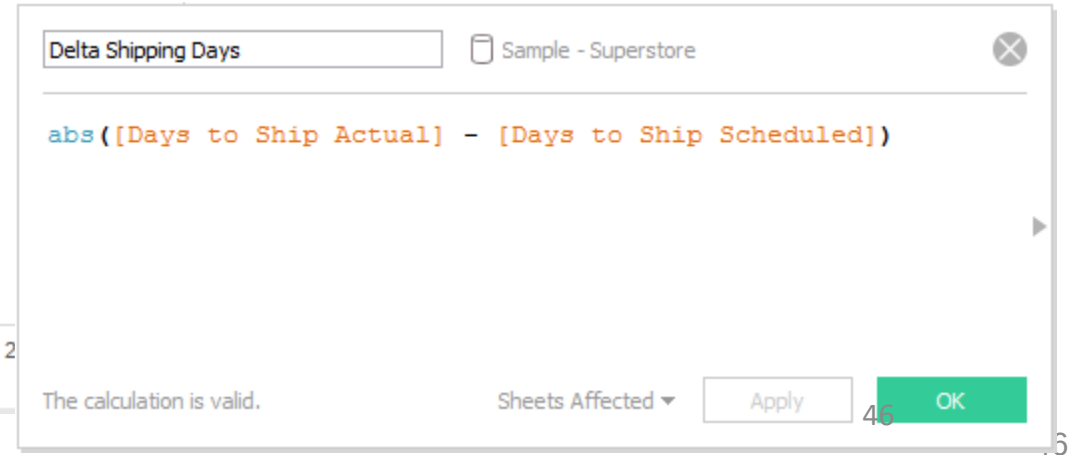

Calculated Field: Aggregated Granularity

The calculated field implicitly defines an aggregation (denoted as a generic AGG)

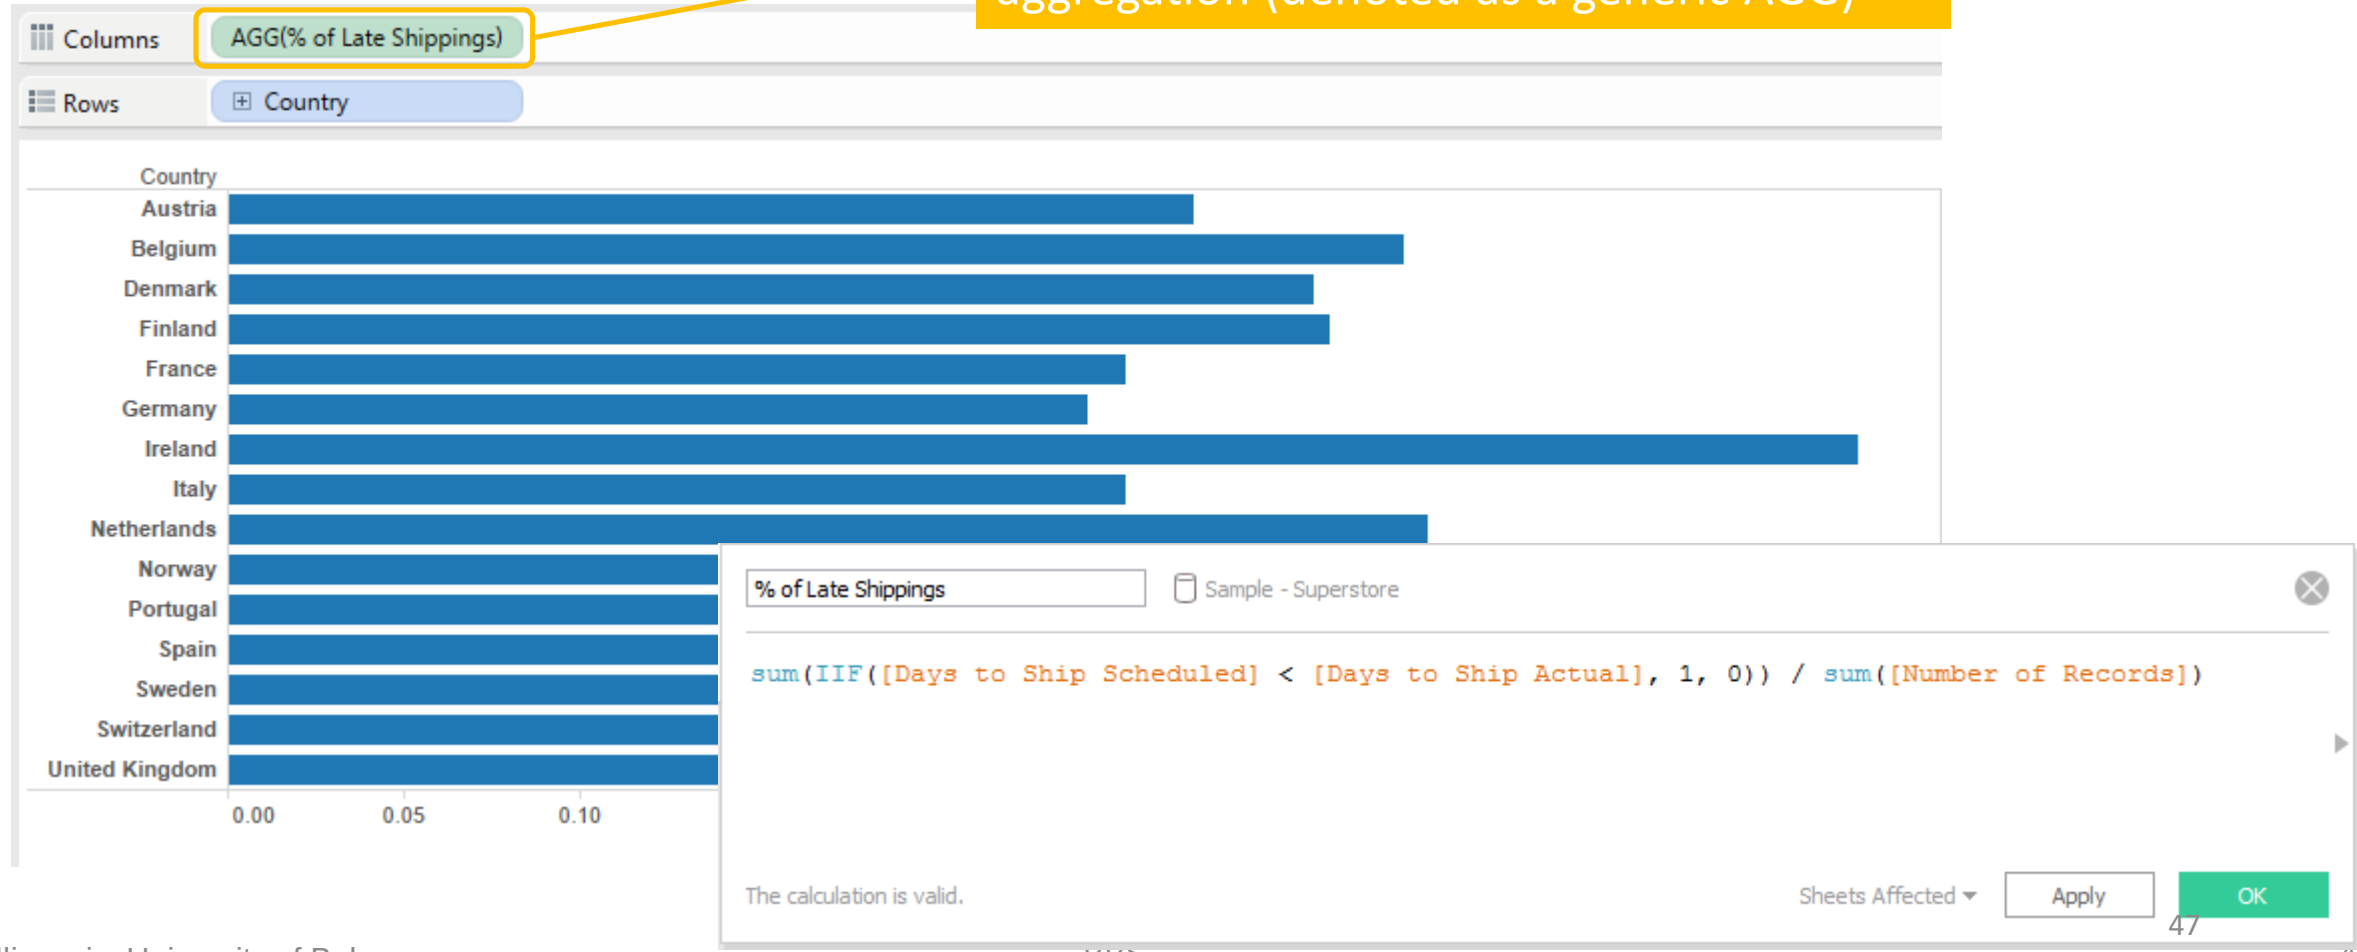

More functionalities: calculated fields

Can be defined at different granularities

- Line granularity: e.g., *[Extendedprice]*[Quantity]*
- Aggregated granularity: e.g., *SUM(IF ([Tax]>0) THEN 1 ELSE 0 END)/COUNT([Quantity])*

Categorical fields can be defined as well

- E.g., *IF ([Tax]>0) THEN 'Taxed' ELSE 'Not taxed' END*

Level-of-detail (LOD) expressions: fix the reference group-by attributes

- Define reference aggregated values
- Useful to create more advanced queries combining more aggregations

More functionalities: LOD fields

Nested aggregation

- Further aggregate the results of a query
- For instance:
  - Calculate the sum(Quantity) by Brand and Nation
  - Calculate the average of the result by Brand
  - TotalCustomerSales =
    {INCLUDE [Nation (Customer)]: SUM([Quantity])}

TotalCustomerSales is defined at a *finer* aggregation level

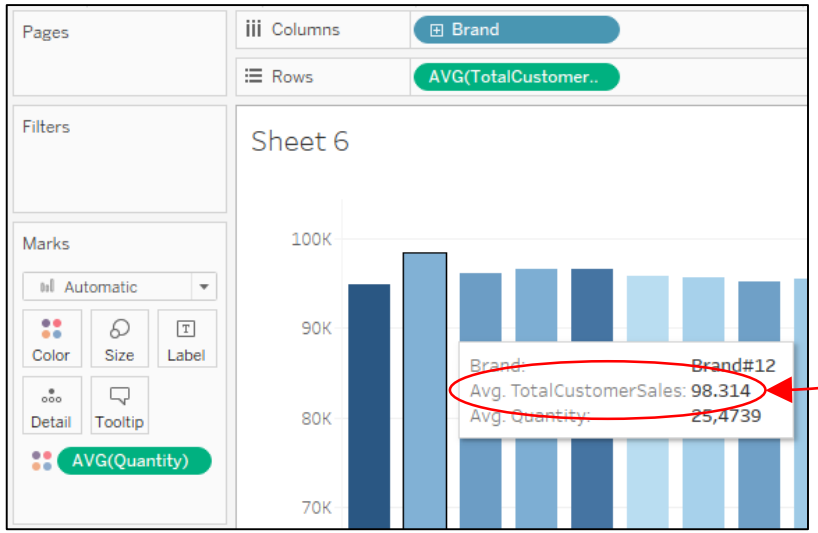

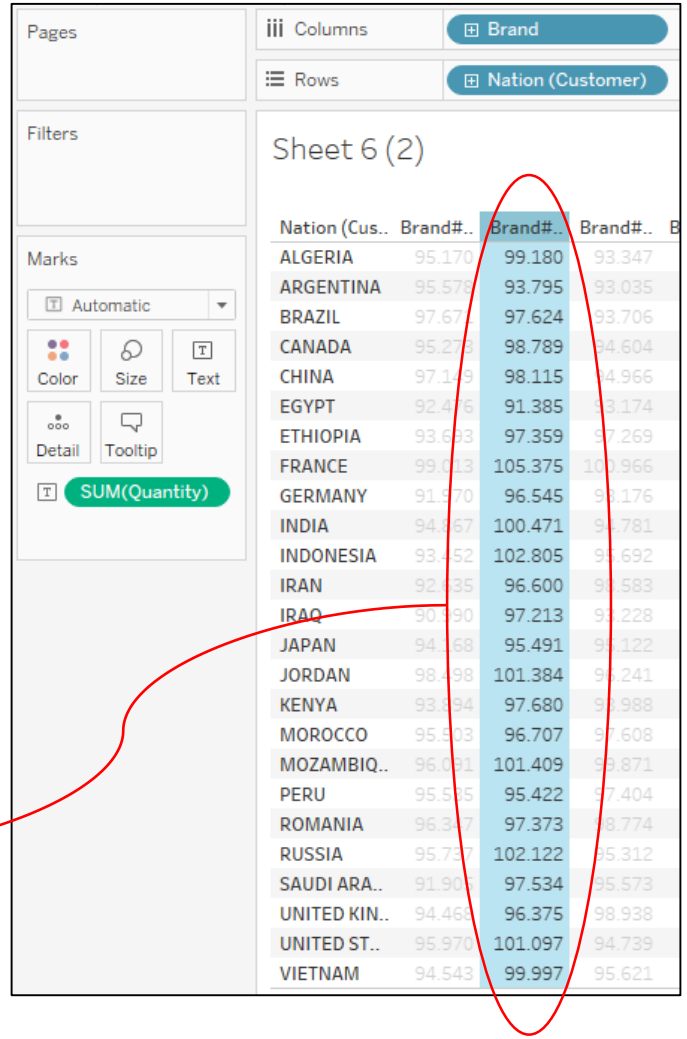

More functionalities: LOD fields

Use coarser data at a finer aggregation level

- For instance:
  - Calculate the sum(Quantity) by Nation (and Region)
  - Divide the result by the sum(Quantity) by Region
  - PercWithinRegion =
    {FIXED [Nation (Customer)]: SUM([Quantity])} / {FIXED [Region (Customer)]: SUM([Quantity])}

PercWithinRegion embeds a calculation at a *coarser* aggregation level

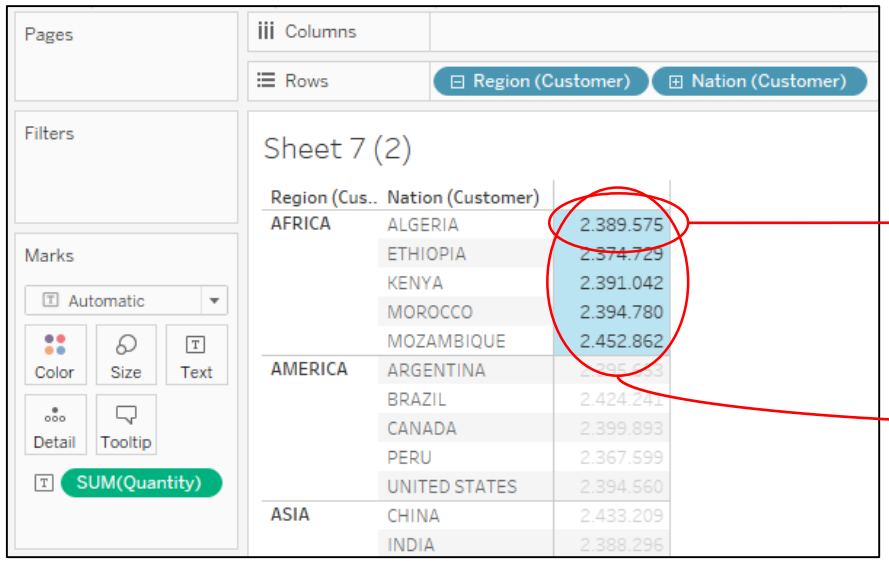

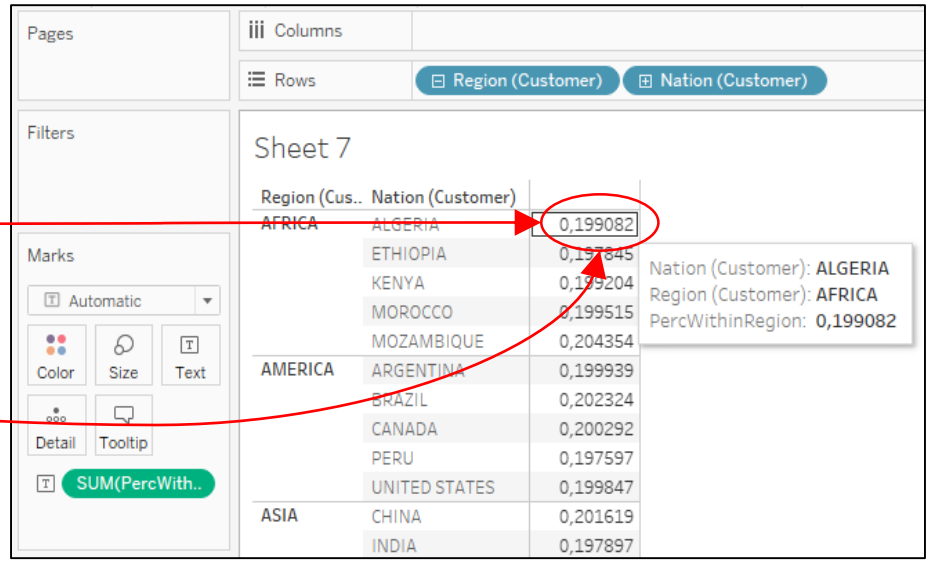

References

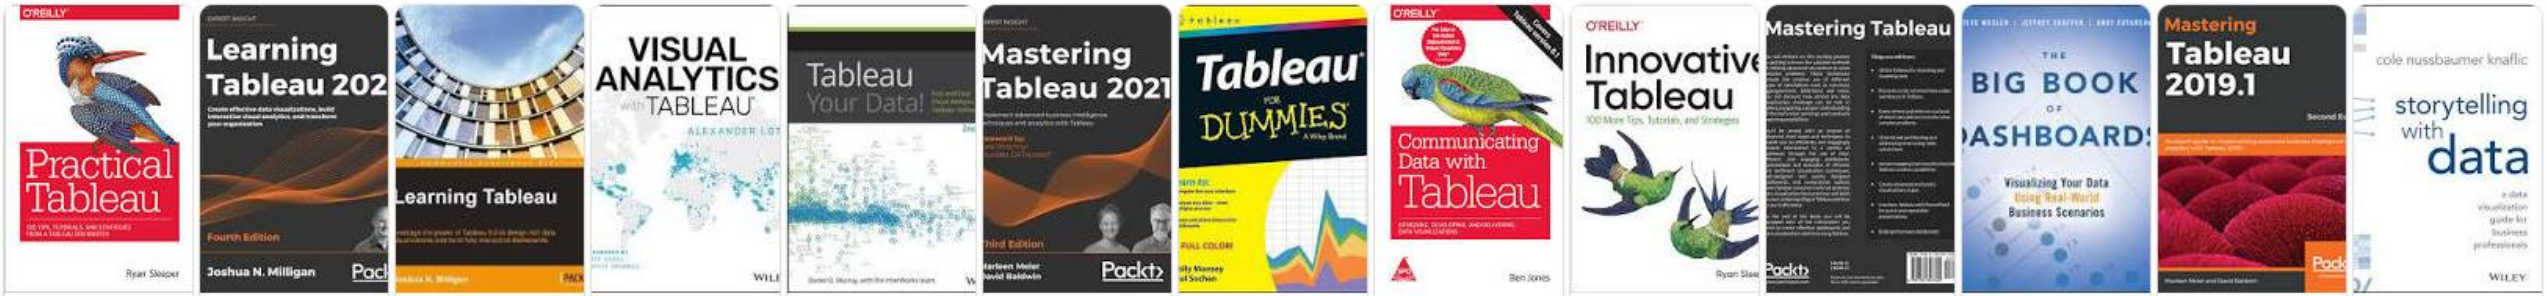

Practical Tableau: 10… Ryan Sleep…

Learning Tableau 20… Joshua N. …

Learning Tableau Joshua N. …

Visual Analytics wi… Alexander …

Tableau Your Data! Fast … Daniel G. M…

Mastering Tableau 20… Marleen Me…

Tableau For Dummies Pavol Socháň

Communic… Data with T… Ben Jones, …

Innovative Tableau: 10… Ryan Sleep…

Mastering Tableau David Bald…

The Big Book of Dashboa… Steve Wexl…

Mastering Tableau 20… Marleen Me…

Storytelling With Data: … Cole Nussb…

Exercises - Part I

Connecting to a source

Connectors

Tableau is able to connect to different data sources

- Text files (DSV) and Excel
- Relational DBMSs (e.g., MySQL, Oracle and SQL Server)
- Big Data (e.g., Hive, Spark SQL and Impala)
- Etc.

Tableau needs a *tabular view* of the data (i.e., a set of tables)

- Different tables from the same source can be linked through joins
- JSON data can be converted into a tablar view as well

Join

Tableau supports different types of join

- *Inner*: only the records that satisfy the given condition are returned
- *Left* (*Right*): corresponds to an inner join where there are also the tuples belonging to the left (right) table that do not satisfy the given condition
- *Full Outer*: corresponds to

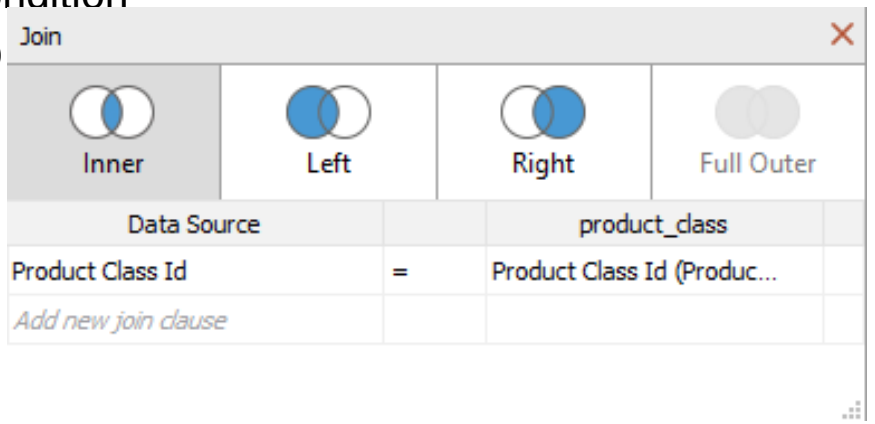

joins

Example – MySQL

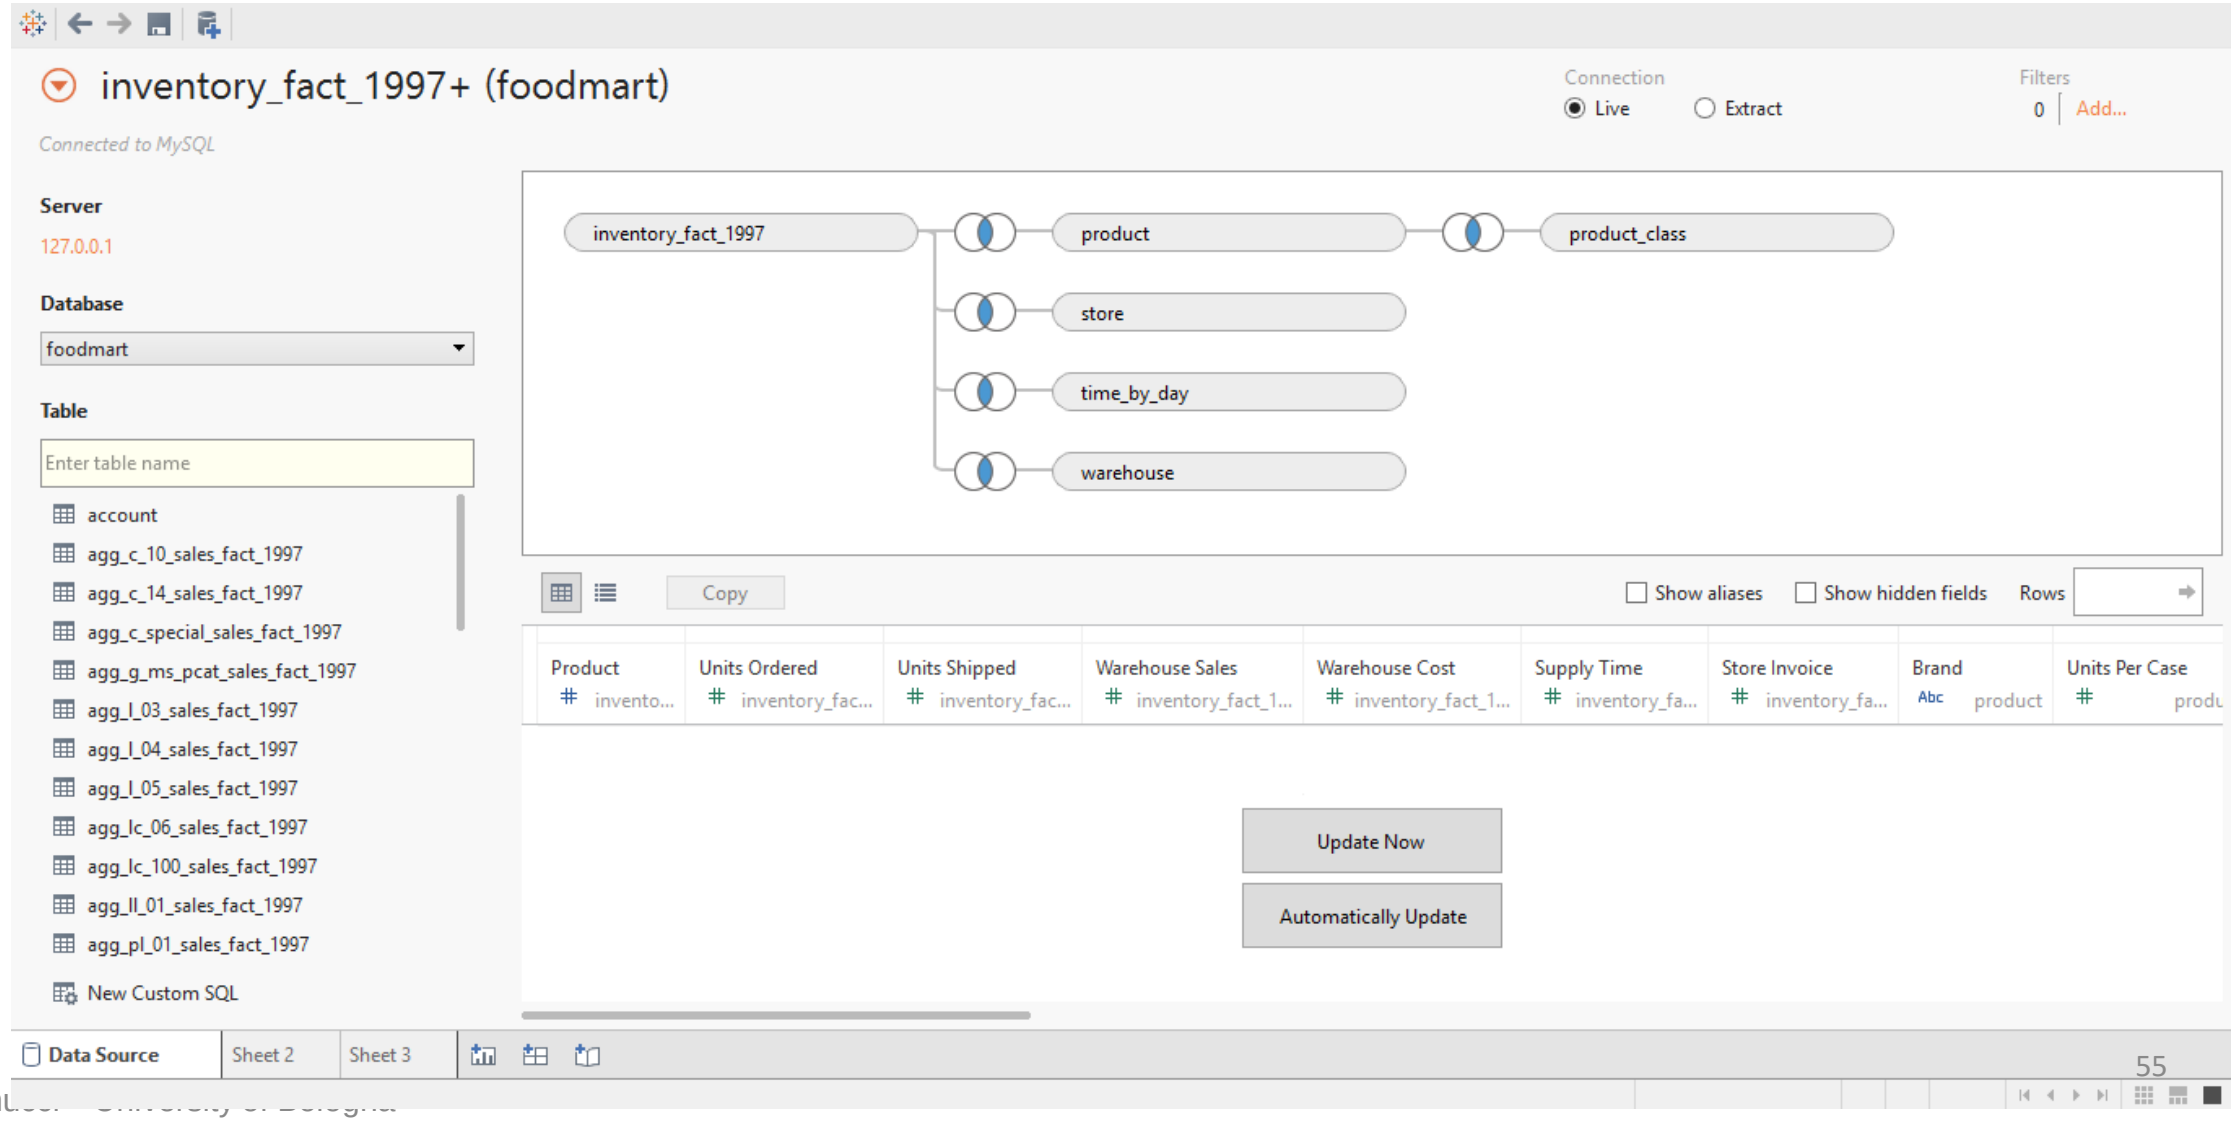

55

Example – Text File (DSV)

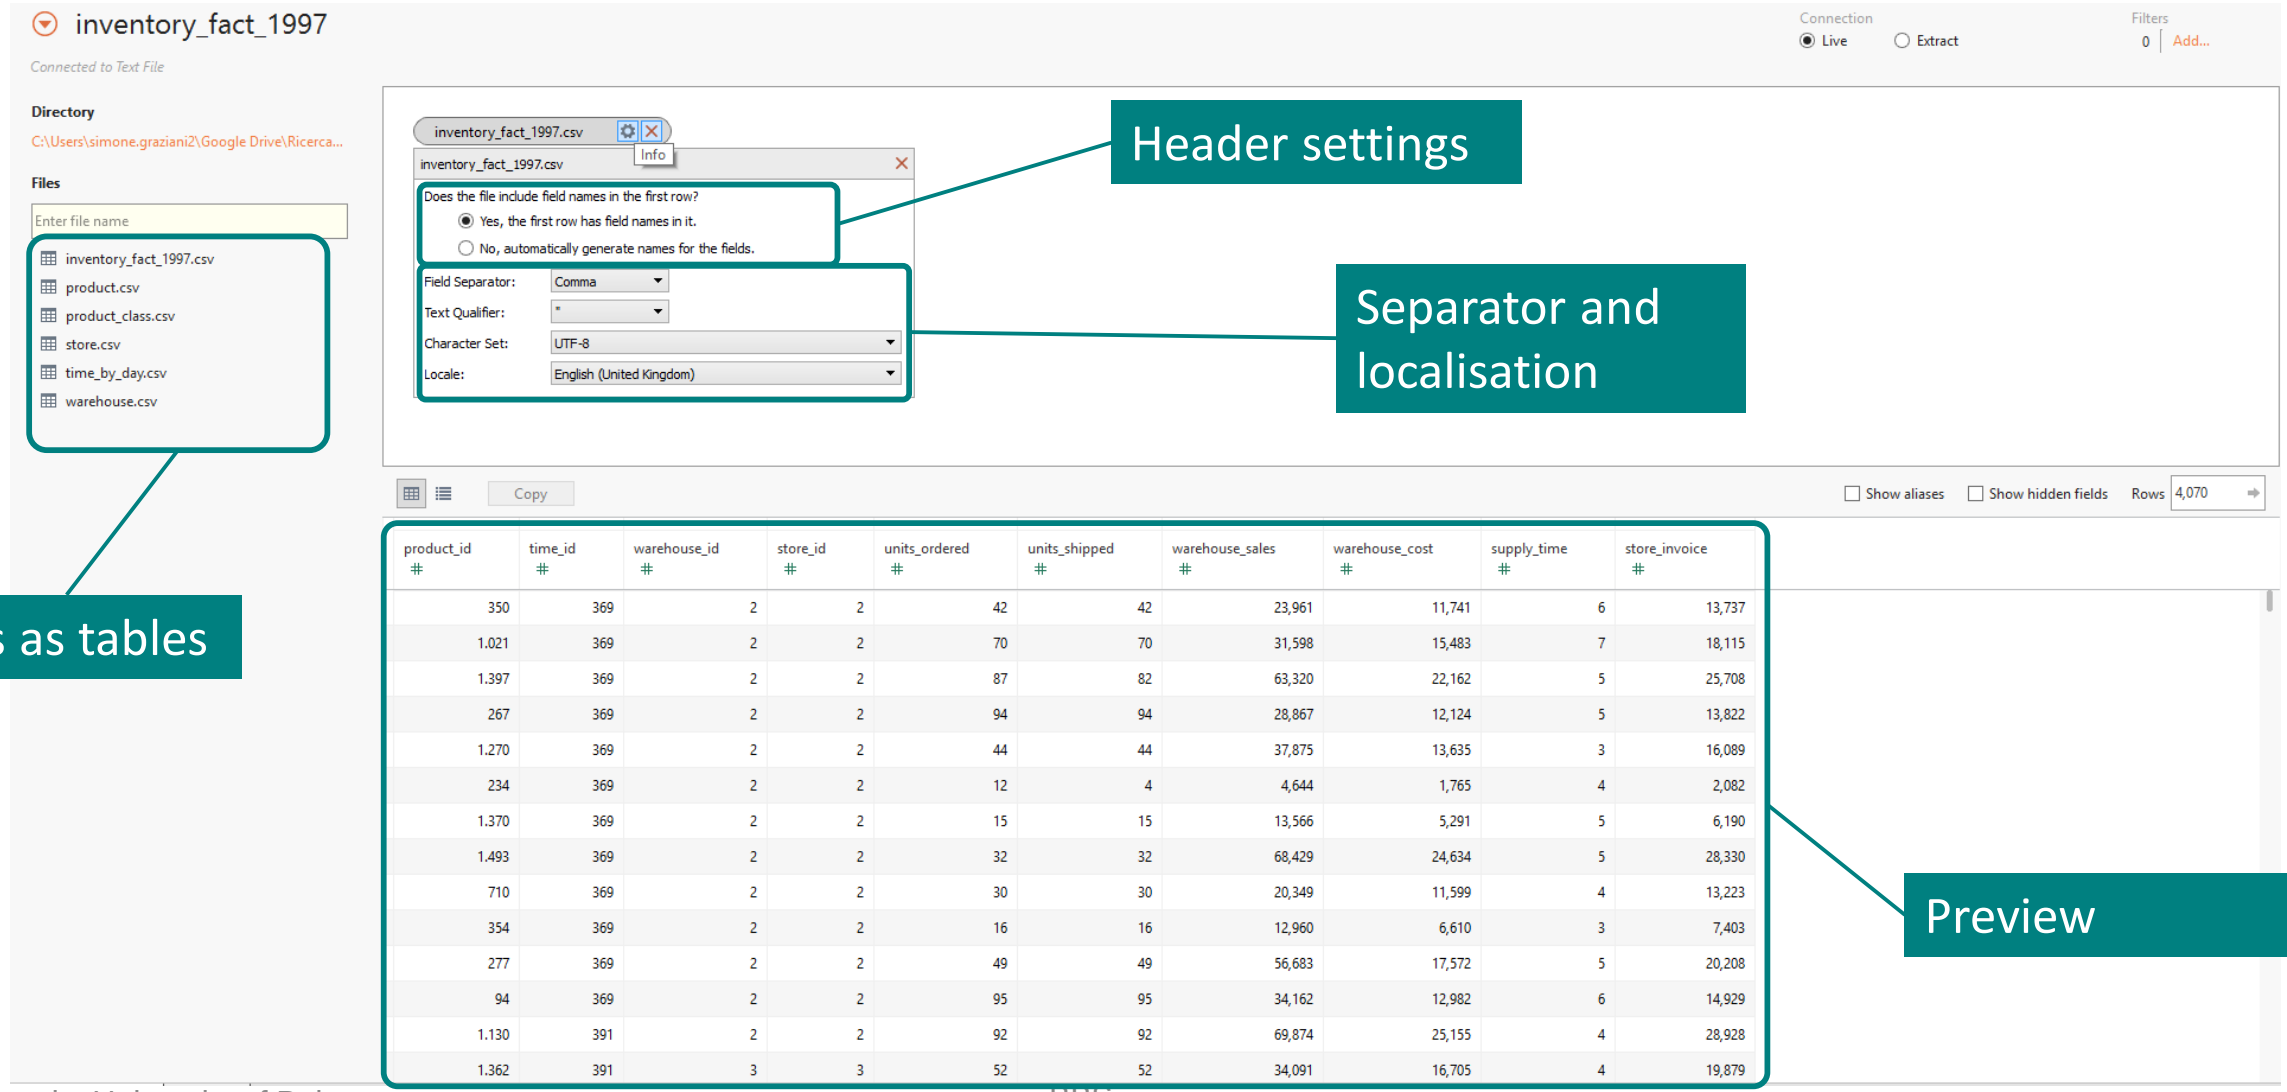

Example – Union of Text Files

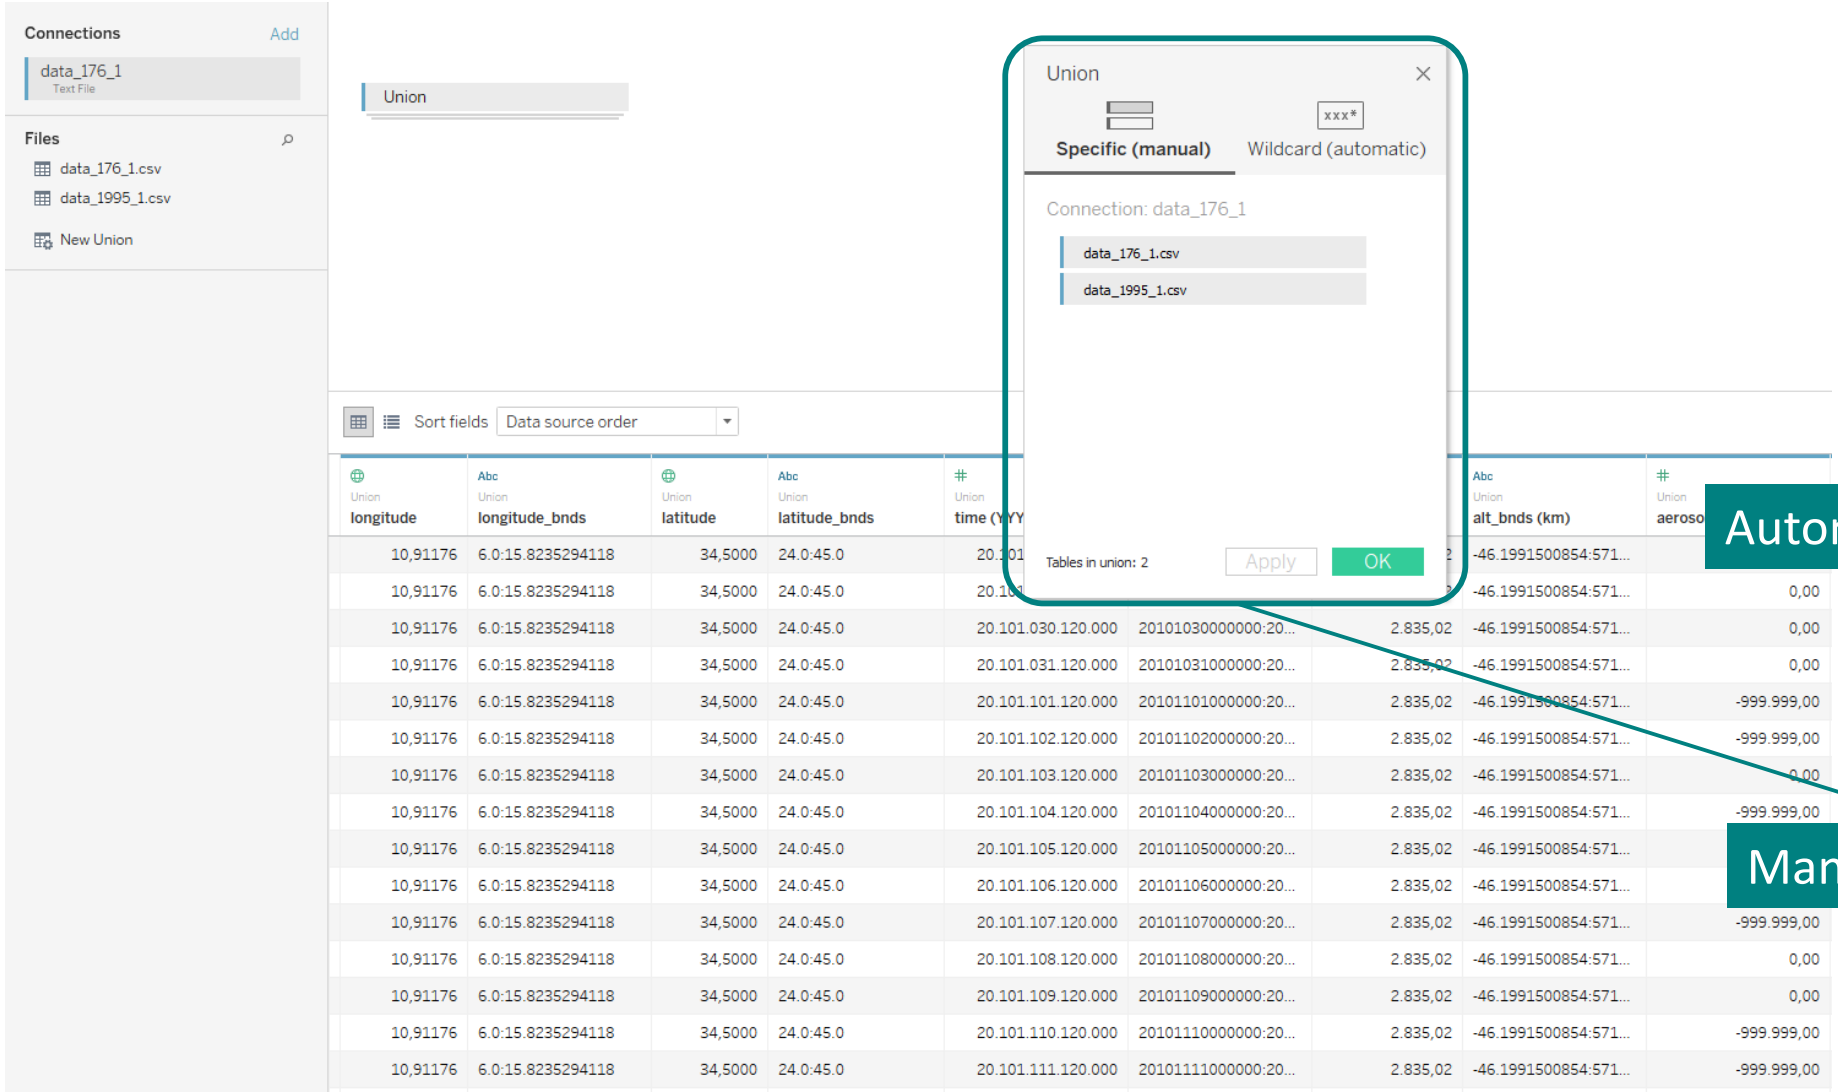

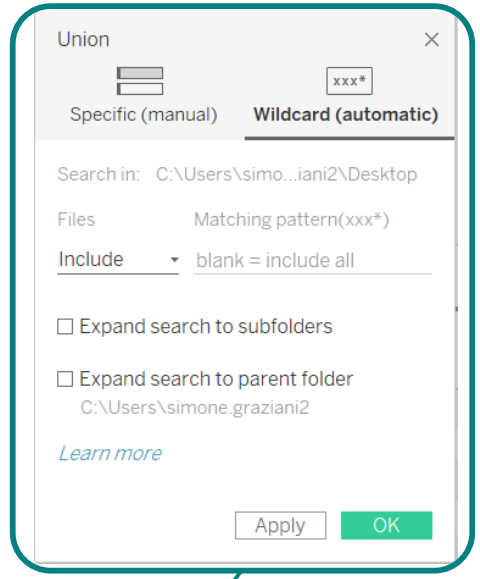

Automatic selection

Manual selection

Custom SQL

With data sources that support SQL it is also possible to define a table through a *custom query*

Table

Live VS Extract

There are two different types of connection

- *Live*: each time that a visualisation is created (or edited) Tableau performs a query to the data source
- *Extract*: Tableau performs a single (big) query at the beginning to pull all the required data locally; by exploiting the (locally) stored data, there is no need to issue additional queries the data source

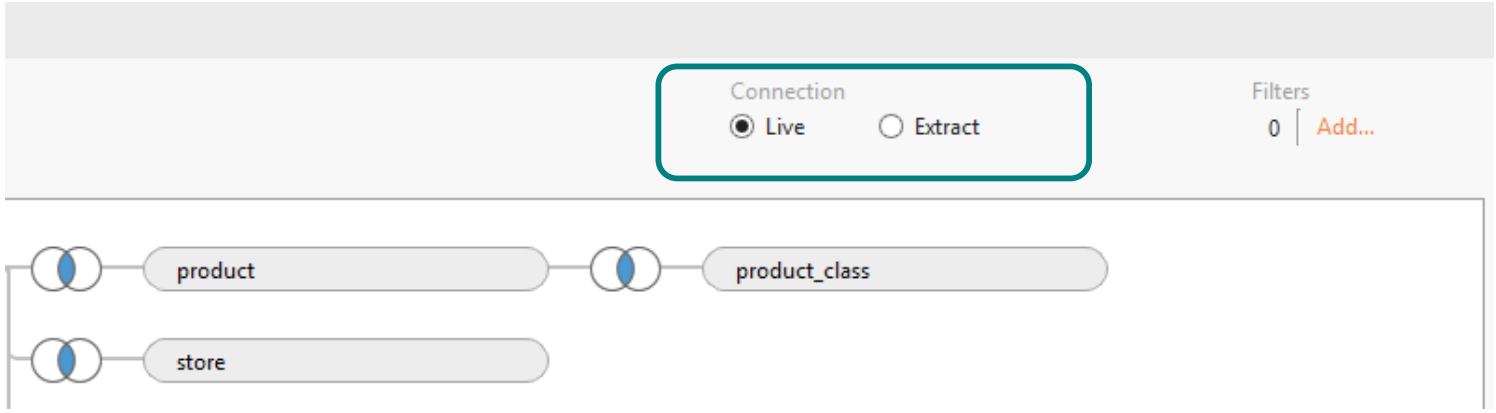

Live VS Extract (2)

A live connection is advised when

- The data source offers high performances (e.g., a typical DW)
- Data freshness is a must
- The amount of data to be analysed is too high to be stored in local

An extract connection is instead advised when

- The data source does not offer fast responses to analytical queries
- The data source must not be overloaded with analytical queries (e.g., an operational data store)
- Off-line data is needed (e.g., performing analyses without a connection to the data source)

From the DFM to Tableau

Tableau offers limited support to directly represent DFM structures. Specifically, Tableau cannot directly represent

- Hierarchies with branches, shared hierarchies and convergences
  - *Workaround*: linearise with attributes duplication
- Descriptive attributes
  - *Workaround*: use a standard field placed outside of hierarchies
- Aggregation constraints
  - It is not possible to force the user to only use sensible aggregation operations (however, proper documentation is usually enough to guide the user)

From the DFM to Tableau – Example

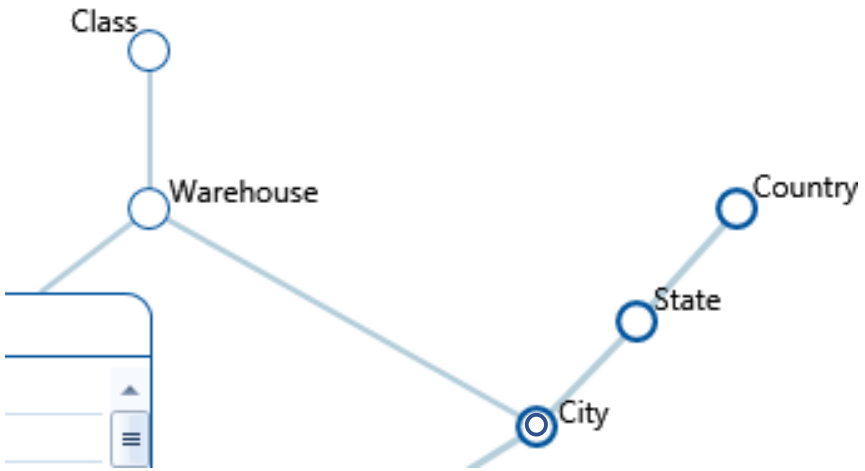

- Warehouse
  - W. Class
    - Class
    - Warehouse (Class)
  - W. Location
    - W. Country
    - W. State
    - W. City
    - Warehouse (Loc)

Exercises - Part II

First queries

Exercise 1

Visualise, through a bar chart, the total sum of *Store Sales* for each *S. State*

- Which one is the state with the highest sales?

Apply a drill-down operation to show the sales at the *S. City* level

- Are there cities whose sales are much lower than the others'?

How many stores (*Store*) are there in each *S. State*? In each *S. City*?

- Tip: Use the *COUNTD* aggregation
- Would it be reasonable to say that cities with fewer stores also have lower total sales?

Visualise the sales for each city (as done before) and represent the number of stores through the color property

Exercise 2

Given the last visualisation of Exercise 1, assign the *Type* field to the color property

- Can you notice any interesting pattern?

Visualise the total (SUM) sales for each *Type*

- Is the result surprising / expected?

Assign the number of stores to the properties color and label

Exercise 3

Visualise, through a line chart, the monthly sales trend

- Any interesting pattern?

Split the previous chart by *S. State* (i.e., an axis for each state)

- Does the previous pattern hold for each state?
- Tip: by default all axes have the same range, if you wish to change this behaviour: right click on any axis > *Edit Axis* > Select *Independent axis…*

Given the previous result, visualise the impact of each *Family* on the total sales while still showing the monthly trends

- Which type of visualisation is best suited?
- Tip: try to assign the *Family* field to one (or more) mark properties; try also different types of marks
- Tip: to change the type of mark use the drop-down menu on the *Marks* panel

Exercise 4

Visualise all stores and sort them in descending order by their total sales

Add the *Type* attribute to the previous visualisation

Add the number of customers and assign it to the color property

Visualise, in descending order, the total sales for each *Type* and *S. State* (put both fields on the rows)

- Verify that the correct order (by descending total sales) is applied

Exercise 5

Visualise the sales for each *Occupation* (*Customer* dimension) and exclude (i.e., filter out) the tuples whose *Store Sales* value is less than 5

- Tip: use a filter on *Store Sales*

Given the previous visualisation, apply another filter (in addition to the previous one) to exclude all occupations whose total sales is less than 80K

- There are still some occupations with sales less than 80K… Why?
  - Tip: to convert a regular filter into a *context filter*, right click > *Apply to Context*

Exercises - Part III

More queries

Exercise 6

Visualise the top ten customers by total sales

- Tip: a *Top N* filter can be appplied by drag-and-dropping a field into the filters shelf (as any other filter) and by using the *Top* tab (right-click > *Edit Filter…* > *Top*)

Given the previous visualisation, add the *Occupation* field (next to the *Customer* one)

Given the previous visualisation, filter by *Occupation* selecting the *Professional* value

- How many customers are there? Can you show the Top N customers whose occupation is *Professional*?
  - Tip: see Exercise 5

Exercise 7

Visualise the distribution of *Store Sales* values (i.e., how many receipts are there such that the total bill is between 0$ and 1$, between 1$ and 2$, etc.?)

- Tip: a histogram can be easily created from the *Show Me* window
- How can you describe the resulting distribution? Normal, uniform, skewed, …
- Which sales does the bin labeled as 0 contain? Does it include only sales with value = 0?
  - Tip: you can use the *View Data* feature

Without using the *Show Me* window, create a bar chart such that: the average *Store Cost* is on the rows shelf, while on the columns shelf there are the bins of *Store Sales* (bin size = 2)

- Tip: to create a bin, from the *Measures* pane, right click on *Store Sales* > *Create* > *Bins…*
- Does this chart clearly show a correlation between sales and cost? Why?
- Edit the chart so that the correlation (if it indeed exists) between sales and cost is more apparent

Exercise 8

Create a set with the Top 500 customers by total *Store Sales*

- Tip: to create a set, right click on *Customer* > *Create* > *Set*

Put the sum of sales on the columns shelf and the set of customers (created at the previous point) on the rows shelf

- What does this visualisation represent?

Edit the previous visualisation by moving the set of customers on the color property and adding *S. State* on the rows shelf

Exercise 9

Visualise the monthly trend of the total profits (Profit = Sales - Costs) for each *Type*

- Tip: to create a *Calculated Field*, from the main menu (on top), *Analysis* > *Create Calculated Field…*

Create a histogram of the number of sales by customers' age (create age bins of size = 5)

- Tip: function *DATEDIFF* returns the difference between two dates (for each function tableau has a brief description accessible directly from the calculated field editor)
- How can you describe the age distribution of the customers?

Given the previous visualisation, instead of the number of sales, visualise the number of sales over the number of customers (i.e., # sales / # customers)

- Any significant difference with the previous chart?

Exercise 10

For each C eir
correspon

- Tip:

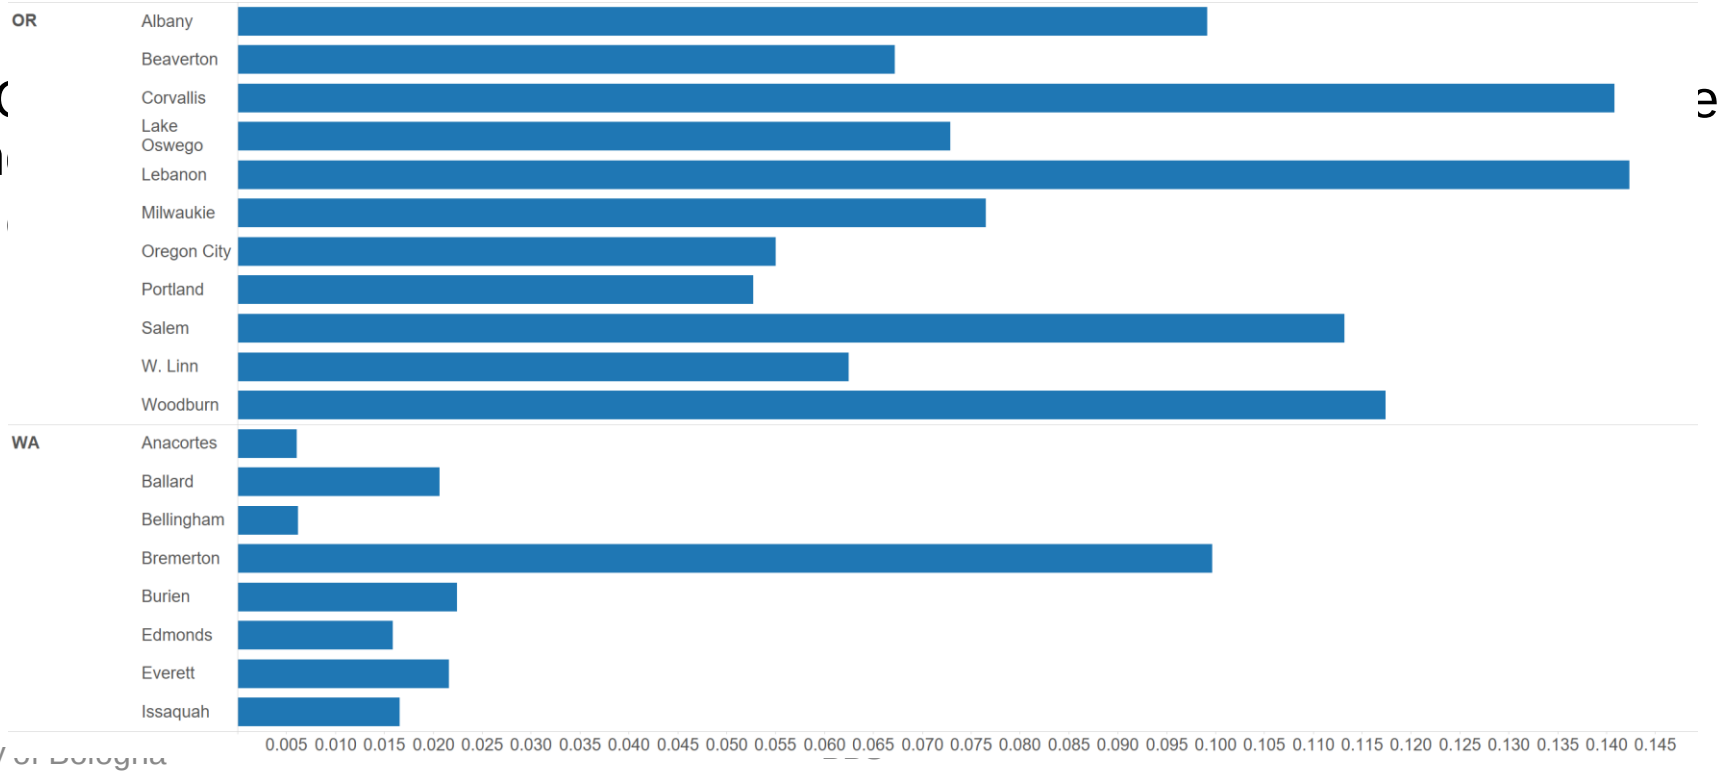

Exercise 11

Compute the distribution of customers over total sales, i.e., show how many customers are there such that the total sales of each one of them is in the ranges 0 – 49, 50 – 99, 100 – 149, …

- The result